

Supporting native_posix on macOS

Zephyr Developer Summit, 2022

Chris Friedt

Embedded SWE, ASIC FW



ABOUT ME

EE / RFE by training, SWE by trade

17 years in mobile, consumer & industrial IoT, SDR

Long time Embedded Linux dev and Open Source advocate

Started in Zephyr as individual contributor in 2019

Picked up by Meta just over 1 year ago

COMMUNITY ROLES

[Linux Plumbers IoT Microconference](#) Organizer

12-14 September, Dublin, Ireland

“IoTs a 4-Letter Word” ([CFP](#))

[Zephyr LTSv2](#) Release Manager

Current LTS is [v2.7.2](#), 145 changes since v2.7.1

Backports for CVE + Bug Fixes for 2 years

6 month overlap between LTS releases ([Process](#))

Zephyr [TSC](#) Member



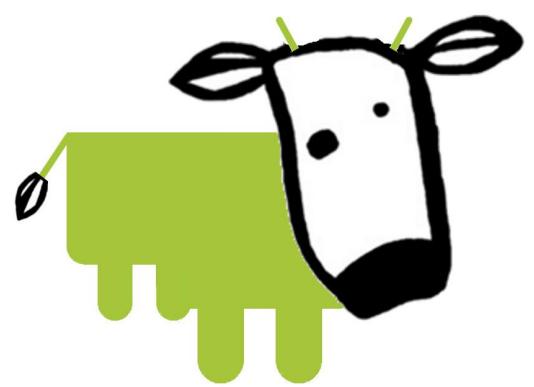
[Summit Spin '22 Jersey](#)





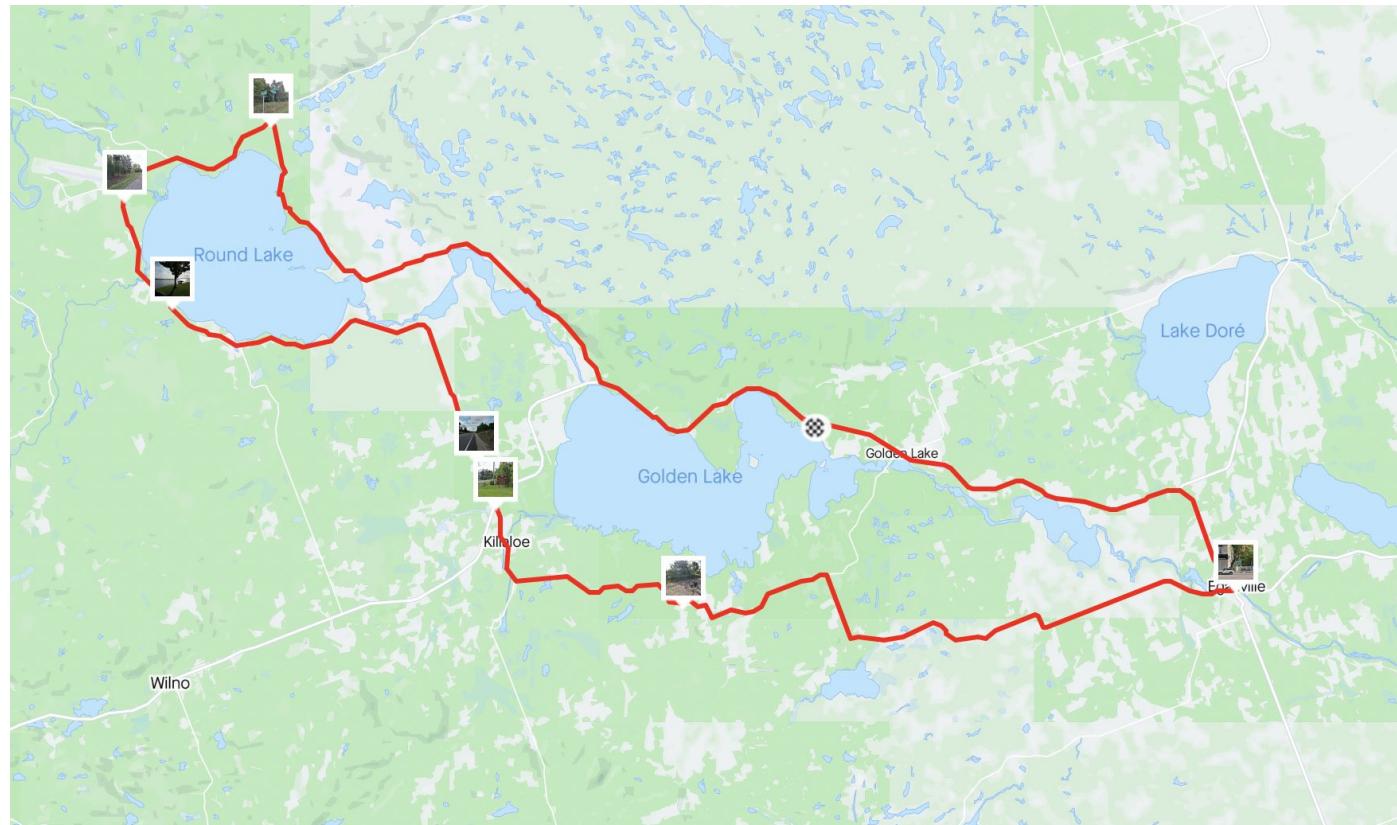
yocto •
PROJECT

posix
bluetooth
spi
cc13xx_launch
greybus
ring_buffer
perf
cc1352r1_launch
release
soc
pinctrl
emul
ieee802154
riscv
atomic
cc26xx
timers
Kconfig
sched
net
dns_sd
pthread
tis
kernel
arm
i2c
mdns
boards
cc1352r_sensortag
Devicetree
gpio
logging
C++
LTS



 **Zephyr®**

Route Map



01 native_posix 🐧

02 macOS 🍎

03 Obstacles 🚧

04 In the garage 🔧

05 The road ahead 🚔

01 native_posix



01 native_posix

What is native_posix?

A convenience boards for testing & development.

A collection of emulated devices and peripherals

The SoC layer contains functionality that would normally be at that layer for a physical chip. For example, special boot code, any special linker sections required, etc.

The POSIX arch layer contains core functionality that, for instance, maps zephyr's threading model to the POSIX model.

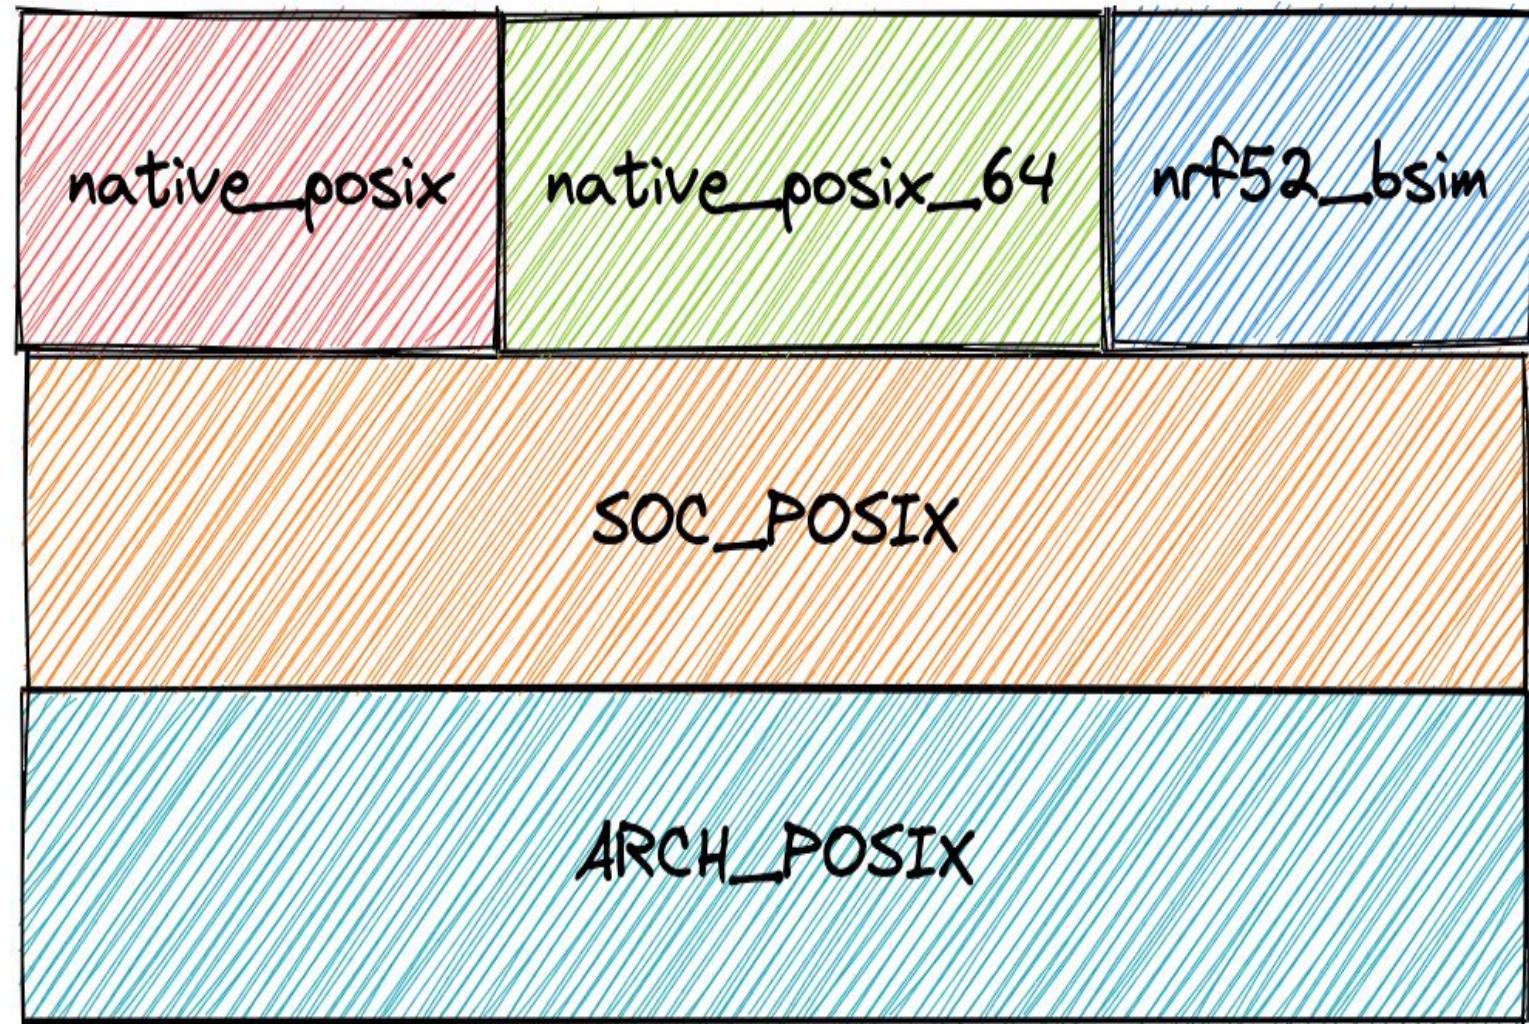
CONFIG_POSIX_API != CONFIG_ARCH_POSIX

CONFIG_POSIX_API is mutually exclusive with CONFIG_ARCH_POSIX

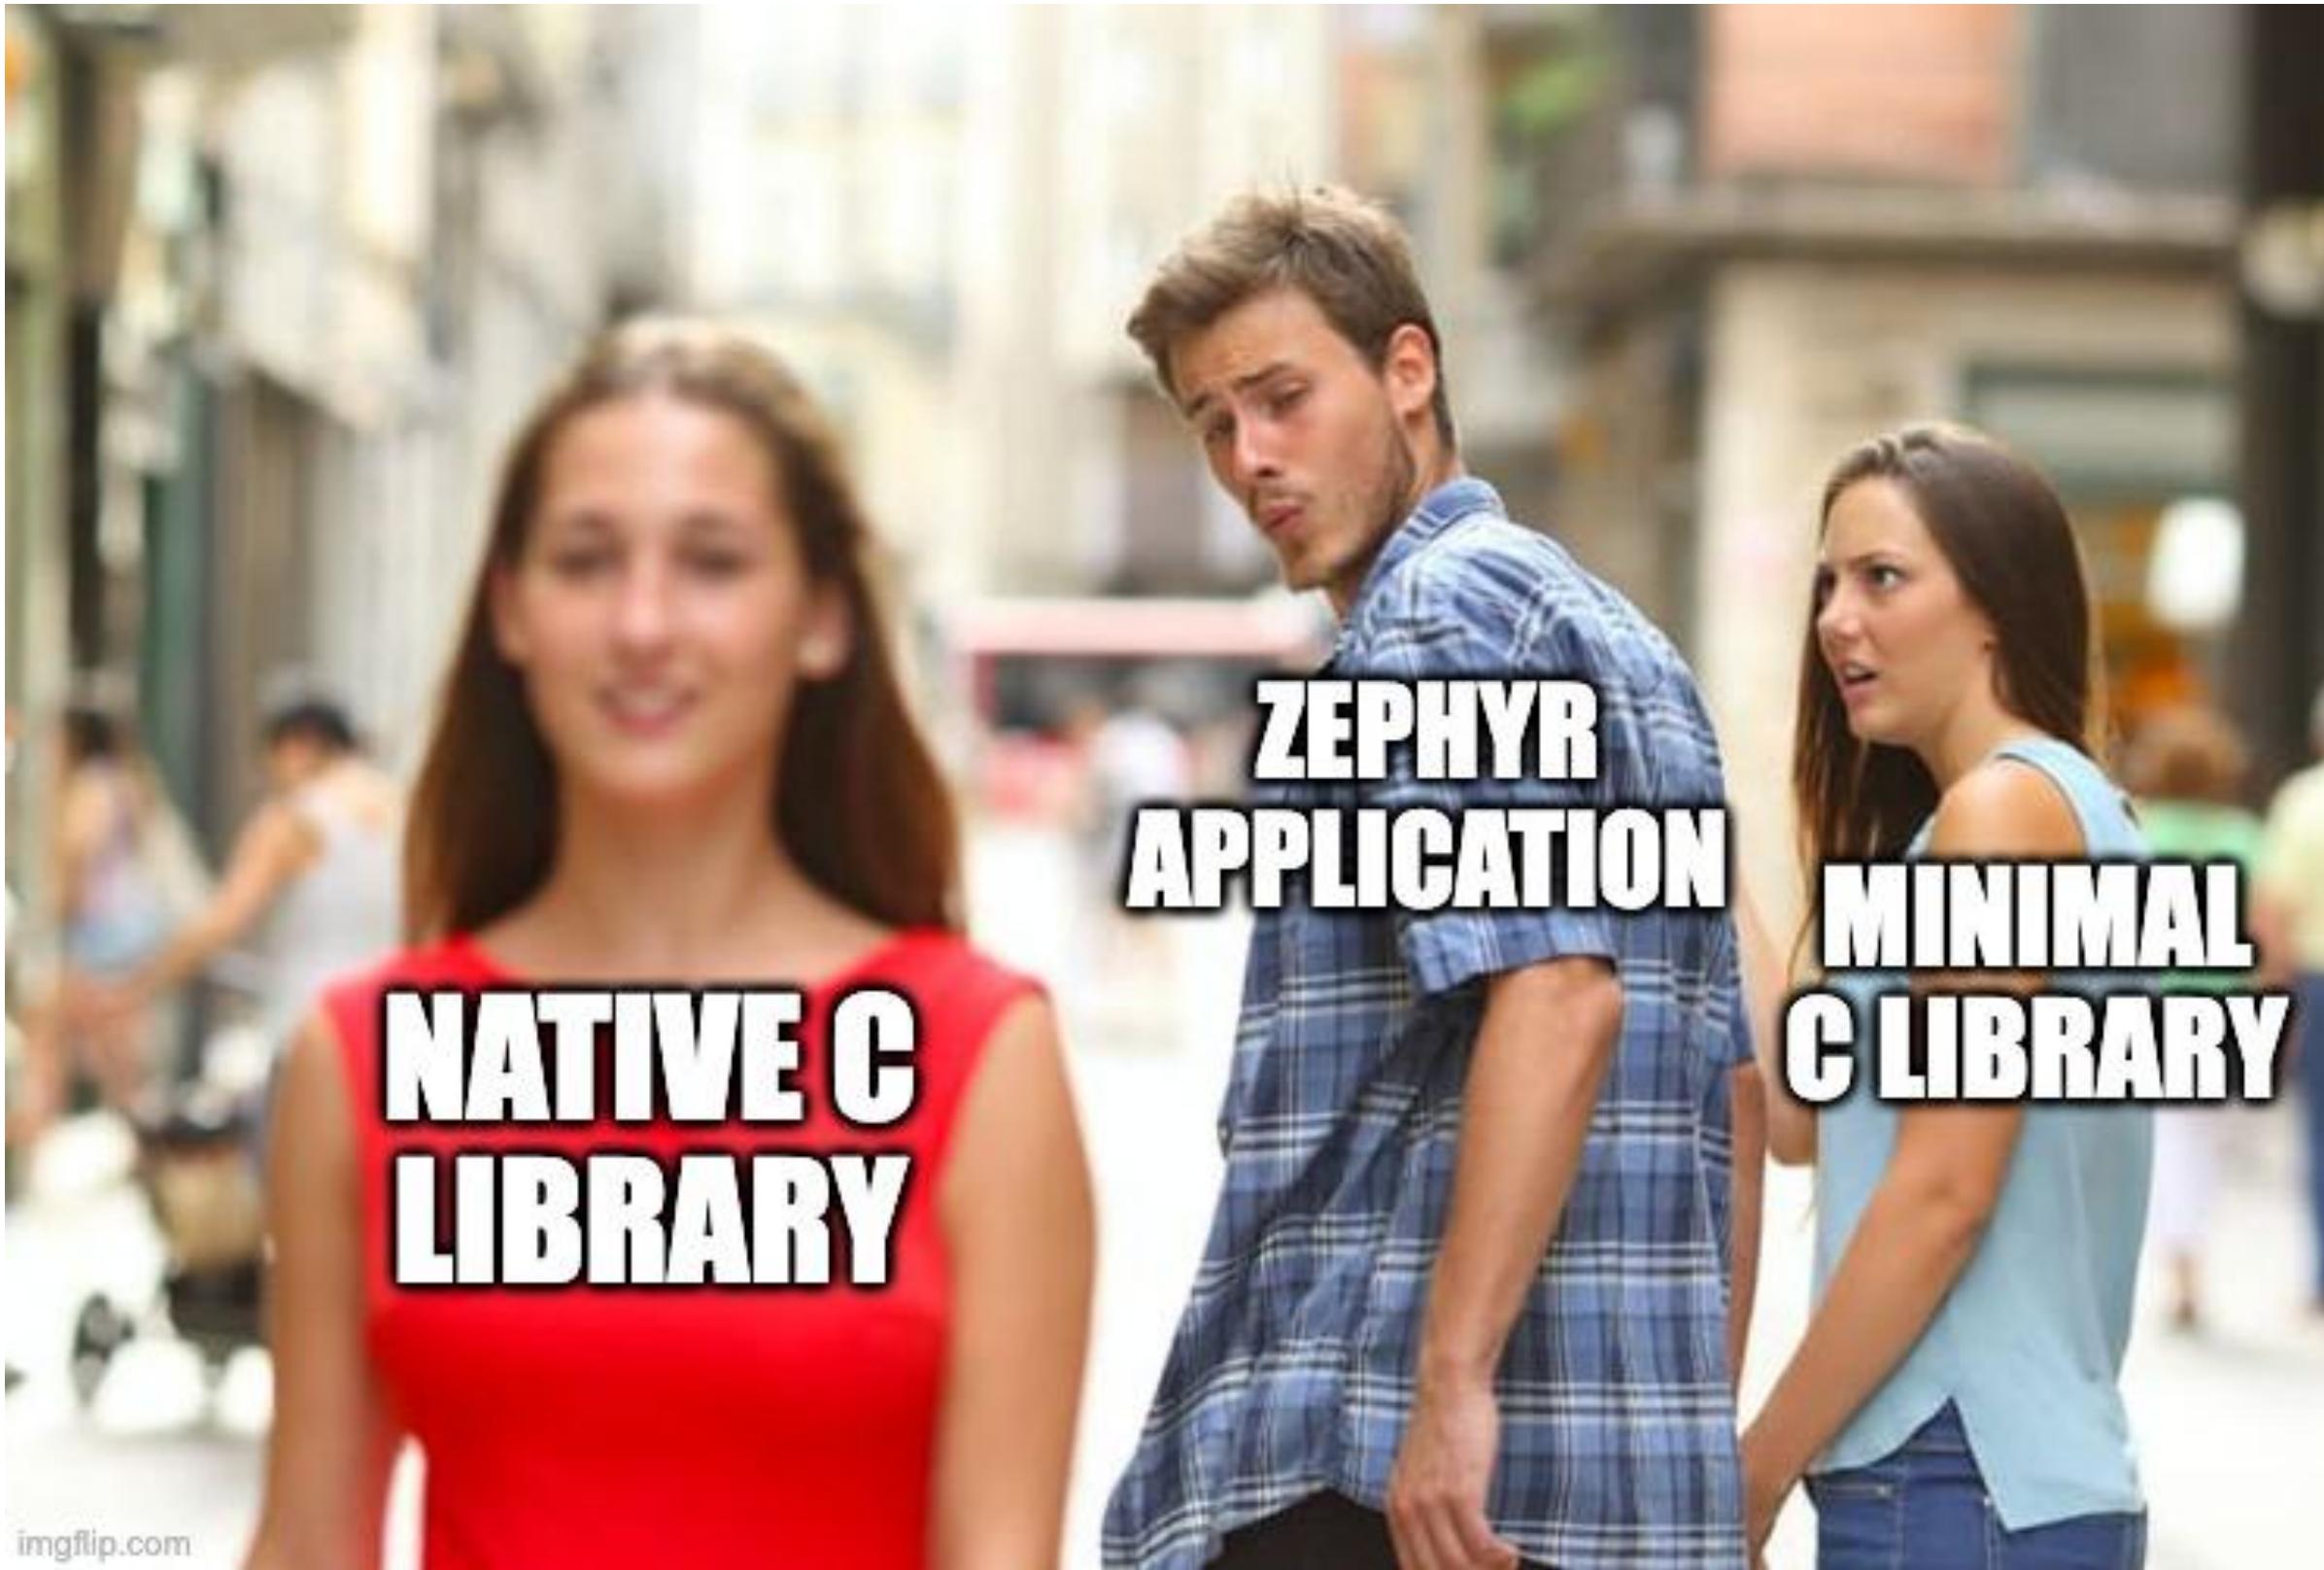
board

soc

arch



01 native_posix



01 native_posix

Why native_posix?

faster development cycle

functionally accurate

more resources (RAM, GHz)

native debugging

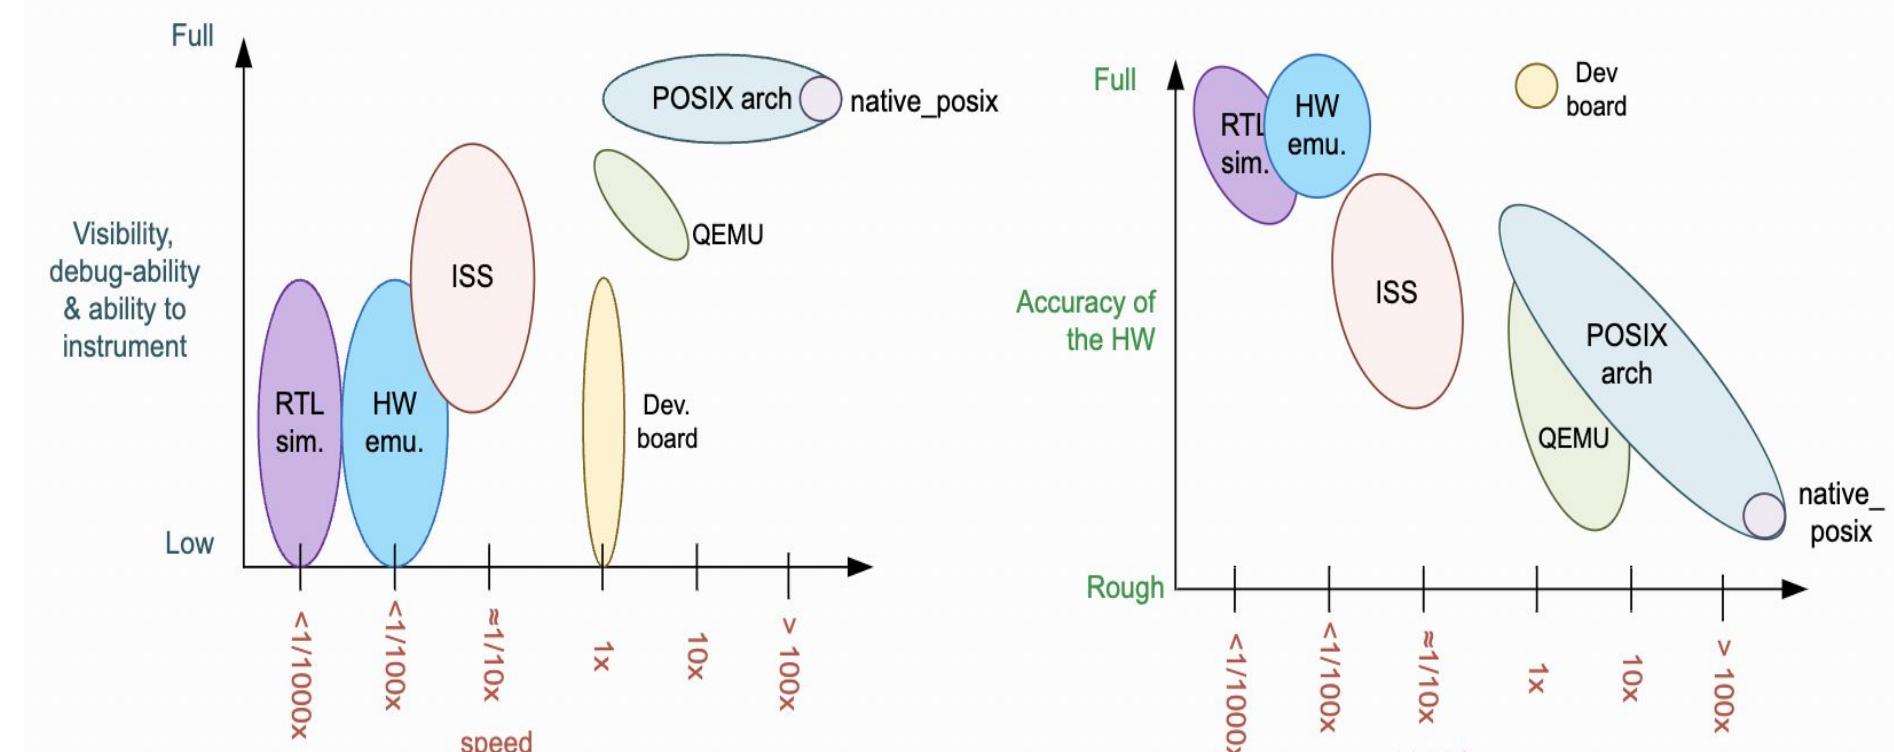
ASan (Address Sanitizer)

UBSan (Undefined Behaviour Sanitizer)

valgrind (memory leak detection)

gcov (code coverage)

perf (flamegraphs!)



ISS: Instruction Set Simulator
Dev. board: Development board
HW emu. : Hardware emulator, e.g. Cadence Palladium
RTL sim. : HW RTL simulations

<https://bit.ly/3z3Zw4f>

01 native_posix

Why native_posix?

Emulated devices and peripherals!

Mentioned several times already this week

Keynote by Simon Glass

Talk by Keith Short

Virtual Talk by Yuval Peress

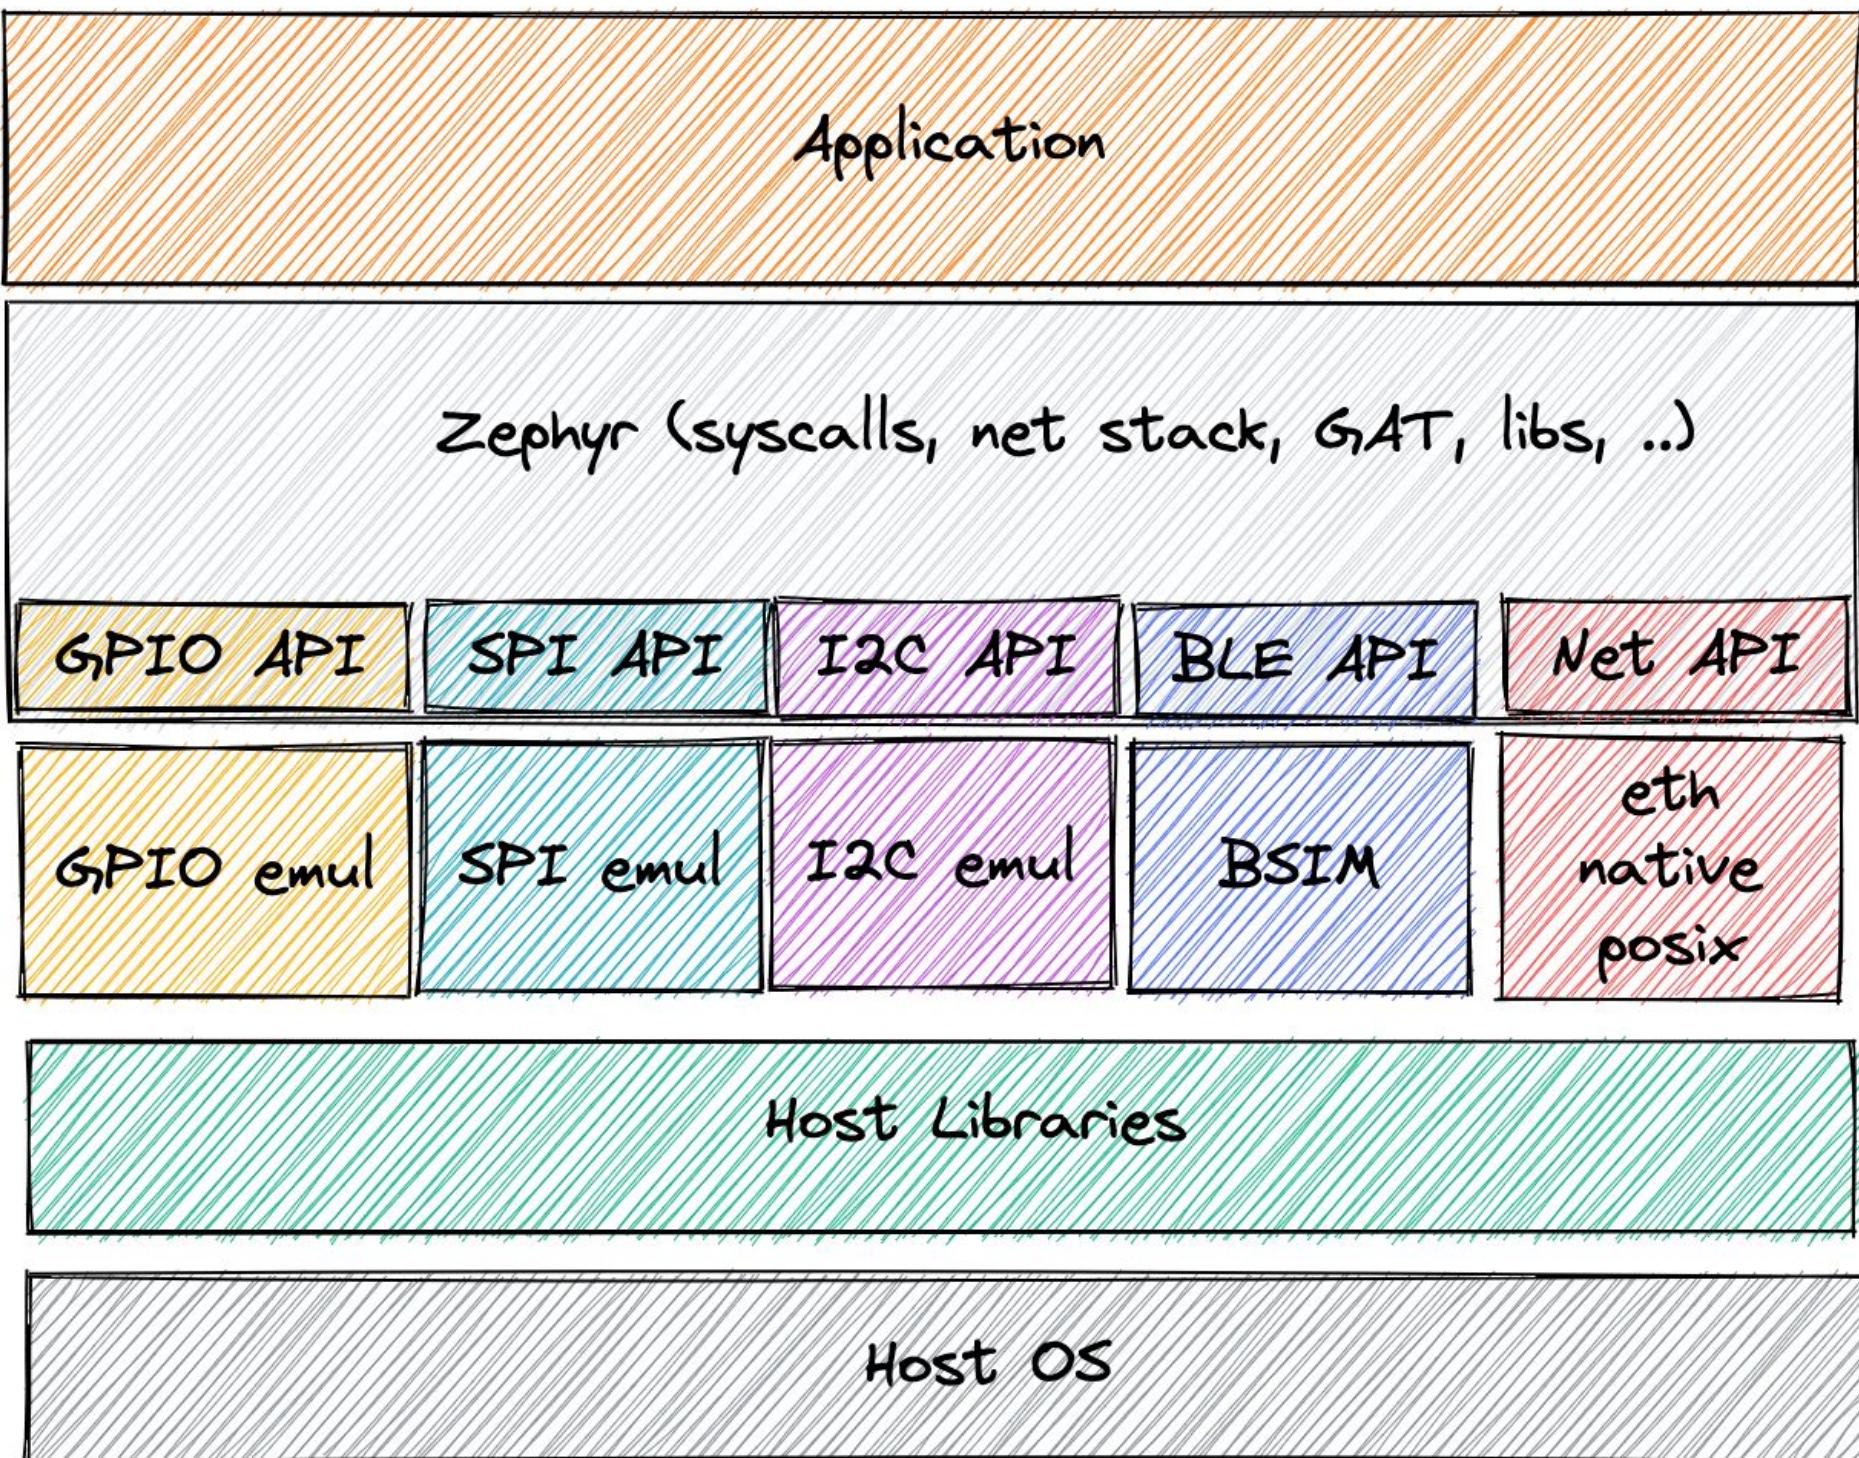
Fail Fast in Testing

Large category of bugs resolved at the API level (HW agnostic)

Testing on real HW takes significantly longer

Test the exact code paths through the app and OS

Inconsistent results with functional & system tests points to the bug



01 native_posix

Why native_posix on macOS?

If we are inclusive in the Zephyr Project, by supporting a POSIX platform other than Linux, then we cast a larger net, and are able to attract the best developers in every community.

Everyone has a preference.

If we can support macOS, which is the only major POSIX OS that does not use ELF, then chances are we can support any POSIX OS.



www.shutterstock.com · 139106594

How native_posix works today

```
NATIVE_TASK(np_add_uart_options, PRE_BOOT_1, 11);
NATIVE_TASK(np_cleanup_uart, ON_EXIT, 99);
```

Selected SW components register a native task with the macro

Function pointer

Level

Priority

Tasks are sorted via a linker script fragment

Specified native tasks are executed prior to Zephyr startup

```

soc > posix > inf_clock > C soc.c > run_native_tasks(int)

221
222 /**
223 * @brief Run the set of special native tasks corresponding to the given level
224 *
225 * @param level One of _NATIVE_*_LEVEL as defined in soc.h
226 */
227 void run_native_tasks(int level)
228 {
229     extern void (*__native_PRE_BOOT_1_tasks_start[])(void);
230     extern void (*__native_PRE_BOOT_2_tasks_start[])(void);
231     extern void (*__native_PRE_BOOT_3_tasks_start[])(void);
232     extern void (*__native_FIRST_SLEEP_tasks_start[])(void);
233     extern void (*__native_ON_EXIT_tasks_start[])(void);
234     extern void (*__native_tasks_end[])(void);
235
236     static void (**native_pre_tasks[]) (void) = {
237         __native_PRE_BOOT_1_tasks_start,
238         __native_PRE_BOOT_2_tasks_start,
239         __native_PRE_BOOT_3_tasks_start,
240         __native_FIRST_SLEEP_tasks_start,
241         __native_ON_EXIT_tasks_start,
242         __native_tasks_end
243     };
244
245     void (**fptr)(void);
246
247     for (fptr = native_pre_tasks[level]; fptr < native_pre_tasks[level+1];
248         fptr++) {
249         if (*fptr) { /* LCOV_EXCL_BR_LINE */
250             (*fptr)();
251         }
252     }
253 }
```

01 native_posix

How native_posix works today

The header file `posix_cheats.h` is included from the command line when preprocessing every source file, and it has two main responsibilities.

The first of those is redefining `main()` to `zephyr_app_main()`. The second is redefining any referenced POSIX functions (e.g. “`zap_open`”). With that, any references to POSIX functions result in undefined symbol errors and the Zephyr binary cannot be linked. In the POSIX SoC and arch layers, the macro `POSIX_NO_CHEATS` disables the default behaviour of “zapping” POSIX symbols.

```
C posix_cheats.h ×  
arch > posix > include > C posix_cheats.h > ...  
43 #ifndef main  
44 #define main(...) zephyr_app_main(__VA_ARGS__)  
45 #endif
```

```
C main.c ×  
boards > posix > native_posix > C main.c > ...  
51 /**  
52 * This is the actual main for the Linux process,  
53 * the Zephyr application main is renamed something else thru a define.  
54 */  
55 int main(int argc, char *argv[]){  
56     run_native_tasks(_NATIVE_PRE_BOOT_1_LEVEL);  
57     native_handle_cmd_line(argc, argv);  
58     run_native_tasks(_NATIVE_PRE_BOOT_2_LEVEL);  
59     hwm_init();  
60     run_native_tasks(_NATIVE_PRE_BOOT_3_LEVEL);  
61     posix_boot_cpu();  
62     run_native_tasks(_NATIVE_FIRST_SLEEP_LEVEL);  
63     hwm_main_loop();  
64     /* This line should be unreachable */  
65     return 1; /* LCOV_EXCL_LINE */  
66 }  
67  
68  
69  
70  
71  
72  
73  
74  
75  
76
```

01 native_posix

How native_posix works today

posix_boot_cpu() spawns a pthread which calls z_cstart()

Multithreading is initialized

The usual Zephyr SYS_INIT() and device init functions are invoked, via ordered iterable section.

arch_swap() calls posix_swap()

Off to the races!

```
C posix_core.c ×
arch > posix > core > C posix_core.c > posix_swap(int, int)
186
187 /**
188 * Let the ready thread run and block this thread until it is allowed again
189 *
190 * called from arch_swap() which does the picking from the kernel structures
191 */
192 void posix_swap(int next_allowed_thread_nbr, int this_th_nbr)
193 {
194     posix_let_run(next_allowed_thread_nbr);
195
196     if (threads_table[this_th_nbr].state == ABORTING) {
197         PC_DEBUG("Thread [%i] %i: %s: Aborting curr.\n",
198                 threads_table[this_th_nbr].thead_cnt,
199                 this_th_nbr,
200                 __func__);
201         abort_tail(this_th_nbr);
202     } else {
203         posix_wait_until_allowed(this_th_nbr);
204     }
205 }
```

02 macOS

02 macOS



macOS History

Macintosh Operating System (v9 in 1999)

“The Best Operating System Ever”

Cooperative multitasking

Mac OS X (2000 to 2016)

Built on technology developed by NeXT, CMU

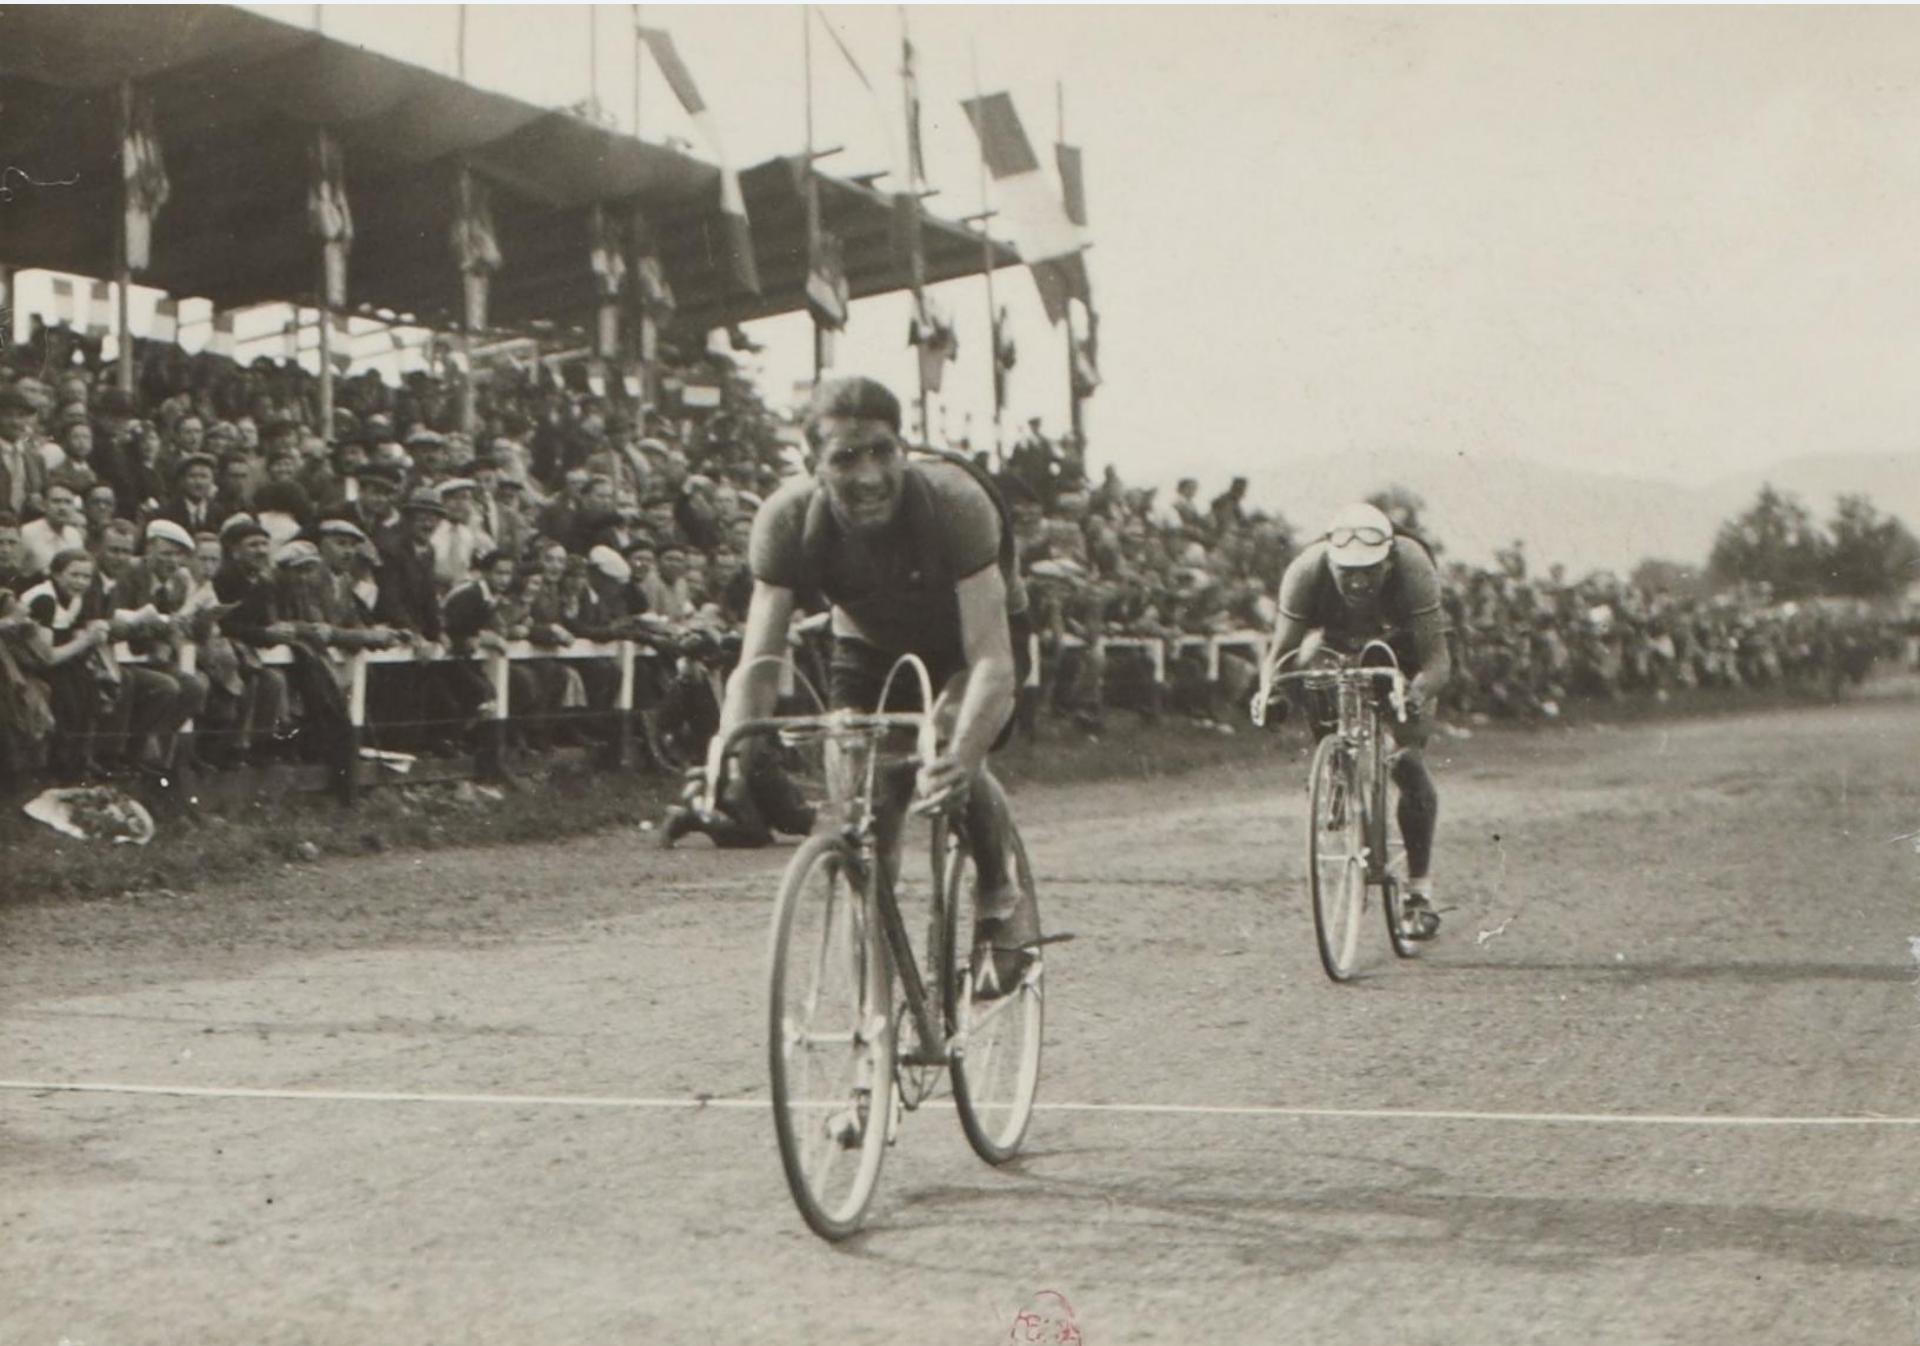
XNU kernel (4.3 BSD kernel + Mach)

Hybrid (monolithic and microkernel components)

FreeBSD + Apple utils, NeXTSTEP, etc

POSIX support, X11, etc

macOS (2016 to present)



Linux vs macOS

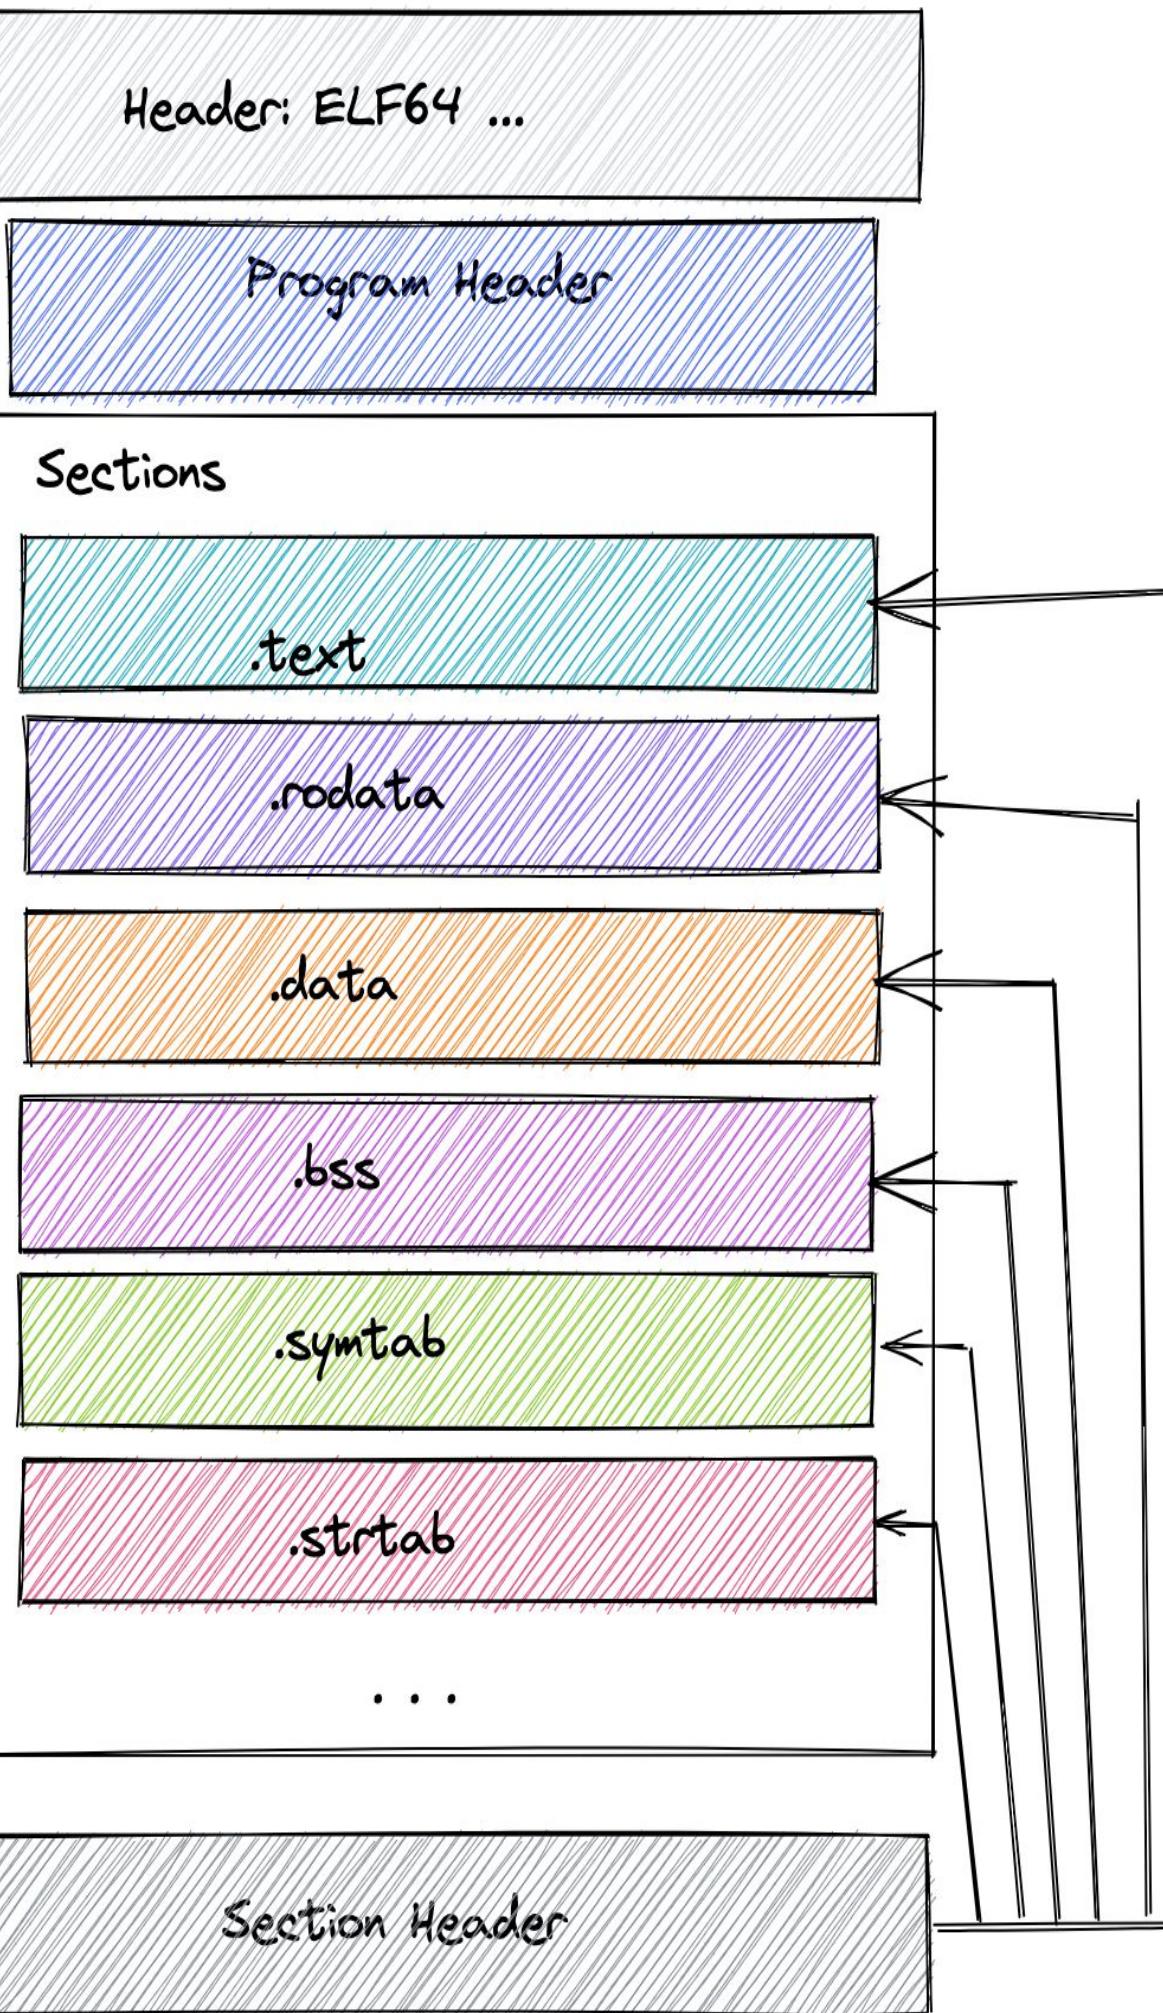
	Linux	macOS
POSIX	Yes! (GNU)	Yes! (BSD)
Kernel	monolithic	~microkernel
Executable Format	ELF	Mach-O
Main library	libc.so.6	~libSystem.B.dylib
POSIX system calls	SWI + syscall nr + regs	fn calls + msg passing

02 macOS

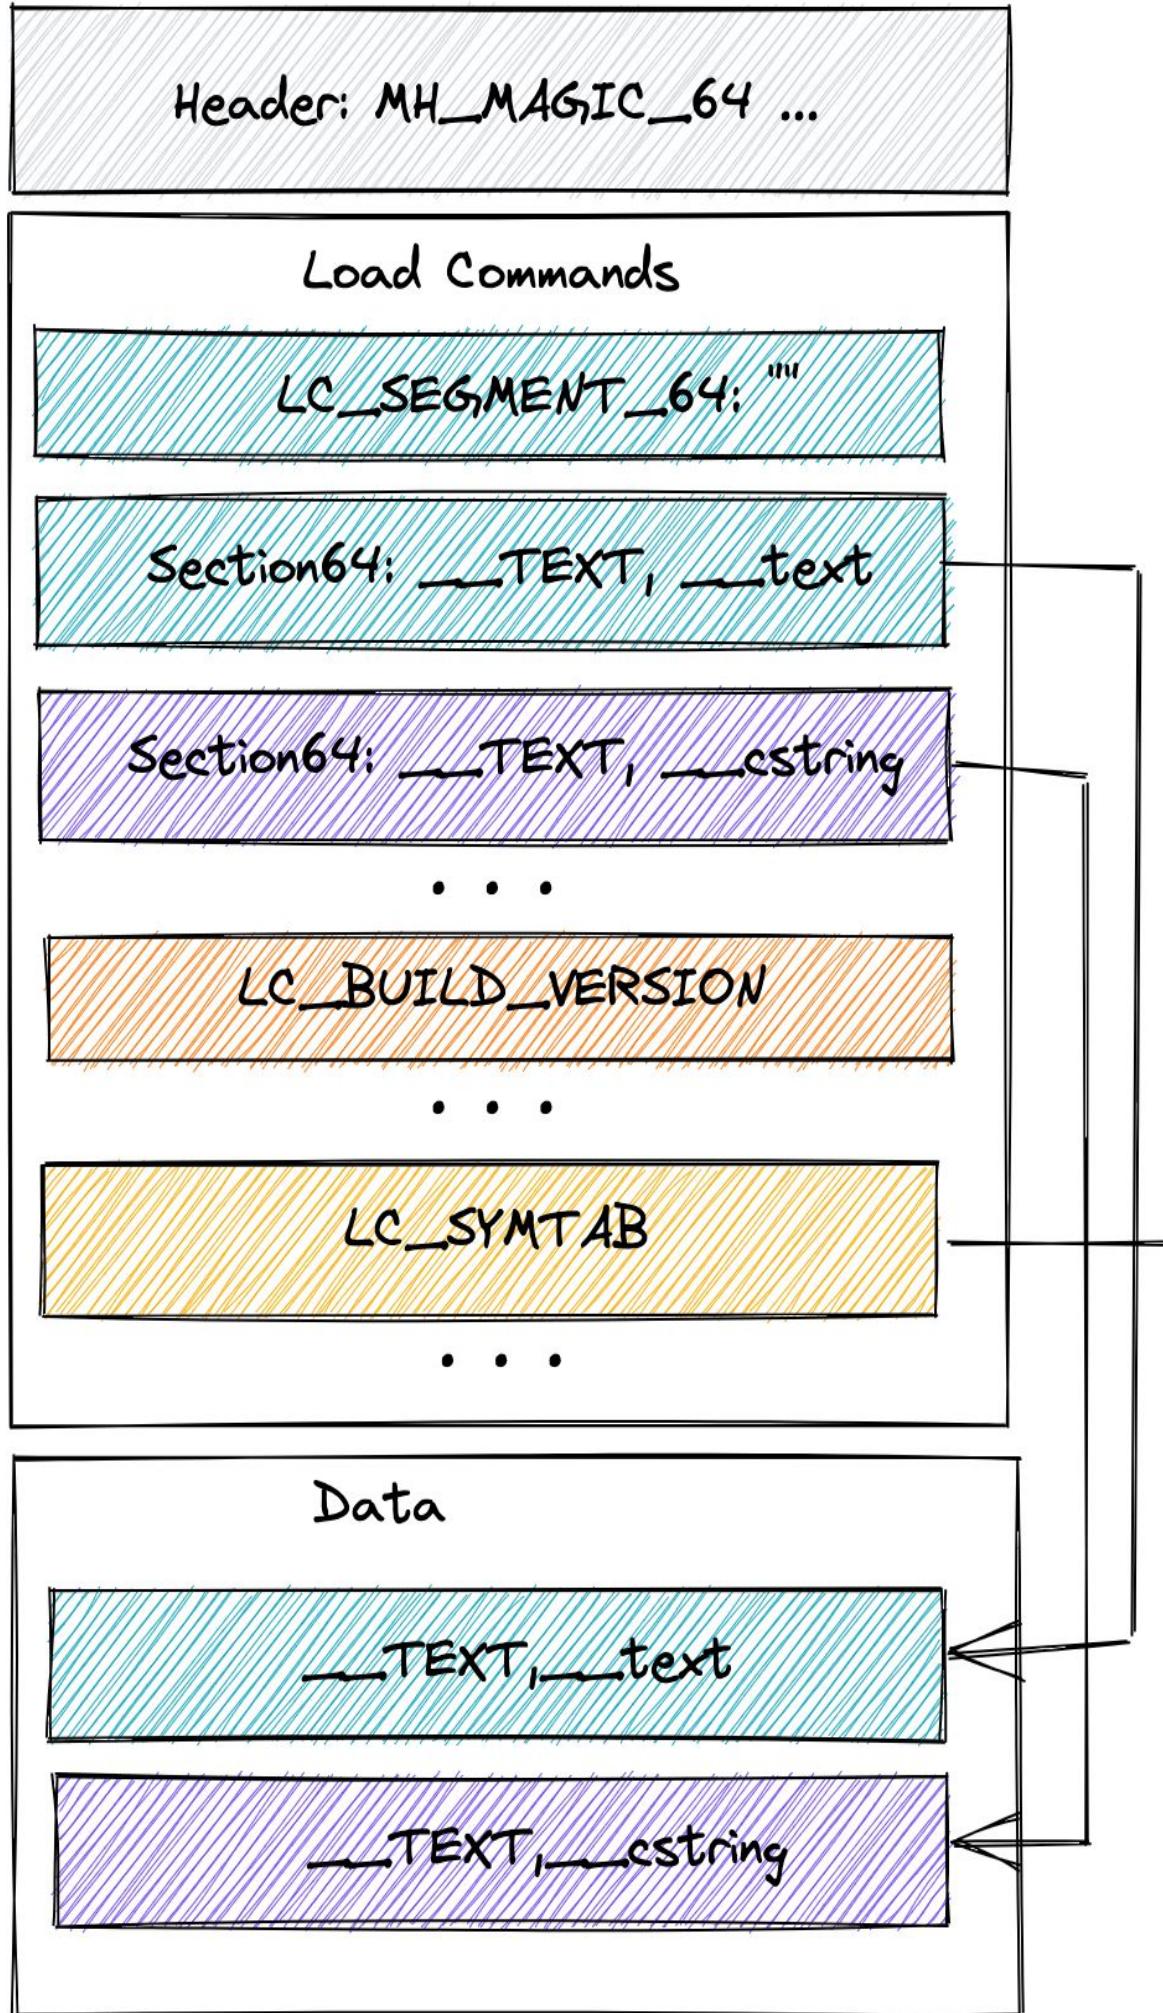
ELF vs Mach-O

Both binary object file formats achieve the same purposes. Mach-O is command oriented whereas ELF is section oriented. Mach-O has a header, ELF has a header. Notably, in Mach-O, each section is associated with a segment, and sections often hold a reference to a distant region within Data, whereas ELF sections are mostly atomic. The section header is also clearly at the end in ELF, and ELF uses back references rather than forward references.

ELF Binary File Format

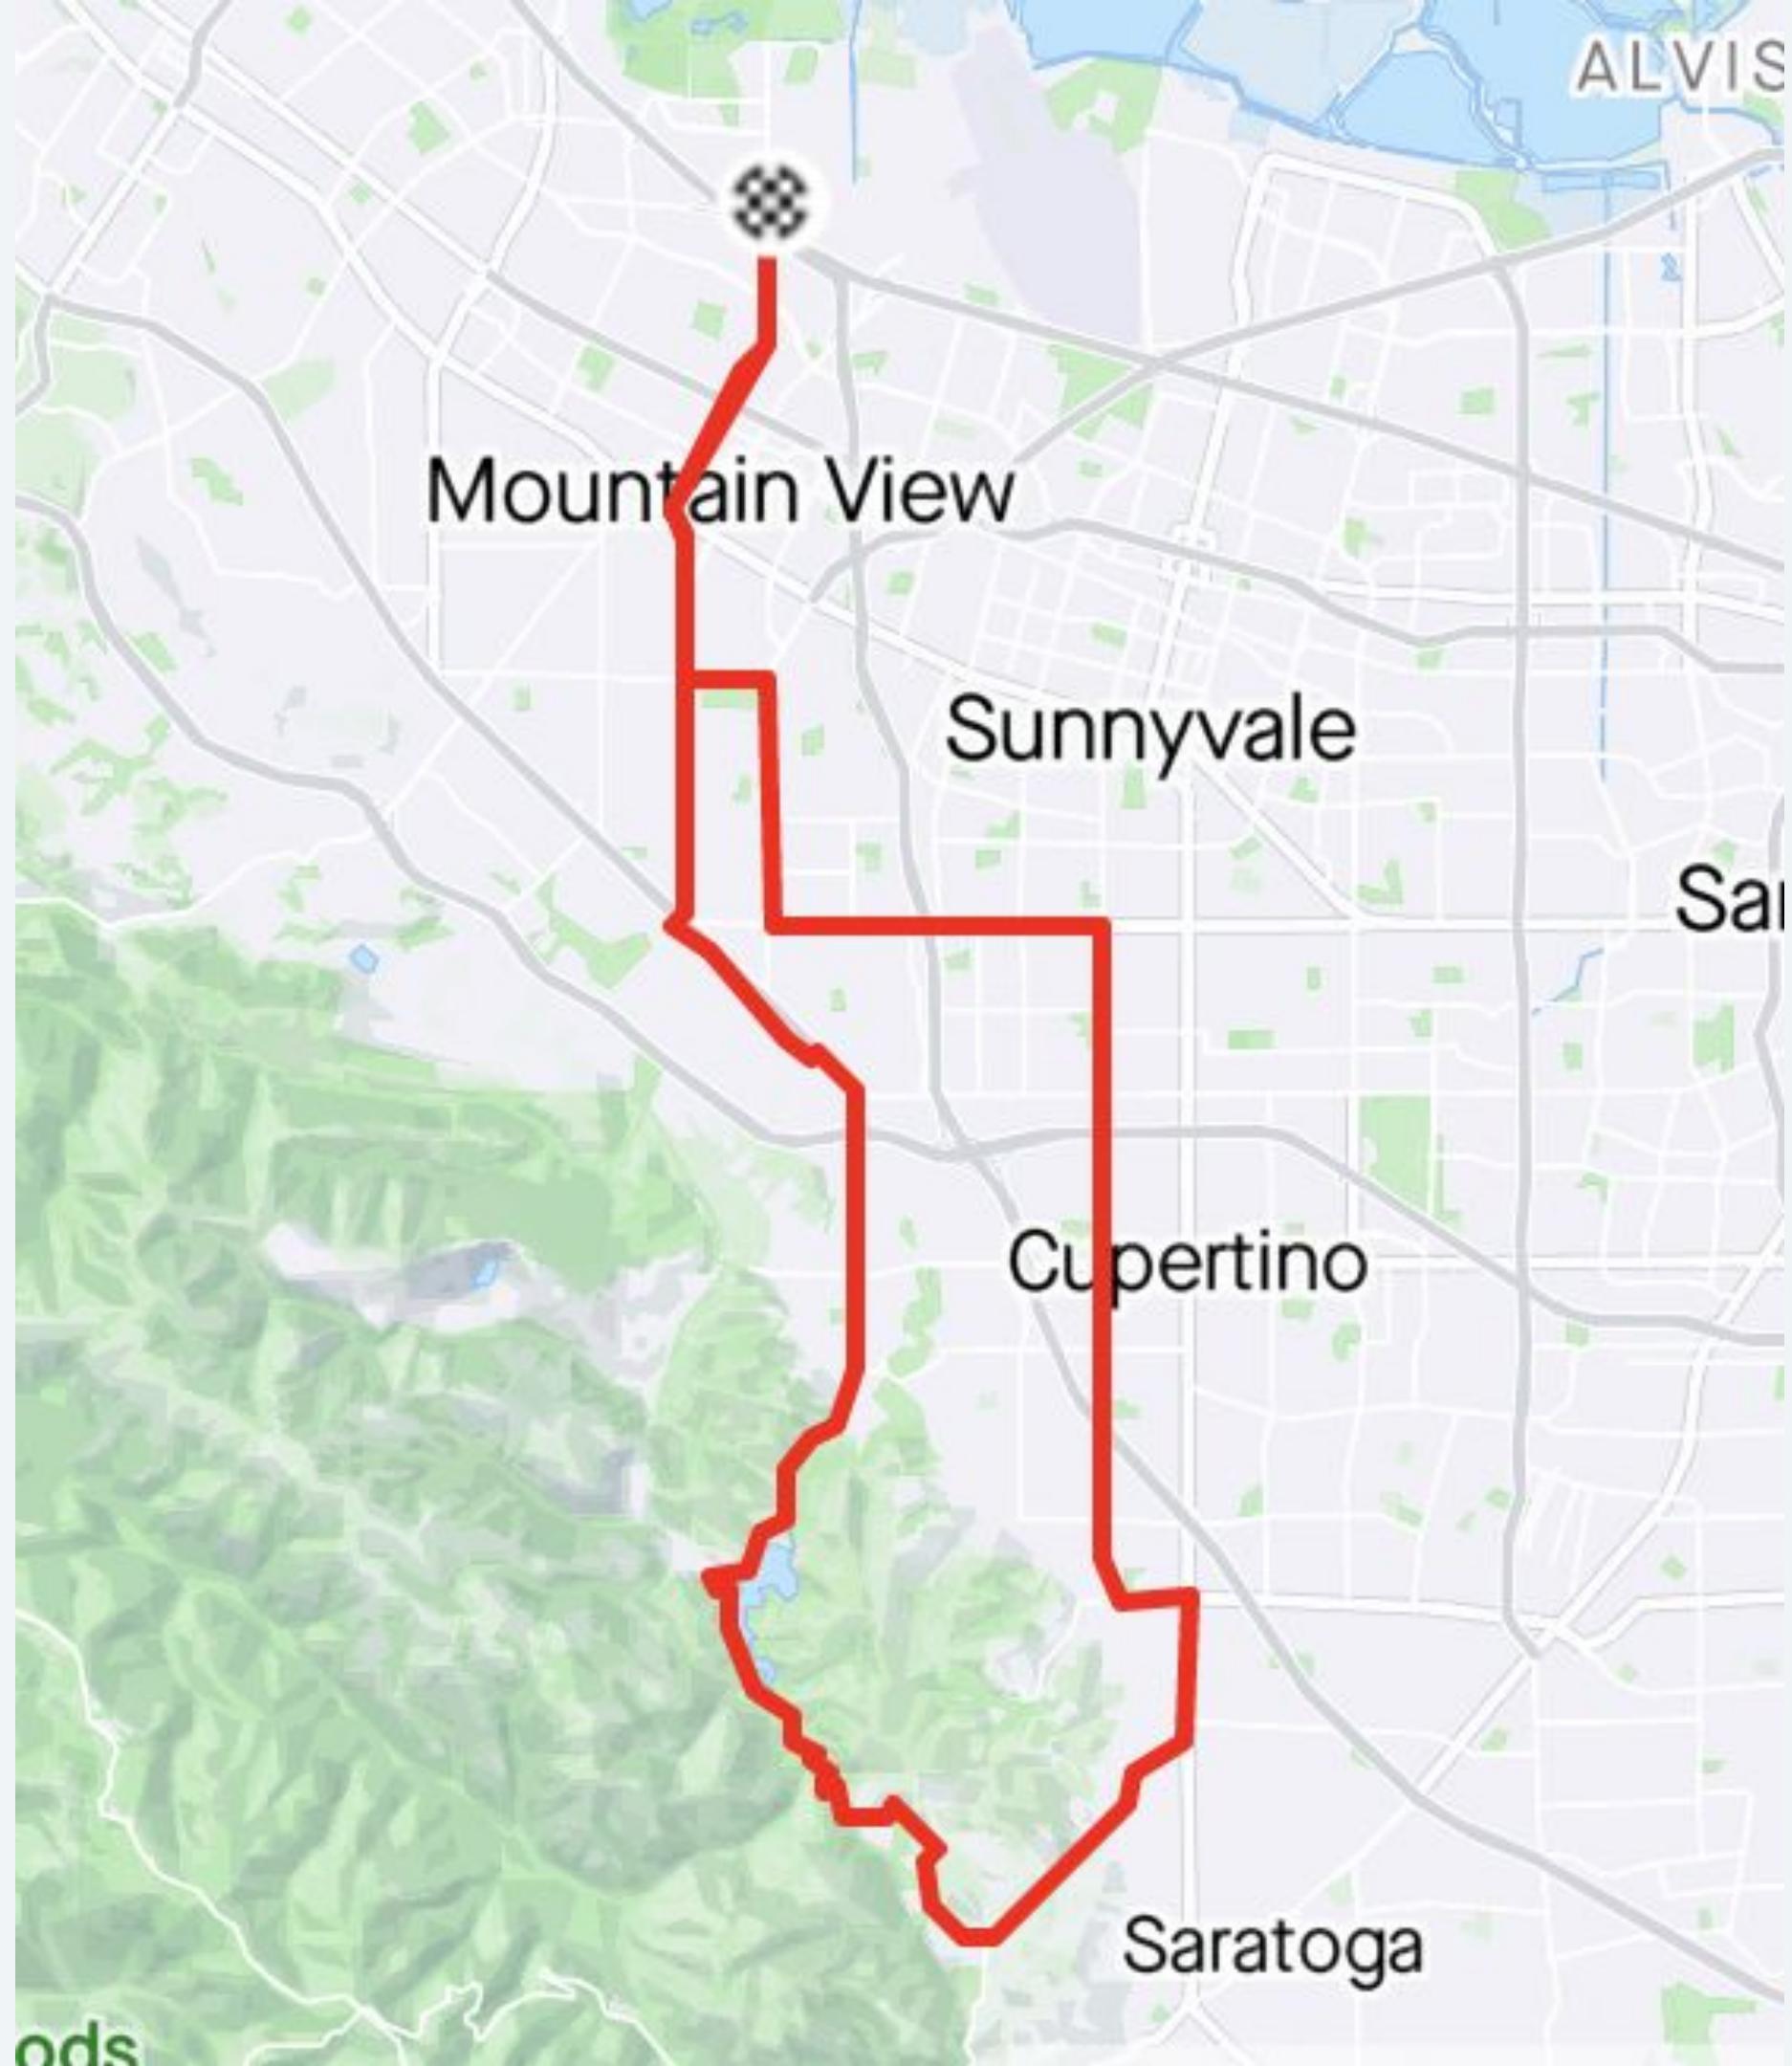


Mach-O Binary File Format



Segment Mapping

ELF	Mach-O
.text	__TEXT, __text
.rodata	__DATA, __const
.bss	BSS, __bss, BSS, __common
.data	__DATA, __data
.got	__DATA, __got



03 Obstacles



The Mach-O Binary Format

Only executable format supported on macOS

Does not support weak symbol visibility

```
warning: '__weak' only applies to pointer types; type here is 'void'  
duplicate symbol '_sys_clock_set_timeout'
```

Does not support symbol aliases

Requires a segment name for every section name

```
error: argument to 'section' attribute is not valid for this target:
```

```
mach-o section specifier requires a segment and section separated by  
a comma
```

Segment and section names must be $1 \leq n \leq 16$ characters

```
error: argument to 'section' attribute is not valid for this target:
```

```
mach-o section specifier requires a section whose length is  
between 1 and 16 characters
```



ELF Hacks

Zephyr's post-compile stages use ELF-hacks to generate constant data

gen_offset_header.py (_OFFSET and _SIZEOF => .h)

gen_handles.py (device handles)

Great for reducing binary size

Mach-O tooling lags far behind ELF tooling (e.g. pyelfutils / machotools)



Toolchain Abstractions

Relatively minor

The .type directive is not supported in assembly

Apple Id does not support --whole-archive

Apple Id does not support GNU linker scripts

does support an order file

not entirely clear if order_file is sufficient

GNU linker scripts also generate symbols

No static linkage (all Apple executables are dynamic)

-order_file file

Alters the order in which functions and data are laid out. For each section in the output file, any symbol in that section that are specified in the order file `file` is moved to the start of its section and laid out in the same order as in the order file `file`. Order files are text files with one symbol name per line. Lines starting with a # are comments. A symbol name may be optionally preceded with its object file leaf name and a colon (e.g. `foo.o:_foo`). This is useful for static functions/data that occur in multiple files. A symbol name may also be optionally preceded with the architecture (e.g. `ppc:_foo` or `ppc:foo.o:_foo`). This enables you to have one order file that works for multiple architectures. Literal c-strings may be ordered by quoting the string (e.g. `"Hello, world\n"`) in the order file.

04 In the garage

Prototype

Bit of a dirty hack ([#37123](#))

Simple workarounds for Toolchain abstractions

Use macros to create structs with `__attribute__((constructor))`

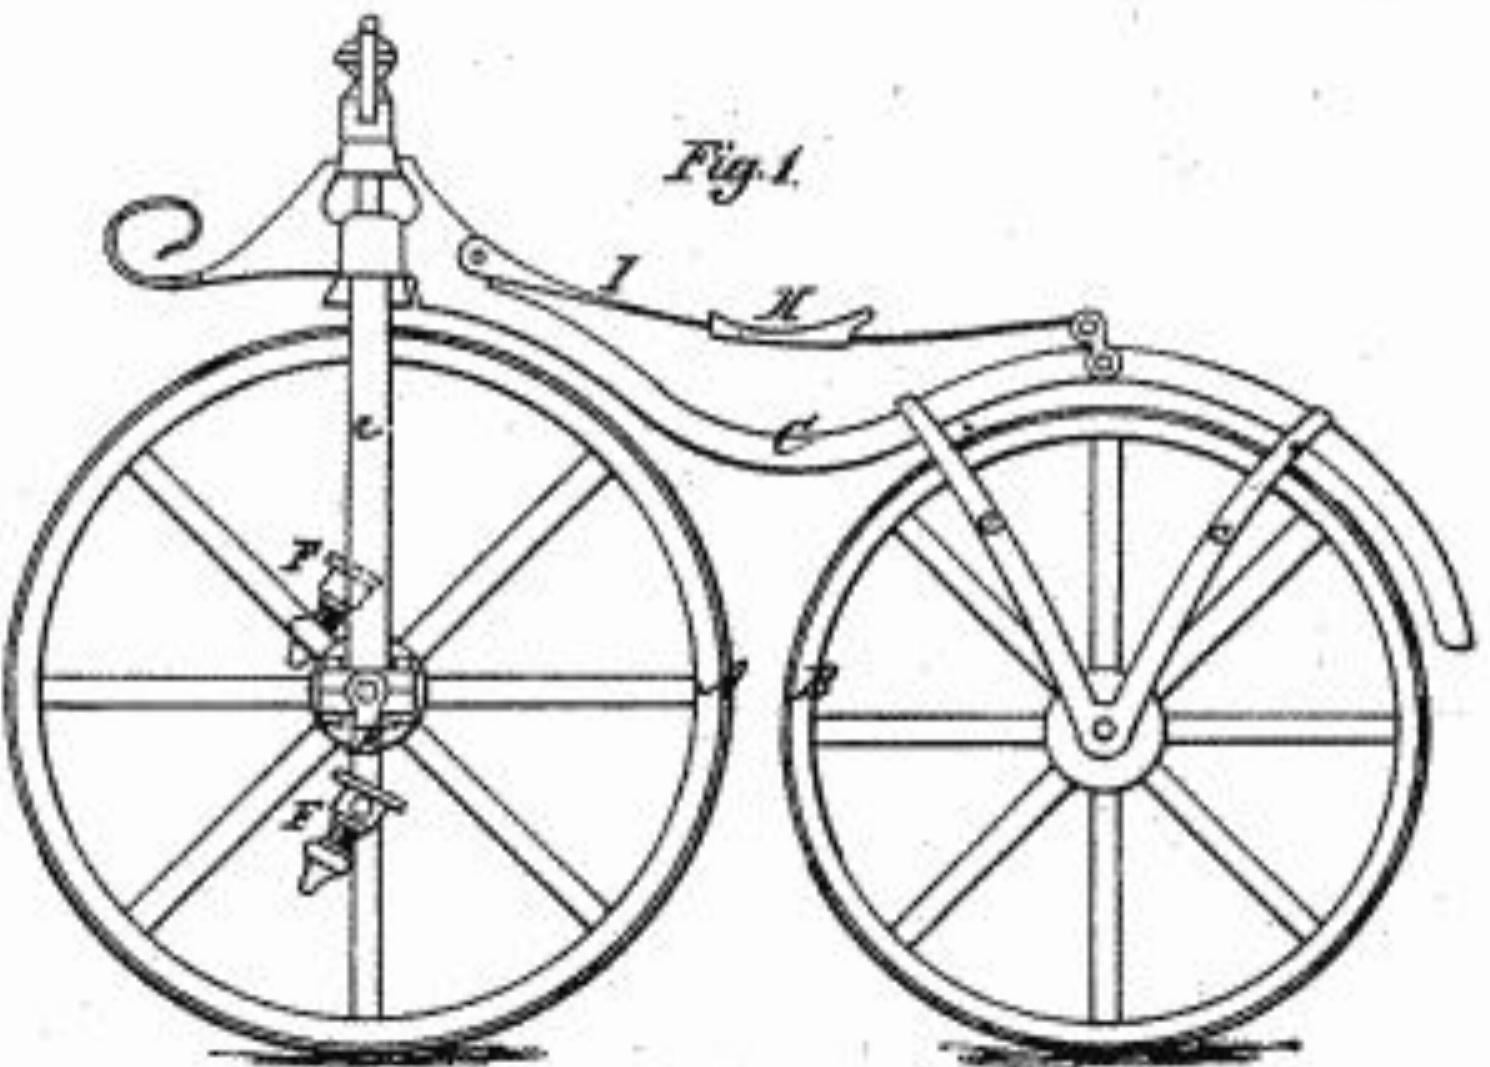
Have each constructor append details to a C++ vector<>

NATIVE_TASK() and SYS_INIT() only

Too many iterable sections to be feasible

Required manually tweaking section names / dumping iterable sections into 1 segment, section

Hard to pedal



Spare Parts Build

Basically the same as the prototype

Try to use Homebrew GCC

Same linker issues (same linker)

Completely ignores weak attribute

⇒ symbol collisions

⇒ at least better than UB

Might work for a ride or two but not reliably



Pro Model

Modify LLVM (clang, clang++) ([Changes](#))

-fhash-section-names=N

-fhashed-section-names=outfile

Sha256 of section name

Base64 Encoding

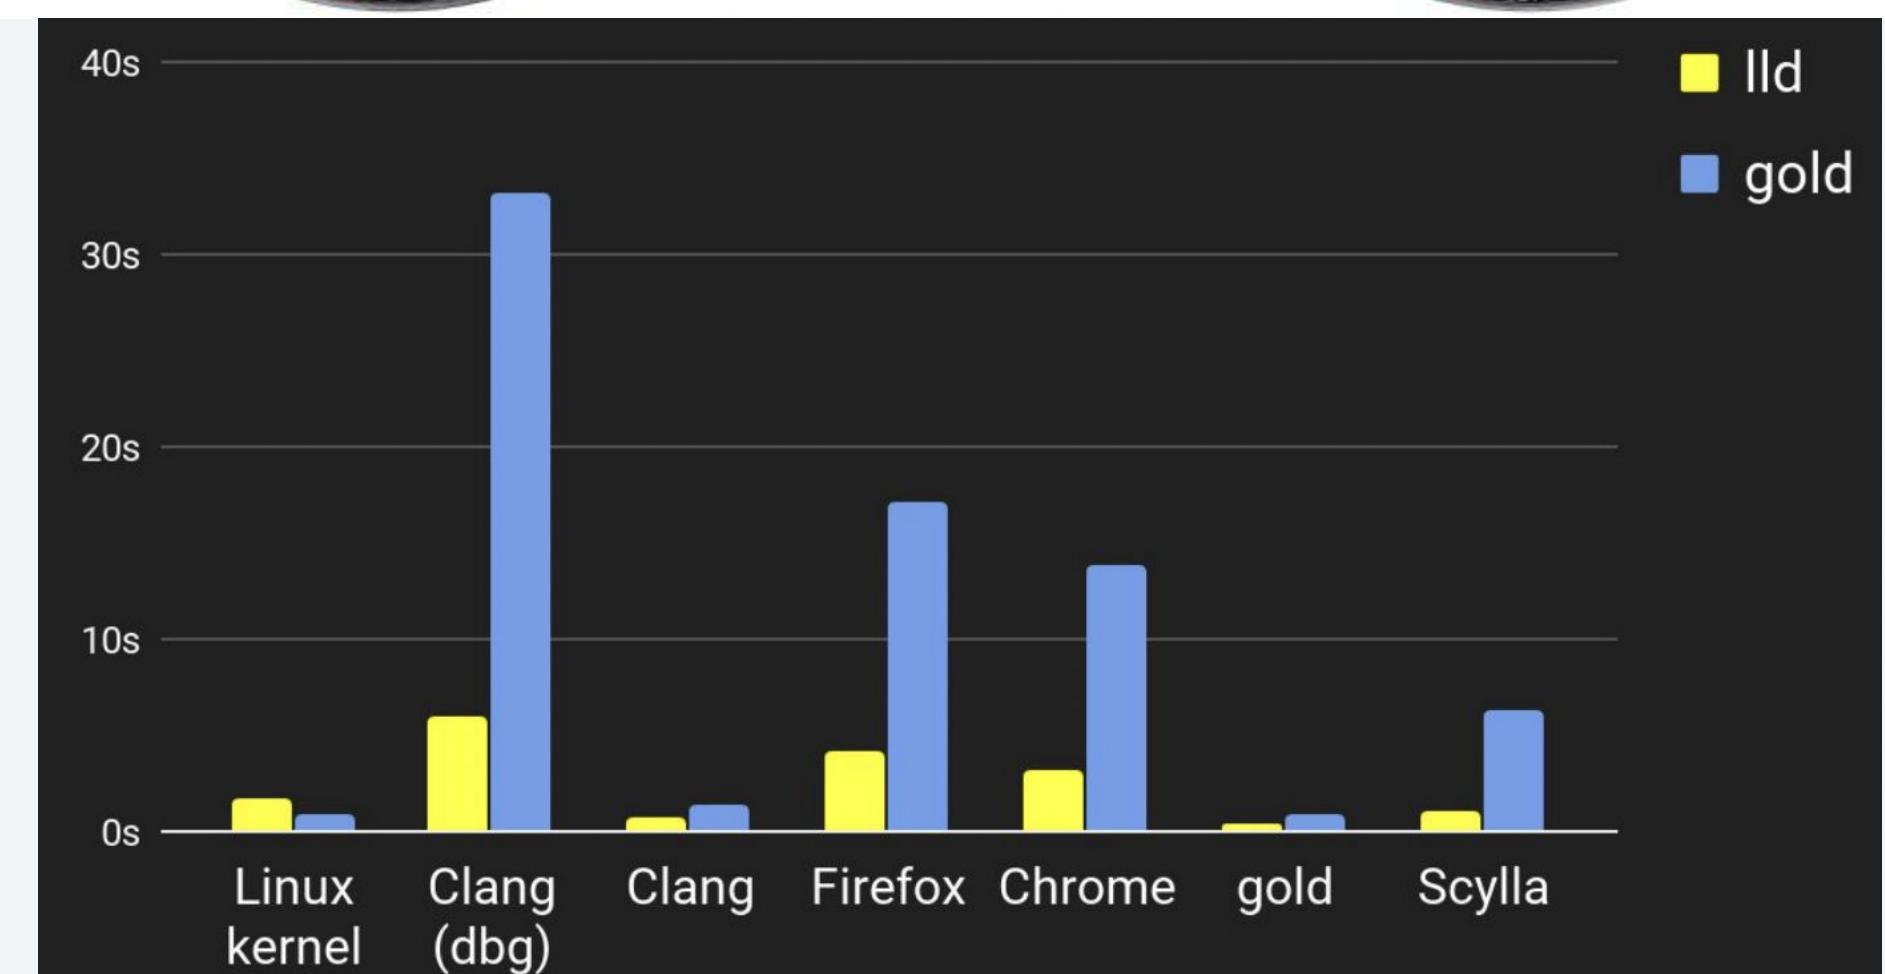
Truncate to N

We likely want LLVM to act as a drop-in replacement for gcc

Including support for GNU Linker scripts on macOS

LLVM's linker, LLD is significantly faster than the GNU linker ([Source](#))

Additional [talk](#): FOSDEM 2019, Peter Smith, Linaro



Promising results

Details in [#40709](#)

Short section names are preserved as-is

Longer section names truncated to 16

Output file shows a clear mapping

```
% objdump --headers build/zephyr/zephyr.exe

build/zephyr/zephyr.exe:      file format mach-o arm64

Sections:
Idx Name          Size   VMA          .k_mutex.static.fdtable_lock    ee505i1iDSJ38lw6
  0 __text        0000cbdc 00000000 .z_init_PRE_KERNEL_199_ PWCrD+v05fLb0TNC
  1 __stubs       000001d4 00000000 .__device_handles_pass1 ISGPhW3oi0+QmD32
  2 __stub_helper 000001ec 00000000 .z_device_PRE_KERNEL_20_         g689RYj/NgI8BJm1
  3 __cstring     00000f55 00000000 .z_init_PRE_KERNEL_20_         KPvctDvcsEzL2GeJ
  4 __const        0000014d 00000000 .z_init_PRE_KERNEL_130_ DSRqIUjHq06d+Nny
  5 __literal8    00000058 00000000 .noinit."ZEPHYR_BASE/kernel/init.c".0   kSBlhTrw6lGfb4bH
  6 __literal16   00000010 00000000 .noinit."ZEPHYR_BASE/kernel/init.c".1   M7LnbayK+Apx90Az
  7 __got         00000090 00000000 .noinit."ZEPHYR_BASE/kernel/init.c".2   2ykHVrz1H6C6zM7x
  8 __const        00000080 00000000 .z_init_PRE_KERNEL_130_ DSRqIUjHq06d+Nny
  9 __la_symbol_ptr 00000138 00000000 .noinit."ZEPHYR_BASE/kernel/mailbox.c".0   AivJBhsGJLg4nCaa
 10 __data        00000460 00000000 .noinit."ZEPHYR_BASE/kernel/mailbox.c".1   0G3gaNFFoLgkTC6u
 11 ee505i1iDSJ38lw6 00000020 00000000 .k_stack.static.async_msg_free yXwPRdgDQM/TZ0qv
 12 X30kufZugPkMRV0u 00000008 00000000 .noinit."ZEPHYR_BASE/kernel/mailbox.c".1   0G3gaNFFoLgkTC6u
 13 gSbov272c0SyitB9 00000008 00000000 .z_init_PRE_KERNEL_130_ DSRqIUjHq06d+Nny
 14 PWCrD+v05fLb0TNC 00000010 00000000 .noinit."ZEPHYR_BASE/kernel/pipes.c".0   k0ons3eUg3xD0NU5
 15 ISGPhW3oi0+QmD32 00000006 00000000 .k_stack.static.pipe_async_msgs      xcc74YkFJp0/l+79
 16 .z_devstate    00000004 00000000 .noinit."ZEPHYR_BASE/kernel/pipes.c".1   66dcYG27Liq2h08v
 17 g689RYj/NgI8BJm1 00000030 00000000 .z_init_PRE_KERNEL_130_ DSRqIUjHq06d+Nny
 18 KPvctDvcsEzL2GeJ 00000010 00000000 .z_init_PRE_KERNEL_130_ DSRqIUjHq06d+Nny
 19 kSBlhTrw6lGfb4bH 00000400 00000000 .noinit."ZEPHYR_BASE/kernel/system_work_q.c".0   U9Ufj7zPekAbDTsr
 20 2ykHVrz1H6C6zM7x 00000800 00000000 .z_init_POST_KERNEL40_   oinUk81sKcoS+CT2
 21 M7LnbayK+Apx90Az 00000100 0000000100015228 DATA
 22 DSRqIUjHq06d+Nny 00000040 0000000100015328 DATA
 23 yXwPRdgDQM/TZ0qv 00000030 0000000100015368 DATA
 24 AivJBhsGJLg4nCaa 00000050 0000000100015398 DATA
 25 0G3gaNFFoLgkTC6u 00000640 00000001000153e8 DATA
 26 xcc74YkFJp0/l+79 00000030 0000000100015a28 DATA
 27 k0ons3eUg3xD0NU5 00000050 0000000100015a58 DATA
 28 66dcYG27Liq2h08v 000004b0 0000000100015aa8 DATA
 29 oinUk81sKcoS+CT2 00000010 0000000100015f58 DATA
 30 U9Ufj7zPekAbDTsr 00000400 0000000100015f68 DATA
 31 __bss          00000555 0000000100016370 BSS
 32 __common       00000239 00000001000168c8 BSS
```

Assisted Riding

[Macho.js](#) a useful learning and visualization tool for exploring Mach-O files

Modifying gen_offset_header.py relatively easy to do in python wrapping around nm ([Source](#))

Some progress with [machotools](#) (Mach-O variant of [pyelfutils](#))

Strategy: make a “toy” linker in Python called LowD ([Source](#))

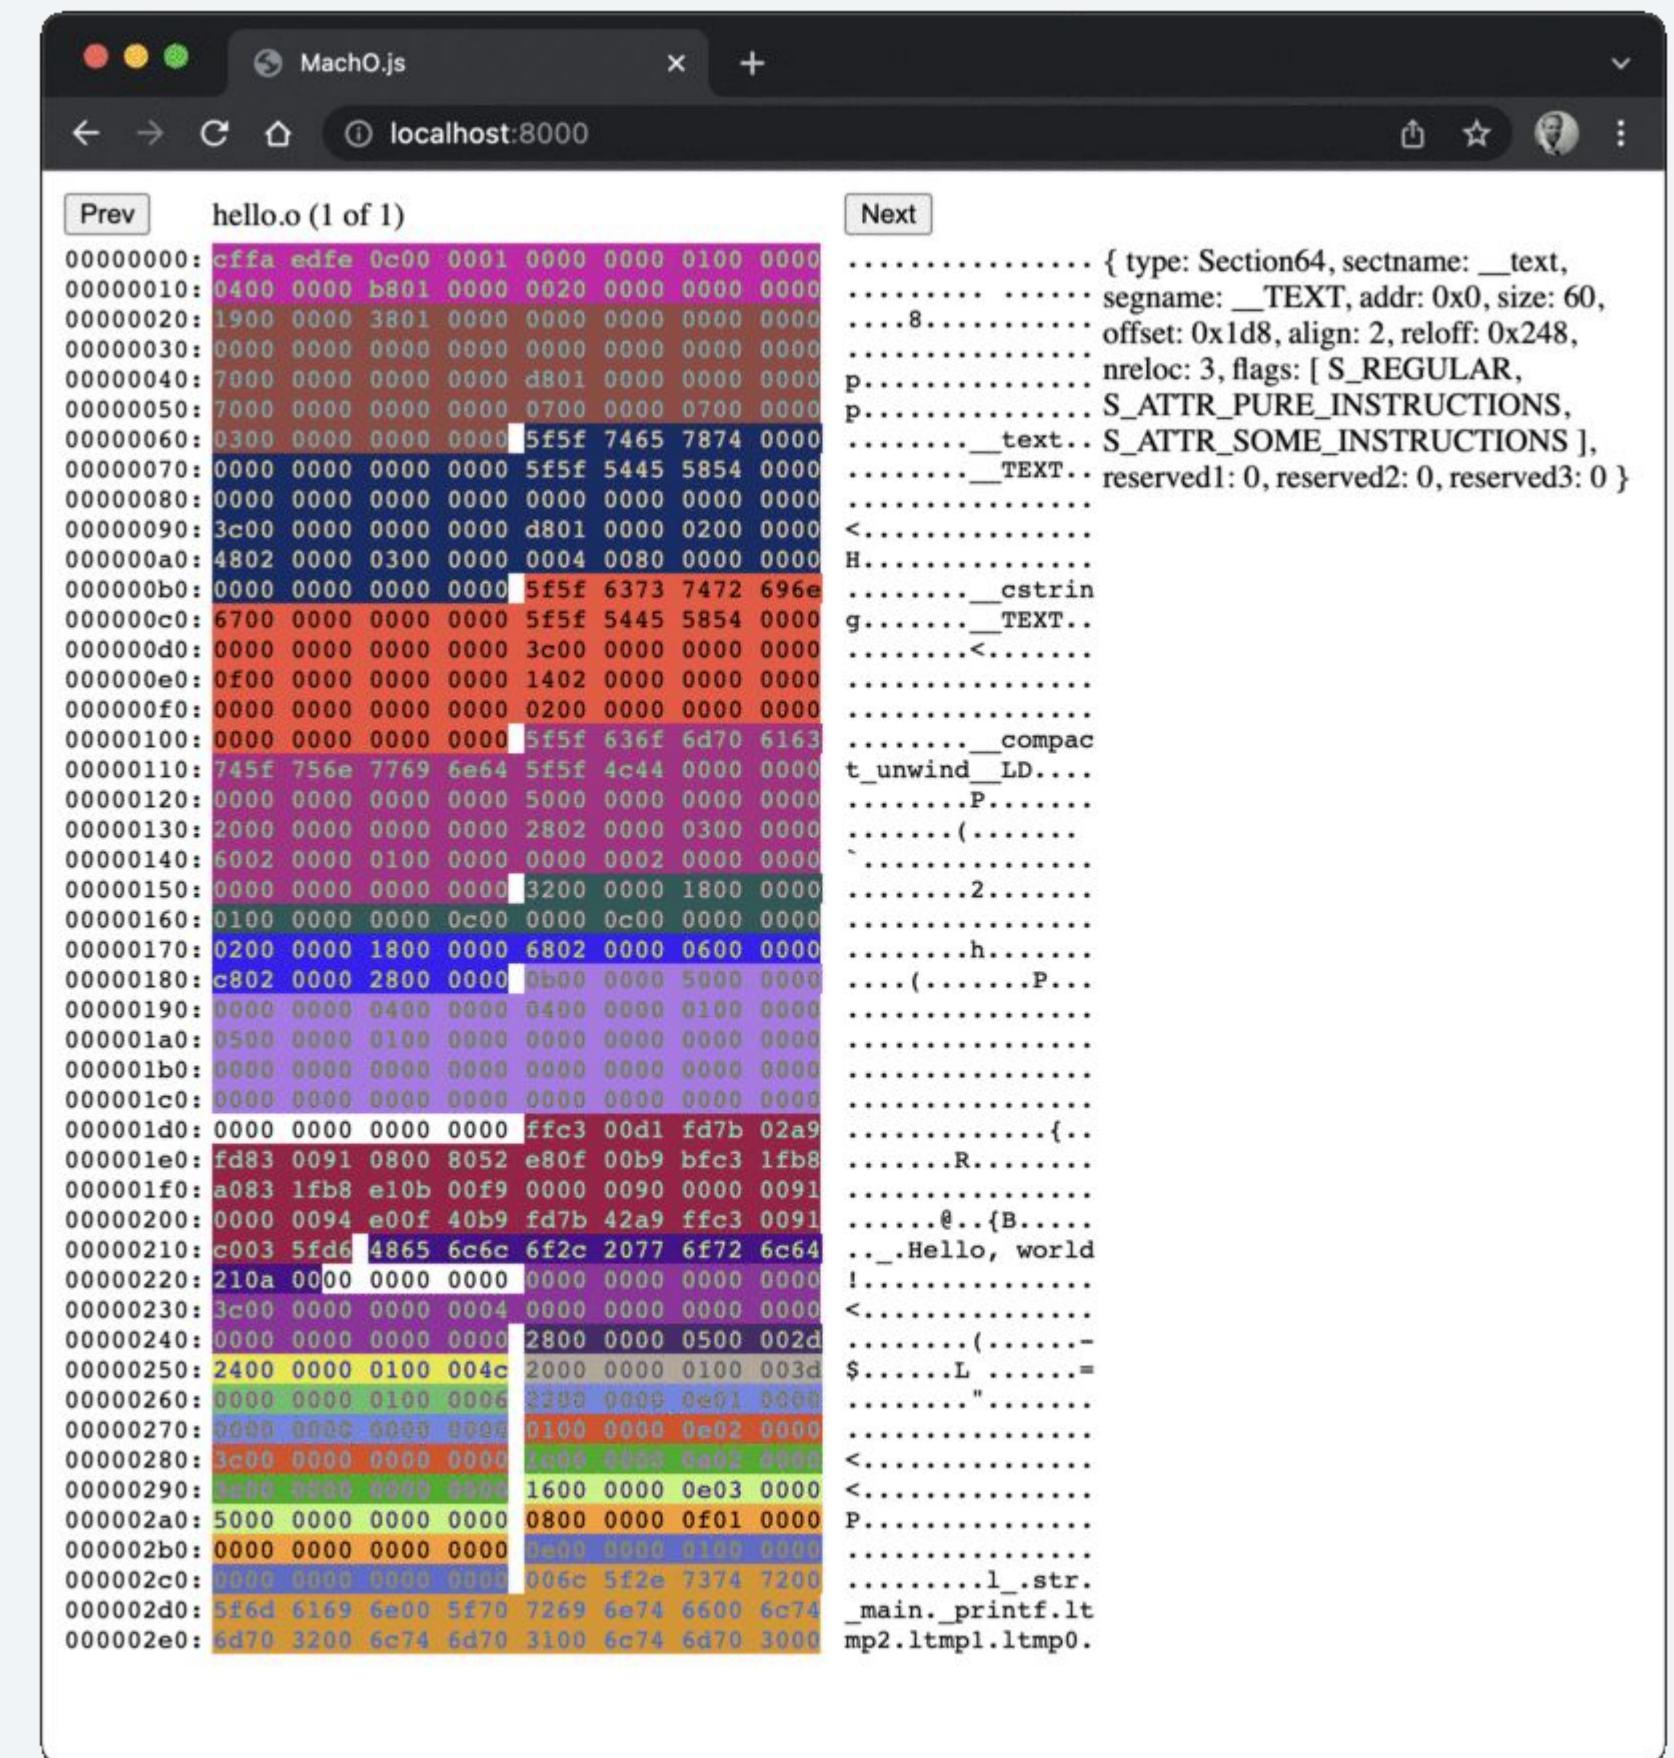
contrived test cases to start (hello.c, etc)

leverage clang’s -fuse-ld=LowD option (actually, requires ld and symlink)

pain point: Apple’s dynamic linker cache as of Big Sur ([workaround](#))

eventually add a new target to LLD and pretend its an ELF target

Weird combination of LD64 (Mach-O) and LLD (ELF)



05 The road ahead



The road ahead

native_posix on macOS

Continue work on LowD in spare time

Retrofit to LLVM, accept -fhashed-section-names=infile

Add section reordering support to machotools

Update gen_handles.py to handle Mach-O binaries via machotools

Zephyr support (cmake integration, include/toolchain)

Make it available for users via the Zephyr SDK [EXPERIMENTAL]

What to do about weak symbols: use Kconfig to configure software 🤔

Beyond native_posix

Qemu network support for macos ([#35515](#))

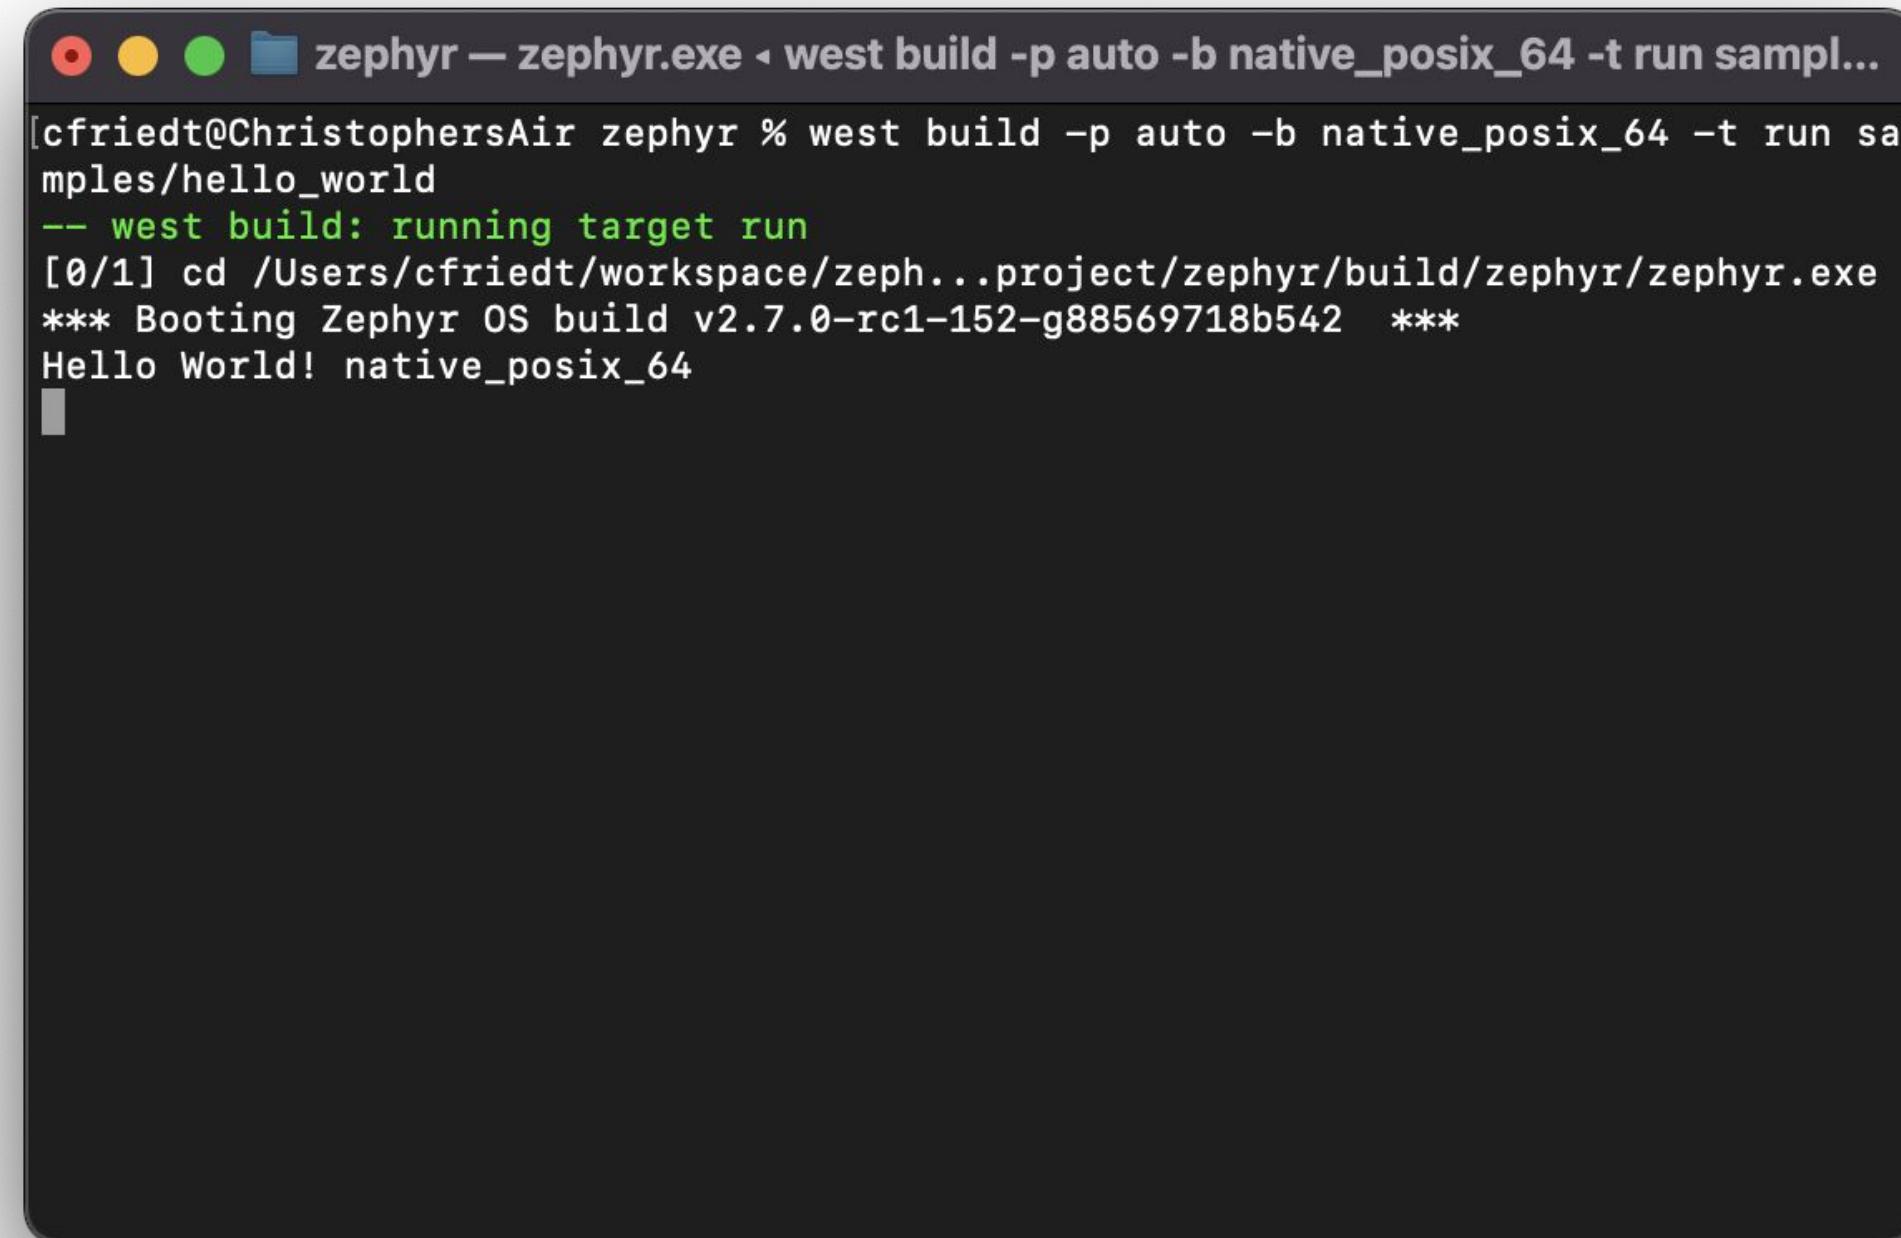
Allow CONFIG_POSIX_API with CONFIG_ARCH_POSIX ([#45100](#))

Actually, the Mach-O section name limits are a surprisingly common rant on GitHub

Other open-source projects could benefit from the work



Thanks You Questions?



```
[cfriedt@ChristophersAir zephyr % west build -p auto -b native_posix_64 -t run samples/hello_world
-- west build: running target run
[0/1] cd /Users/cfriedt/workspace/zeph...project/zephyr/build/zephyr/zephyr.exe
*** Booting Zephyr OS build v2.7.0-rc1-152-g88569718b542 ***
Hello World! native_posix_64
```

The logo consists of a blue infinity symbol followed by the word "Meta" in a dark gray sans-serif font.

∞ Meta