



Combining Zephyr & C++20 to build safe, low-footprint, event-based applications

Ioannis Papamanoglou
ioannis.papamanoglou@zonneplan.nl

Zonneplan

9th June 2022

Outline

- Introduction
- Motivation
- Demonstration
- Implementation
- Conclusions
- Appendix

Introduction

About me

- ▶ Ioannis Papamanoglou
- ▶ Rotterdam - The Netherlands
- ▶ lead embedded developer @Zonneplan
since 2021



About me

- ▶ working with Zephyr since 2019
- ▶ currently running 2.7.0 (LTS2)
- ▶ Zephyr code contributions
 - ▶ first version of DT_INST_FOREACH
 - ▶ drivers (gpio expander)
 - ▶ bugfixes (flash)

About me

- ▶ mobility
 - ▶ car sharing
 - ▶ bicycle lock for bicycle sharing
 - ▶ smart cabins
- ▶ energy
 - ▶ smart grid (solar panels, batteries, ...)
 - ▶ ev charger



About Zonneplan

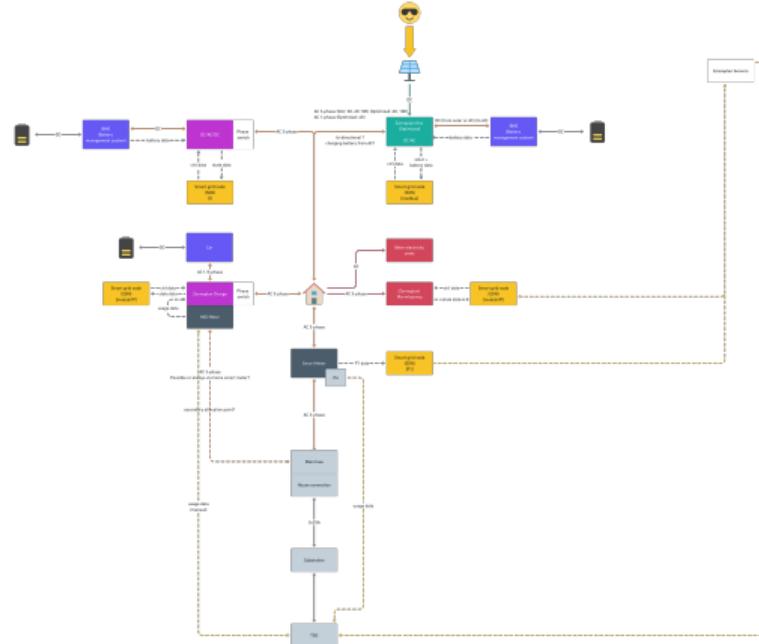
- ▶ founded 2012 as solar panel expert focused on b2c
- ▶ expanded to market leader in the Netherlands (> 100k customers)
- ▶ transitioning to smart energy business
- ▶ growing energy tech department
- ▶ all firmware running on event-based, Zephyr-powered architecture
- ▶ ~ 200 employees → 22 developers → 5 embedded

zonneplan•



About Zonneplan

- ▶ currently designing and building a smart-grid platform
 - ▶ Airnode
 - ▶ IoT device that connects energy assets
 - ▶ plug-and-play for user
 - ▶ goal: coordinate the huge amount of energy consumers/producers/storage



Abstract

- ▶ motivate event-based embedded applications
- ▶ compare to main-loop & threaded-loops
- ▶ demonstrate potential implementation

Preview

```
1 struct FooEventData {
2     int foo;
3 };
4 using FooEvent = GenEvent<BaseEvent, FooEventData>;
5
6 class FooGenerator
7     : public EventGeneratorFull<FooEvent>
8 {
9 public:
10     FooGenerator() {}
11
12     void do_something(int val) {
13         FooEvent event;
14         event.foo = val+1;
15         NOTIFY(event);
16     }
17 };
```

Listing: event generator example

Preview

```
19 class FooListener
20     : public EventListenerWithAsyncStreamQueue<FooEvent, 5>
21 {
22     public:
23         FooListener(FooEvent::GeneratorType& generator)
24             : EventListenerWithAsyncStreamQueue<FooEvent, 5>(
25                 generator, RingBufferPolicy::KEEP_LAST, WorkerQSys)
26         { }
27
28         void handle(FooEvent& event, FooEvent::GeneratorType& generator)
29         override {
30             LOG_DBG("Got foo event: %d", event.foo)
31         }
32     };
```

[Listing: event listener example](#)

Motivation

Typical embedded example

- ▶ power: mains
- ▶ interfaces: modbus rtu, pwm, adc
- ▶ network: none
- ▶ memory: 16k
- ▶ flash: 128k
- ▶ application: smart electric vehicle charger



Typical embedded example

- ▶ communicates to EV via PWM & ADC
 - ▶ → maximum charge current
 - ▶ ← EV status
- ▶ communicates to host via Modbus RTU
 - ▶ → current state, ...
 - ▶ ← charging current, charging on/off
- ▶ control power relays (GPIO)
- ▶ check safety systems
 - ▶ leakage current (GPIO)
 - ▶ relay weldcheck (GPIO)
 - ▶ temperature (I^2C sensor)

Typical embedded example

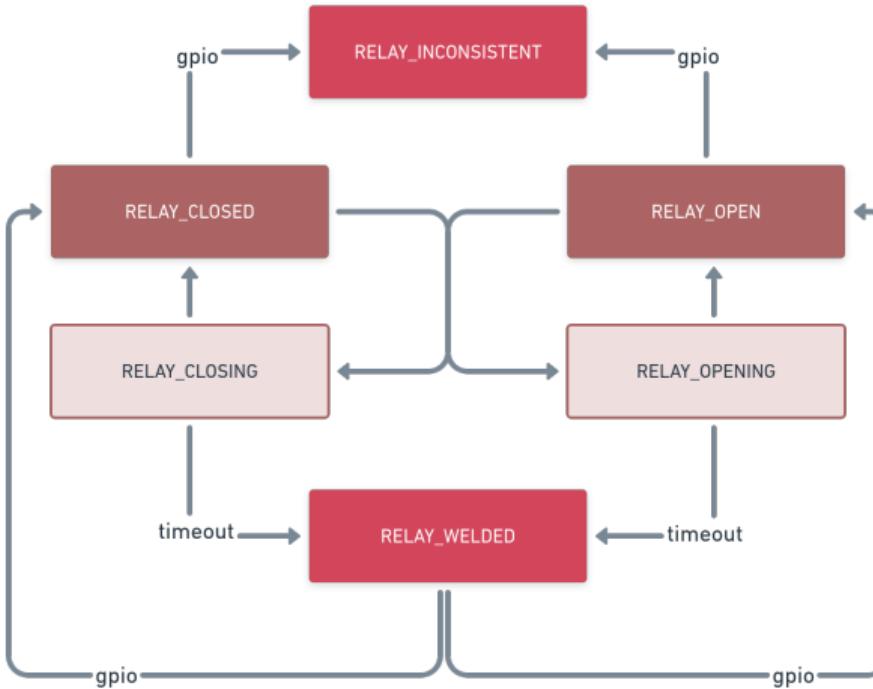


Figure: states single component (relay)

Typical embedded example

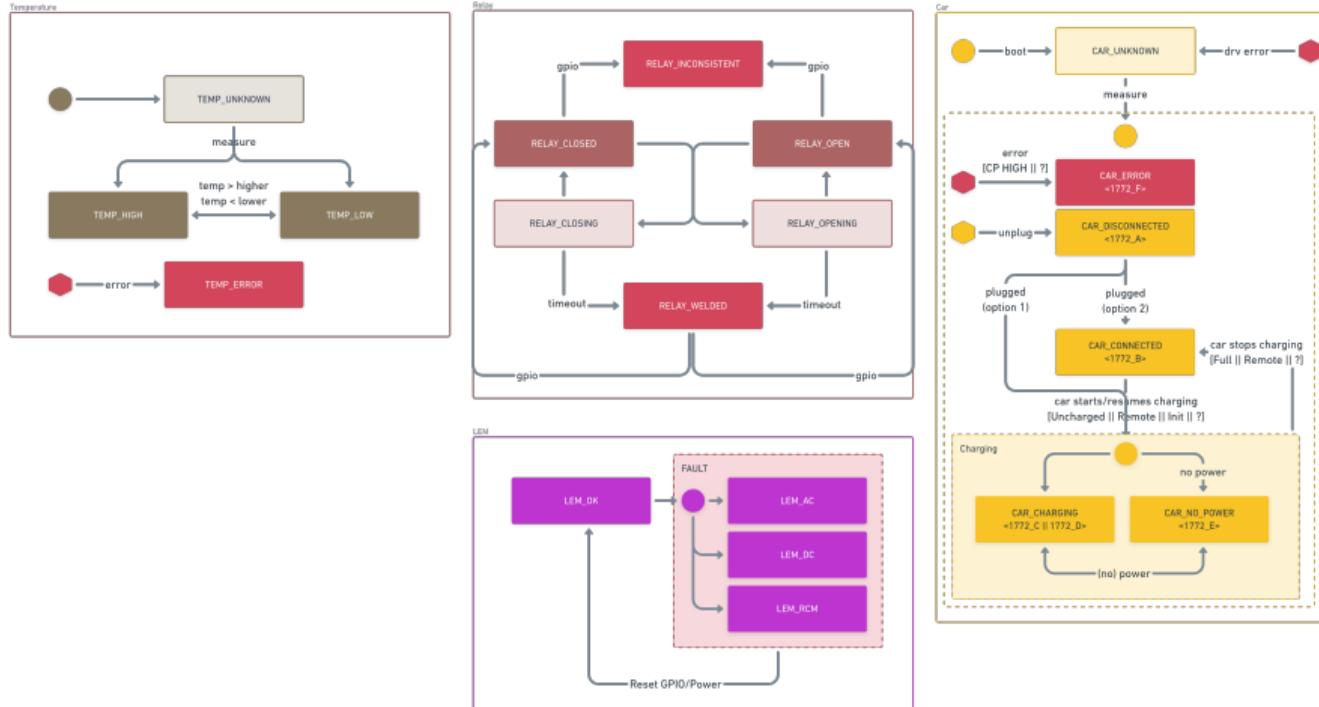


Figure: states components

Typical embedded example

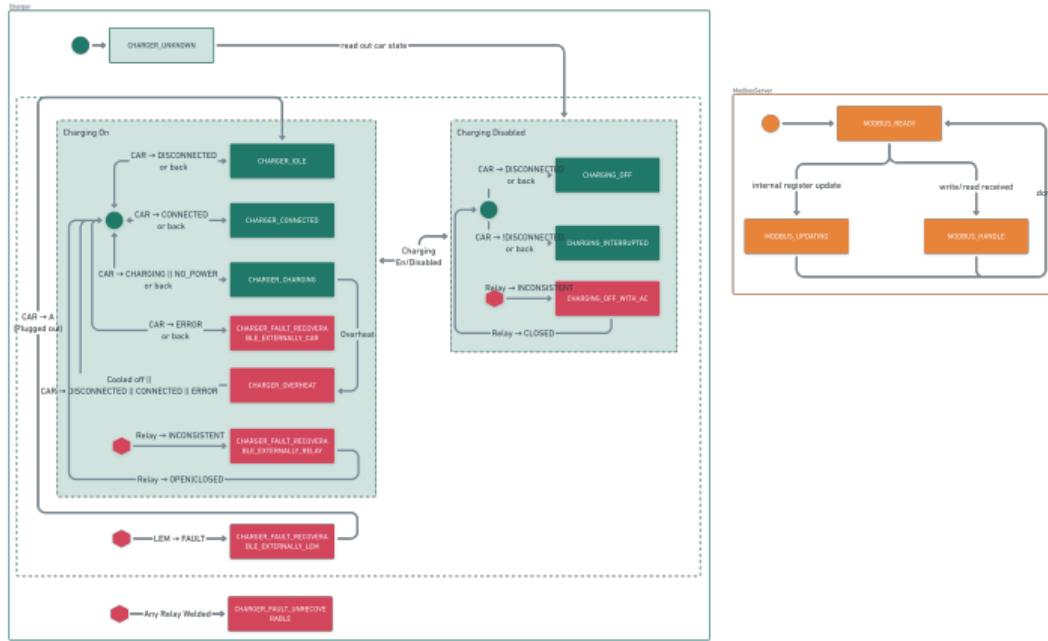


Figure: states app modules

Typical embedded example

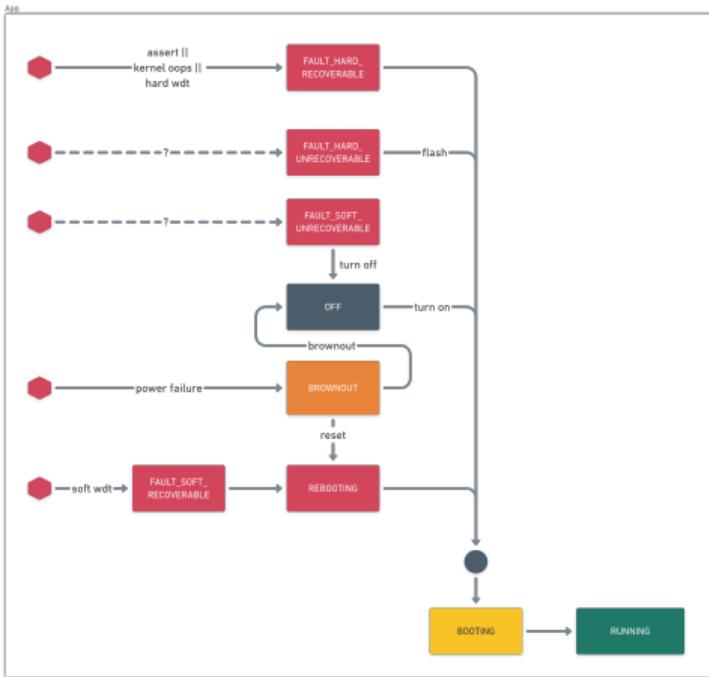


Figure: states app

Analysis

- ▶ several components with their own states
- ▶ state transitions triggered by external events or internal calls
- ▶ state changes themselves are events for higher-level modules
- ▶ events/calls can happen at the same time
- ▶ certain event/call handlers require higher priorities than others

Design pattern comparison

- ▶ main-loop
- ▶ threaded-loops
- ▶ event-based

main-loop

- ▶ single user thread
- ▶ loop
 - ▶ checks current state
 - ▶ checks current pending events
 - ▶ acts
 - ▶ switches state

main-loop

```
21 K_SEM_DEFINE(timing, 1, 1);
22 //define timer that gives timing sem every 100ms
23 //...
24
25 while(k_sem_take(&timing, K_FOREVER) == 0) {
26     // check lem events
27     LemState lem_state = lem.get_state();
28     if (lem_state != LEM_OK) {
29         this->try_switch_state(...);
30     }
31     else if (state == RECOVERABLE_LEM) {
32         this->try_switch_state(...);
33     }
34     // check relay events ...
35     // check ev events ...
36 }
```

main-loop

```
21 K_SEM_DEFINE(timing, 1, 1);
22 //define timer that gives timing sem every 100ms
23 //...
24
25 while(k_sem_take(&timing, K_FOREVER) == 0) {
26     // check lem events
27     LemState lem_state = lem.get_state();
28     if (lem_state != LEM_OK) {
29         this->try_switch_state(...);
30     }
31     else if (state == RECOVERABLE_LEM) {
32         this->try_switch_state(...);
33     }
34     // check relay events ...
35     // check ev events ...
36 }
```

- ▶ Only deterministic if execution paths take less than 100ms

- ▶ conclusion
 - ▶ lots of intertwined logic (e.g down calls)
 - ▶ lots of state to keep track of
 - ▶ priorities implicit
 - ▶ difficult to read, extend and maintain
 - ▶ data passing difficult

threaded-loops

- ▶ multiple threads (for each module/priority)
- ▶ per-module state
- ▶ loop
 - ▶ check module state
 - ▶ check pending events for module
 - ▶ acts
 - ▶ switches state
 - ▶ communicate to other threads/modules

threaded-loops

```
1 //LEM THREAD
2 K_SEM_DEFINE(timing, 1, 1);
3 //define timer that gives timing sem every 100ms
4 //...
5
6 while(k_sem_take(&timing, K_FOREVER) == 0) {
7     // check lem events
8     LemState lem_state = this->get_state();
9     if (lem_state != LEM_OK) {
10         this->try_switch_state(...);
11     }
12     else if (state == RECOVERABLE_LEM) {
13         this->try_switch_state(...);
14     }
15 }
16
17 void LEM::try_switch_state(LemState new_state) {
18     // ... do state machine stuff
19
20     this->parent.notify_lem_state_change(new_state);
21 }
```

threaded-loops

```
1 //LEM THREAD
2 K_SEM_DEFINE(timing, 1, 1);
3 //define timer that gives timing sem every 100ms
4 //...
5
6 while(k_sem_take(&timing, K_FOREVER) == 0) {
7     // check lem events
8     LemState lem_state = this->get_state();
9     if (lem_state != LEM_OK) {
10         this->try_switch_state(...);
11     }
12     else if (state == RECOVERABLE_LEM) {
13         this->try_switch_state(...);
14     }
15 }
16
17 void LEM::try_switch_state(LemState new_state) {
18     // ... do state machine stuff
19
20     this->parent.notify_lem_state_change(new_state);
21 }
```

- ▶ explicit upcall
- ▶ synchronization in parent module needed

threaded-loops

```
27 void LEM::try_switch_state(LemState state) {
28     // do state machine stuff
29
30     k_msgq_put(&this->parent_queue,
31                 &state, K_FOREVER);
32 }
33
34
35 //PARENT THREAD
36 int main() {
37     // check all queues
38     int rc;
39     rc = k_msgq_get(&this->lem_queue,
40                     &lem_state, K_NO_WAIT);
41     if (rc == 0) {
42         //handle lem event
43     }
44     rc = k_msgq_get(&this->relay_queue,
45                     &relay_state, K_NO_WAIT);
46     //...
47 }
```

threaded-loops

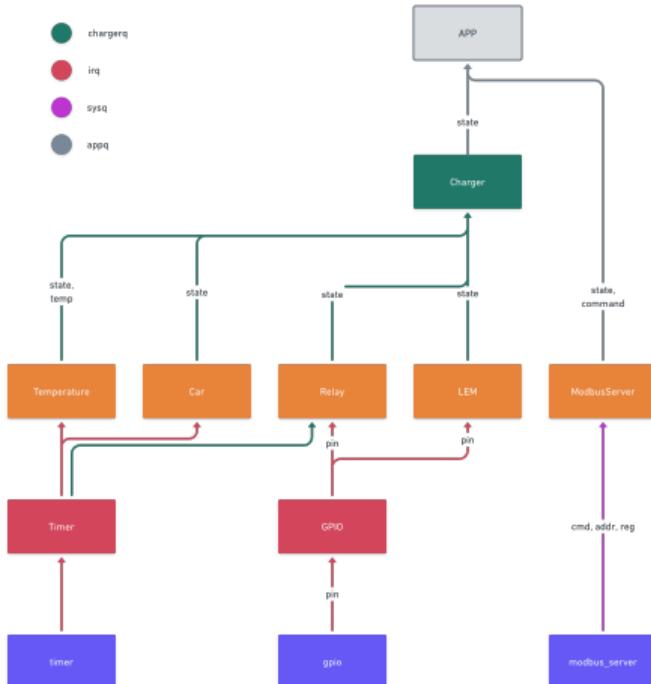
```
27 void LEM::try_switch_state(LemState state) {
28     // do state machine stuff
29
30     k_msgq_put(&this->parent_queue,
31                 &state, K_FOREVER);
32 }
33
34
35 //PARENT THREAD
36 int main() {
37     // check all queues
38     int rc;
39     rc = k_msgq_get(&this->lem_queue,
40                     &lem_state, K_NO_WAIT);
41     if (rc == 0) {
42         //handle lem event
43     }
44     rc = k_msgq_get(&this->relay_queue,
45                     &relay_state, K_NO_WAIT);
46     //...
47 }
```

- ▶ more abstract, but still upcall
- ▶ priority inversion in parent
- ▶ parent implicit priorities
- ▶ type unsafe (can be fixed with C++)

- ▶ conclusion
 - ▶ more modular
 - ▶ explicit but static priorities per module
 - ▶ for proper separation a lot of overhead
 - ▶ data passing expensive
 - ▶ software-engineering wise still not perfect

event-based

- ▶ shared workqueue threads
- ▶ per-module state
- ▶ “loop”
 - ▶ react to incoming event
 - ▶ switch state
 - ▶ generate event



event-based

```
2 void Lem::handle(GPIOEvent &event, GPIOEvent::GeneratorType &generator) {
3     switch (state) {
4         case LEM_OK:
5             this->try_state_change(LEM_FAULT);
6             break;
7         //...
8     }
9
10 void Lem::try_state_change(LemState state) {
11     // handle state change
12     //...
13
14     LemEvent event;
15     event.new_state = this->state;
16     NOTIFY(event);
17 }
18
```

event-based

```
20 class Charger : public EventListener<LemEvent> //...
21 private:
22     EventStreamQueue<LemEvent, 1> lem_event_q;
23     //EventStreamDirectSync<LemEvent> lem_event_q; //direct-call alternative
24     void handle(LemEvent &event, LemEvent::GeneratorType &lem) override;
25     Charger(WorkerQ &workerq)
26         : //...
27         , lem_event_q(*this, lem, RingBufferPolicy::KEEP_LAST, workerq)
28         //, lem_event_q(*this, lem) //direct-call alternative
29     {}
30 };
31
32
33 // runs in workerq or irq depending on line 22/23
34 void Charger::handle(LemEvent &event, LemEvent::GeneratorType &lem) {
35     if (event.new_state != LEM_OK) {
36         this->try_state_change(CHARGER_FAULT_RECOVERABLE_EXTERNALLY_LEM);
37         return;
38     }
39 }
```

- ▶ conclusion
 - ▶ data passing lowest overhead and easy
 - ▶ modular
 - ▶ explicit priorities and thread separation
 - ▶ no tuning of buffers needed
 - ▶ no up or downcalls needed, just a link/relationship
 - ▶ type-safe

Demonstration

Sync

```
1 struct Radio1EventData { int fm1; };
2 using Radio1Event = GenEvent<BaseEvent, Radio1EventData>;
3
4 class Radio1 : public EventGeneratorFull<Radio1Event> {
5     void tune1(int fm1) { // ~30 instructions
6         Radio1Event e;
7         e.get<Radio1Event>().fm1 = fm1;
8         NOTIFY(e);
9     }
10 };
11
12 class BaseListener : EventListener<BaseEvent> {
13     EventStreamDirectSync<BaseEvent> event_q;
14     int event_cnt = 0;
15     // ~10 instructions (including event_q.handle)
16     void handle(BaseEvent &event, BaseEvent::GeneratorType &generator) {
17         event_cnt++;
18     }
19     BaseListener(BaseEvent::GeneratorType &target)
20         : event_q(*this, target) {}
21 };
```

Sync

```
25 int demo_main(const struct device *_)
26     Radio1 radio;
27     BaseListener listener(radio);
28     PRINT_SIZE(radio); //16
29     PRINT_SIZE(listener); //36
30
31     LOG_DBG("Count: %d", listener.event_cnt); // 0
32     radio.tune1(50);
33     LOG_DBG("Count: %d", listener.event_cnt); // 1
34
35     return 0;
36 }
```

Async

```
1 struct Radio1EventData { int fm1; };
2 using Radio1Event = GenEvent<BaseEvent, Radio1EventData>;
3
4 class Radio1 : public EventGeneratorFull<Radio1Event> {
5     void tune1(int fm1) { // ~30 instructions
6         Radio1Event e;
7         e.get<Radio1Event>().fm1 = fm1;
8         NOTIFY(e);
9     }
10 };
11
12 class BaseListener : EventListener<BaseEvent> {
13     EventStreamQueue<BaseEvent, 1> event_q;
14     int event_cnt = 0;
15     // ~40 caller, ~30 callee thread
16     void handle(BaseEvent &event, BaseEvent::GeneratorType &generator) {
17         event_cnt++;
18     }
19     BaseListener(BaseEvent::GeneratorType &target)
20         : event_q(*this, target, RingBufferPolicy::KEEP_LAST, WorkerQSys) {}
21 };
```

Async

```
25 int demo_main(const struct device *_)
26     Radio1 radio;
27     BaseListener listener(radio);
28     PRINT_SIZE(radio); //16
29     PRINT_SIZE(listener); //112
30
31     LOG_DBG("Count: %d", listener.event_cnt); // 0
32     radio.tune1(50);
33     LOG_DBG("Count: %d", listener.event_cnt); // 0
34     k_msleep(1000);
35     LOG_DBG("Count: %d", listener.event_cnt); // 1
36
37     return 0;
38 }
```

Implementation

Requirements

- ▶ no heap allocation
- ▶ infinite listeners per generator
- ▶ infinite generators per listener
- ▶ asynchronous event handling - generation
 - ▶ dynamic queue sizes
 - ▶ different dropping strategies
- ▶ every listener receives every event
- ▶ listeners can access generator
- ▶ listeners know if and how many events are dropped
- ▶ as much as possible done/checked during compile time
- ▶ events are not necessarily specific to generators
- ▶ hierarchical events
- ▶ events can pass data

Overview

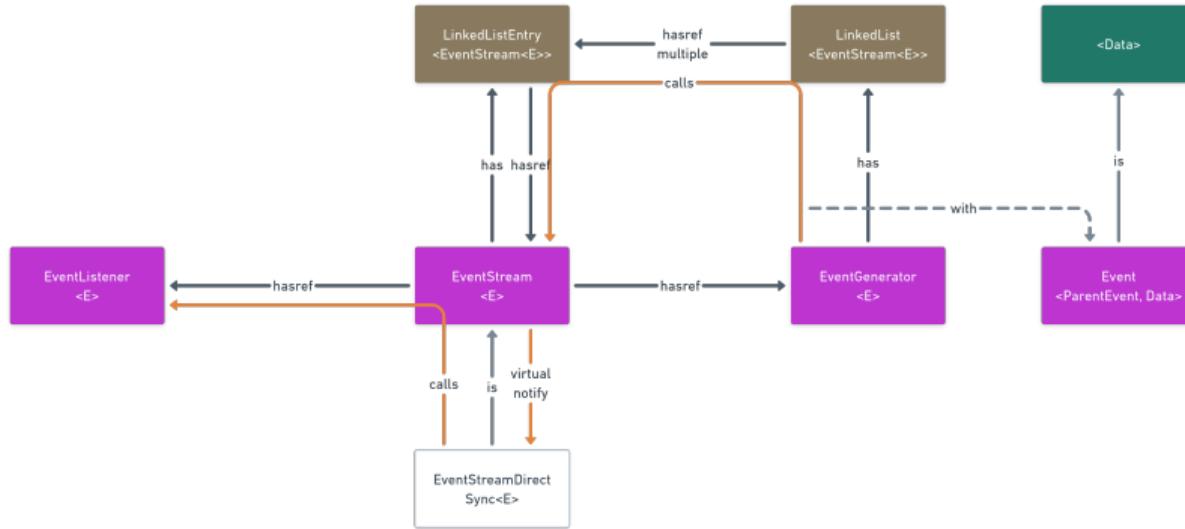


Figure: event-system with direct-call

Overview

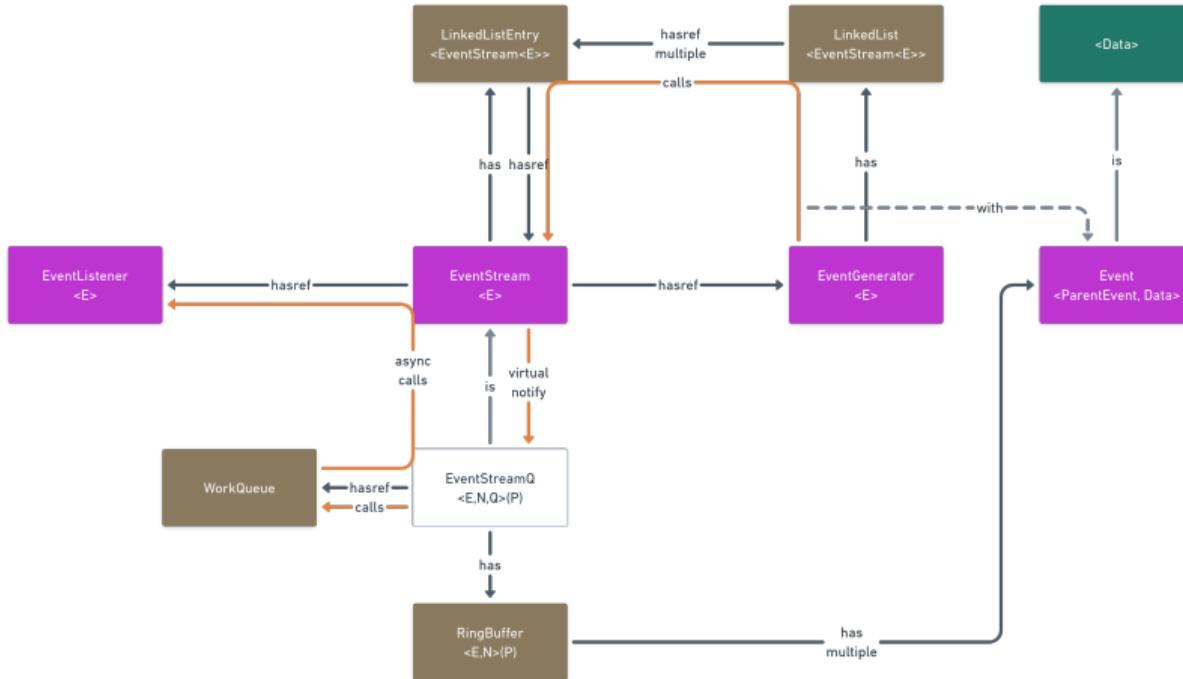


Figure: event-system with async-call

Instantiation

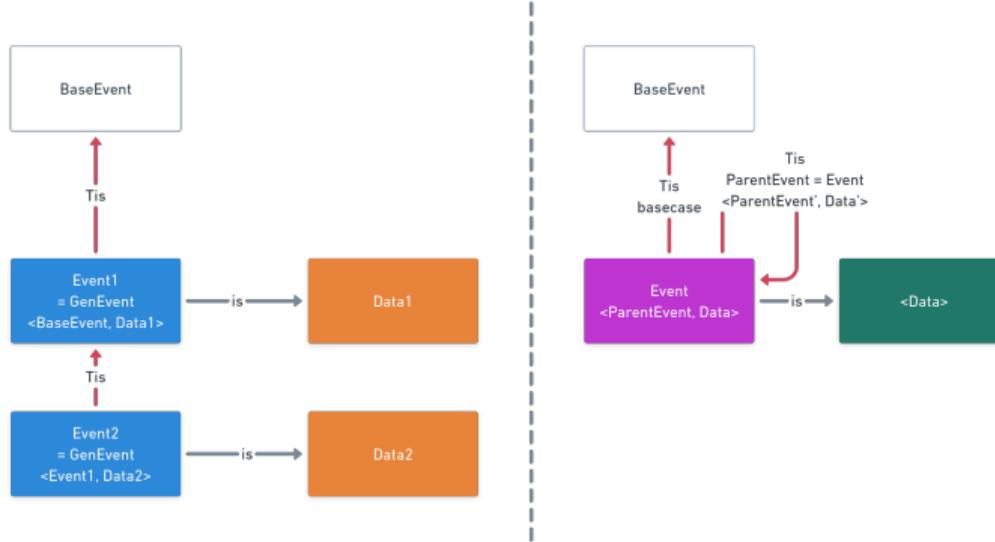


Figure: event instantiation

Instantiation

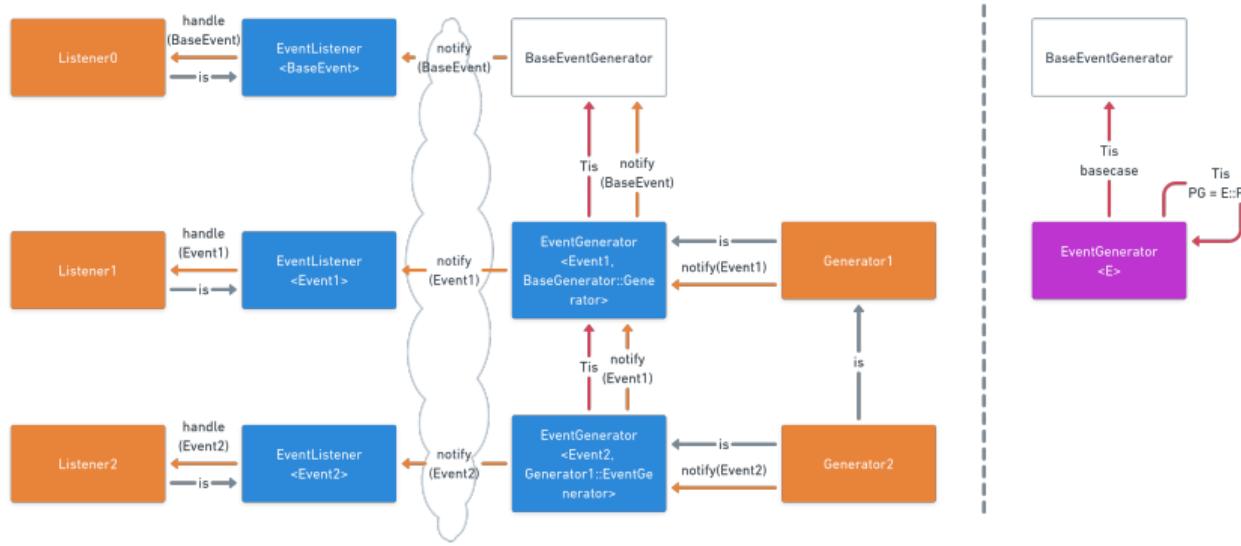


Figure: generator/listener instantiation

Conclusions

Conclusions

- ▶ event-based architectures can be useful to create safe, low-footprint embedded applications

Conclusions

- ▶ event-based architectures can be useful to create safe, low-footprint embedded applications
- ▶ events cooperate well with RTOS and statemachines

Conclusions

- ▶ event-based architectures can be useful to create safe, low-footprint embedded applications
- ▶ events cooperate well with RTOS and statemachines
- ▶ C++20 and Zephyr make it reasonably easy to implement event-based architectures

THANK YOU!

Questions?



<https://sched.co/10CEa>

Appendix

Event inheritance

```
1 struct Radio1EventData { int fm1; };
2 struct Radio2EventData { int fm2; };
3 struct Pillow1EventData { int sft1; };
4
5 using Radio1Event = GenEvent<BaseEvent, Radio1EventData>;
6 using Radio2Event = GenEvent<Radio1Event, Radio2EventData>;
7 using Pillow1Event = GenEvent<BaseEvent, Pillow1EventData>;
8
9 class Radio1 : public EventGeneratorFull<Radio1Event> //...
10 class Radio2 : public EventGeneratorExtendFull<Radio1, Radio2Event> //...
11 class Pillow1 : public EventGeneratorFull<Pillow1Event> //....
12
13 class Pillow1Radio2withBar
14     : public EventDiamond<
15         Pillow1,
16         Radio2,
17         EventGeneratorFull<BarEvent>
18     > //...
```

Implementation - Events - template inheritance

```
1  template<isEvent T_BASE, typename T_EXT>
2  struct GenEvent : public T_EXT {
3      using ParentType = T_BASE;
4      using GeneratorType = Generator<GenEvent<T_BASE, T_EXT>>;
5      T_BASE m_base_event;
6
7      template<isEvent T_TARGET>
8      T_TARGET& get() {
9          if constexpr (std::is_same_v<GenEvent<T_BASE, T_EXT>, T_TARGET>) {
10              return *this;
11          }
12          else if constexpr (std::is_same_v<T_BASE, T_TARGET>) {
13              return this->m_base_event;
14          }
15          else if constexpr (std::is_same_v<BaseEvent, T_BASE>) {
16              static_assert(!std::is_same_v<BaseEvent, T_BASE>, "not a base of this event");
17              return *((T_TARGET*)nullptr);
18          }
19          else {
20              return this->m_base_event.template get<T_TARGET>();
21          }
}
```

Implementation - Generators

```
25 template <isEvent E>
26 class Generator {
27 private:
28     using EventStreamType = EventStream<E>;
29     LinkedList<EventStreamType> event_streams;
30 public:
31     using ParentGenerator = E::ParentType::GeneratorType;
32     using ParentType = ParentGenerator;
33     using GeneratorType = Generator<E>;
34
35     bool subscribe(LinkedListEntry<EventStreamType>& event_stream_node) {
36         this->event_streams.add_prev(event_stream_node);
37         return true;
38     }
39
40     template<typename GeneratorObjectPtr>
41     void notify(E& e) {
42         //...
43     }
44 };
```

Implementation - Generators - recursive notify

```
46 template<typename GeneratorObjectPtr>
47 void notify(E& e) {
48     // check that we are passed something that might actually be a child of this
49     static_assert(std::is_base_of_v<
50         GeneratorType,
51         typename std::remove_pointer<GeneratorObjectPtr>::type
52     >);
53
54     // notify all listeners indirectly by feeding the event streams
55     for (EventStreamType& event_stream : this->event_streams) {
56         event_stream.handle(e);
57     }
58
59     // pass further to top
60     if constexpr (!std::is_same_v<ParentGenerator, void>) {
61         static_cast<GeneratorObjectPtr>(this)->ParentGenerator
62             ::template notify<GeneratorObjectPtr>(e.get_base_event());
63     }
64 }
```

Implementation - Streams - async

```
66 template<isEvent E, size_t N>
67 class EventStreamQueue : public EventStream<E> {
68     RingBuffer<E,N> events;
69     WorkQ<&EventStreamQueue<E,N>::handle_async> async;
70     EventStreamQueue(EventListener<E> &listener, EventStream<E>::G &generator,
71                     RingBufferPolicy replacement_policy, WorkerQ &q)
72         : EventStream<E>(listener, generator)
73         , events(replacement_policy)
74         , async(q, *this) {}
75
76     void handle_async(void *obj) {
77         std::optional<E> eventopt = events.pop();
78         this->listener.handle(*eventopt, this->generator);
79     }
80
81     void handle(E const& event) override {
82         if(!events.try_push(event)) {
83             return;
84         }
85         this->async.schedule(K_NO_WAIT, false);
86     }
}
```