



Clock Management in Zephyr RTOS

Daniel DeGrasse

NXP Semiconductors

A night-time photograph of the Space Needle in Seattle, Washington. The tower is illuminated with white lights and has a red light at its peak. The surrounding city lights are visible, and the lights reflect on the water in the foreground. The sky is a deep blue with some clouds.

Agenda

- Overview of current Clock Control framework
- Discussion of framework requirements
- Proposed Implementation
 - SOC perspective
 - Driver API
- Framework Implementation Details

Overview of Clock Control Framework



Clock Control Framework

- Present in Zephyr since 1.0
- Based around hardware specific “subsystems”
- Supported operations
 - Get clock frequency as integer
 - Set rate via opaque driver specific parameter
 - Get clock status
 - Gate/enable a clock (optionally can do so asynchronously)
 - Configure a clock via opaque pointer
- Besides reading rates, all clock data is driver specific
- Drivers can read clock settings from devicetree to setup clocks at boot time
 - Configure API can also be used for runtime changes

Clock Control Usage Example- STM32



- SOC series specific clock drivers
- STM32 specific macro "STM32_DT_CLOCKS" used to setup configuration structure for clock data
- Clock control configure API permits devices to enable/configure their clocks
 - Device specific clock data encodes register offset and bits to set
- Multiplexer selections are made via devicetree node with "st,stm32-clock-mux" compatible, which applies selections via clock configure API at boot time
- Reads clock controller node devicetree properties to initialize root clocks
 - PLLs set up via this method
 - Validates that each resulting frequency will be within range supported by SOC at build time

Why Rework Clock Control?

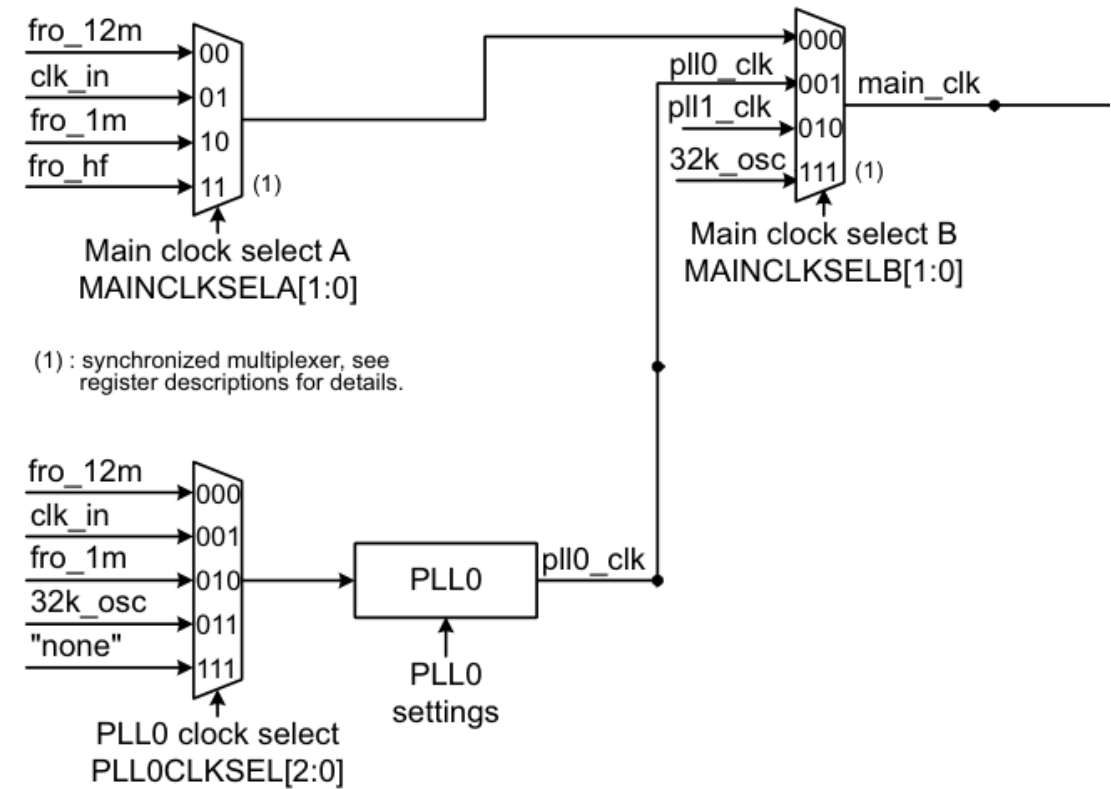
- Better code reuse
 - Dividers and multiplexers are often reused across SOC lines, so we should reuse drivers
- Runtime clock configuration
 - Limited ability to configure complicated sources like PLLs at runtime
 - Usually use settings on clock controller node
 - No facility for clock callbacks exists in clock control framework
- Clock control framework leaks device specific data
 - Need vendor specific macros in driver to initialize clock data
 - Details about clock driver leak into device driver implementations
- Set rate API
 - Flexible, but high runtime/flash overhead to implement this feature

Clock Control Requirements



Clock Hardware

- Gates
 - Often one per peripheral- enable or disable clock signal for peripheral device
 - Can also gate root clock sources
- Sources
 - Usually fixed output frequency, sometimes can be powered off at runtime
- Dividers/Multipliers
- Multiplexers
- PLLs
 - Often very vendor/SOC specific (analog IP)
 - Include multipliers and dividers, as well as vendor specific settings
 - Example: spread spectrum generator on LPC55sxx series
 - Significant power usage- need to be powered down in idle modes



Framework Goals

- Feature parity with clock control
- Ability to configure clock settings at runtime for power savings
 - Support powering off clock sources
- Notify clock consumers when parent changes rate
- Hide clock implementation details from device drivers
- Describe clock hierarchy in devicetree
- Maximize code reuse across SOCs
 - Digital IP like multiplexers and dividers is often reused across SOC lines
- Minimize flash/ram footprint
 - Ideally similar footprint to existing driver with clock control framework
 - Runtime rate calculation should be optional- many devices won't need this feature

Framework Proposal



- Clock driver layer
 - Based on Linux common clock framework
 - Each clock is represented by a clock object
 - Clock objects hold references to parent (for rate calculation) and children (for callback notification)
- Clock management layer
 - Consumers only interact with this layer
 - Clock configurations are described as “states” in devicetree
 - Series of clock nodes with hardware specific configuration data as specifiers
 - Clock outputs can have their frequency queried directly
 - Setting rate can be accomplished via “output clock” nodes
 - Support for rate change callbacks

Clock Driver API

- Heavily inspired by Linux common clock framework
- Clock elements like multiplexers, dividers, and PLLs are considered “nodes”
 - Each clock node will have a clock structure describing it
 - Structures will have private hardware data, as well as clock API
- Clock nodes will implement clock API
 - notify- inform clock that parent is going to change rate
 - get_rate: reads clock rate in Hz
 - configure: configures clock with hardware specific data
 - set_rate: Sets clock rate in Hz
 - round_rate: Determine closest supportable clock rate
- Like device model, clock nodes can reference one another
 - Allows clock node to read parent rate and calculate its own rate
- External clock generators can also be supported
 - No distinction between on die and off die clock sources
- Set rate and round rate APIs are enabled via Kconfig-
devices that don't need this support can disable it

```
/* Notify clock that parent will be reconfigured */
int clock_notify(const struct clk *clk,
                const struct clk *parent,
                uint32_t parent_rate);
/* Read clock rate */
int clock_get_rate(const struct clk *clk);
/* Configure clock */
int clock_configure(const struct clk *clk, void *data);
/* Get closest rate clock can support */
int clock_round_rate(const struct clk *clk,
                    uint32_t req_rate);
/* Set clock to given rate */
int clock_set_rate(const struct clk *clk, uint32_t rate)
```

Clock Management API

- Abstract clock settings to “states”– similar idea to pin control
 - Run and idle states defined by default
 - Drivers can define additional custom states
- Clock rate can be queried from clock outputs
- Devices can register for clock reconfiguration notifications
 - If a parent of a driver’s clock is reconfigured, the driver will be notified
- Driver clock API:
 - `clock_mgmt_get_rate`: Read the rate of a clock
 - `clock_mgmt_apply_state`: Apply a clock state
 - `clock_mgmt_set_callback`: Set a callback for a clock output
- No support for runtime rate setting
 - Drivers can define a state for a given frequency; devicetree implementer decides how to configure clocks

Clock Management SOC Devicetree

- SOC clock tree will be defined hierarchically
- Nodes reference a clock node in SOC devicetree to query rate
- Node references determine clock children and parent relationships in clock framework

```
/* SOC devicetree */
fcclksel0: fcclksel0@400002b0 {
    compatible = "nxp,syscon-clock-mux";
    #clock-cells = <1>;
    /* SYSCON::FCCLKSEL0[SEL] */
    reg = <0x400002b0 0x3>;
    offset = <0x0>;
    input-sources = <&mainclkselb &pll0div
                    &fro_12m &frohfdiv &fro_1m
                    &mclk_in &rtcosc32ksel
                    &no_clock>;
    #address-cells = <1>;
    #size-cells = <1>;

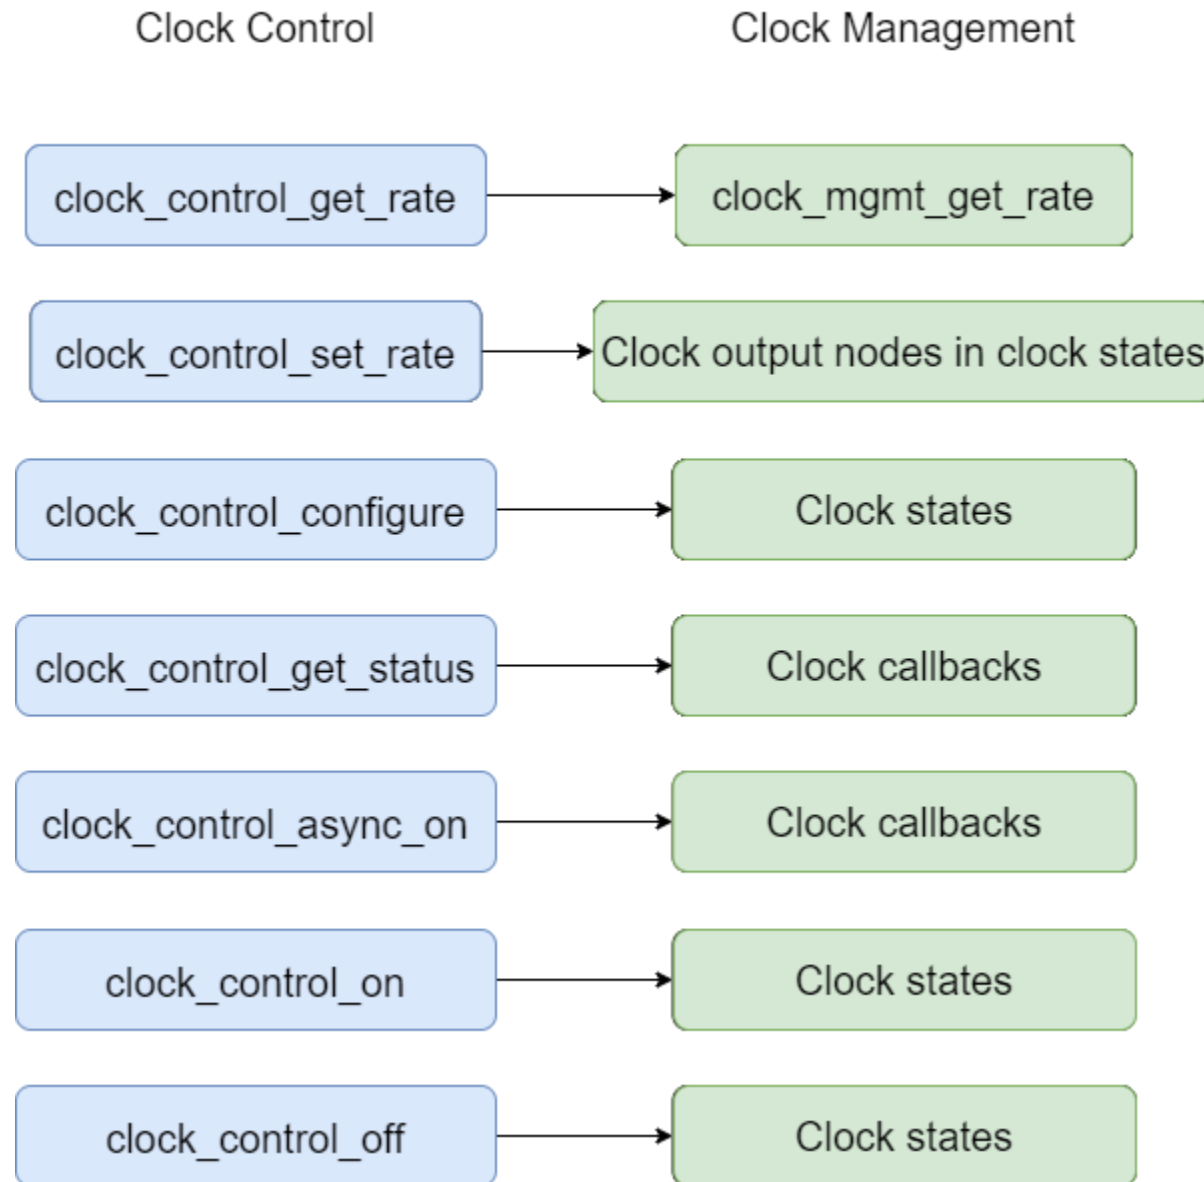
    frgctrl0_mul: frgctrl0-mul@40000320 {
        compatible = "nxp,syscon-flexfrg";
        #clock-cells = <1>;
        /* SYSCON::FLEXFRG0CTRL[MULT] */
        reg = <0x40000320 0x8>;
    };
};
```


Clock Management States

- States will be defined within devicetree, using specifiers on clock nodes
 - Specifiers will be specific to each clock node
 - Passed to clock configure API at runtime to setup clock
- Supporting set_rate with this API
 - Special “clock output” nodes will accept one specifier- the requested clock rate. This will allow drivers to configure clock rates directly if desired
- Nodes can define as many states as needed, and can configure any clock within a state

```
/* Board configuration */
&flexcomm0 {
    status = "okay";
    clock-outputs = <&frgctrl0_mul>;
    /* Setup multiplier and divider selection */
    clock-state-0 = <&fcclksel0 0 &frgctrl0_mul 0>;
    /* State 1 uses lower power clock source */
    clock-state-1 = <&fcclksel0 4 &frgctrl0_mul 0>;
    clock-state-names = "default", "sleep";
};
```

Clock Control versus Clock Management



Framework Implementation



Clock Driver Implementation

```
int syscon_clock_div_get_rate(const struct clk *clk) {
    const struct syscon_clock_div_config *config = clk->hw_config;
    int parent_rate = clock_get_rate(config->parent);
    uint8_t div_mask = GENMASK(0, (config->mask_width - 1));

    if (parent_rate <= 0) {
        return parent_rate;
    }

    /* Calculate divided clock */
    return parent_rate / ((*config->reg & div_mask) + 1);
}
```

Clock Driver Implementation

```
int syscon_clock_div_configure(const struct clk *clk, const void *div) {
    const struct syscon_clock_div_config *config = clk->hw_data;
    uint8_t div_mask = GENMASK(0, (config->mask_width - 1));
    uint32_t div_val = (((uint32_t)div) - 1) & div_mask;
    int parent_rate = clock_get_rate(config->parent);
    uint32_t new_rate = (parent_rate / ((uint32_t)div));

    clock_notify_children(clk, new_rate);
    (*config->reg) = ((*config->reg) & ~div_mask) | div_val;
    return 0;
}
```

Clock Driver Implementation

```
int syscon_clock_div_notify(const struct clk *clk,  
                           const struct clk *parent,  
                           uint32_t parent_rate) {  
    const struct syscon_clock_div_config *config = clk->hw_data;  
    uint8_t div_mask = GENMASK(0, (config->mask_width - 1));  
    uint32_t new_rate = (parent_rate / ((*config->reg & div_mask) + 1));  
  
    return clock_notify_children(clk, new_rate);  
}
```


Clock Management Subsystem Implementation

- Need a method to couple clock management to clock driver API
 - Getting output rate is simple, we already have references to the clock node and an API to do so
- Configuring setpoints will be done via clock_configure API
 - Clock drivers will supply macros to define clock specific data based on node specifiers
- Output clocks and state data will be stored in struct defined by a macro in driver
 - Like pin control, this struct will be passed to clock management functions
- Clock callbacks are implemented by defining a minimal clock structure for the driver consuming a clock
 - This will be considered a child of the clock the driver references
 - Will have the "notify" API called on it when clock is reconfigured
 - Clock structure's private data field will point to clock configuration structure

```
fcclksel0: fcclksel0@4000002b0 {
    compatible = "nxp,syscon-clock-mux";
    #clock-cells = <1>;
};

&flexcomm0 {
    status = "okay";
    clock-state-0 = <&fcclksel0 3>
    clock-state-names = "default";
};

/* Defines any data structure needed for mux */
#define Z_CLOCK_MGMT_NXP_SYSCON_CLOCK_MUX_DATA_DEFINE(node_id, prop, idx)
/* Get value that will be passed to clock configure API */
#define Z_CLOCK_MGMT_NXP_SYSCON_CLOCK_MUX_DATA_GET(node_id, prop, idx) \
    DT_PHA_BY_IDX(node_id, prop, idx, multiplexer)
```

Current Status

- Basic Implementation completed for LPC55S69
 - No support for `set_rate` or `round_rate`
 - Able to configure clock tree to run basic hello world application
- Flash usage
 - Currently about 1 KB more flash than clock control framework
 - Could make notification support optional
 - Save roughly 700KB by disabling clock notifications
- RFC
 - In progress, need documentation and completed implementation

Alternative Options

- Setpoints could be implemented in SOC specific manner
 - Less code reuse, but potentially more optimization
 - Similar to how pin control handles states
 - SOC's would be responsible for implementing callback architecture
 - RFC for this approach: [\[RFC\] Introduce Clock Management Subsystem \(PR #70467\)](#)
- We could implement this using the current clock control subsystem, and build clock management API on top of it
 - Reusing device struct to describe clocks has negative utilization impacts when 20+ structures are defined for basic clock trees
 - Set_rate API is not optional, limiting optimization possibilities
 - Extra API fields in each clock add to flash usage
 - No facility for rate change callbacks



[nxp.com](https://www.nxp.com)

| Public | NXP and the NXP logo are trademarks of NXP B.V. All other product or service names are the property of their respective owners. © 2024 NXP B.V.

Backup Slides



- Clocks must have bidirectional references
 - Parent reference needed to query rate
 - child reference needed to issue callbacks
- Can't reference children directly- they may not exist at link time
 - We also want to avoid referencing child structs unless they are needed to save flash
 - Solution- use DT ordinals and two stage link (like device dependency handles)
 - Python scripting replaces ordinal numbers with clock handles in second link stage
 - Result is that each clock has a pointer to its parent(s), and a reference to an array of children clock handles
- Generic callback helper supplied to forward callback to all children of a clock

Enabling Clocks

- Nodes are all set to “status=okay”
- Clock nodes will all be compiled in, and linker will discard unreferenced clocks
 - Clock structures must be referenced by their children
 - The clock structure created for a clock consumer by the clock management framework will hold references to all parents
- This avoids need to explicitly enable/disable clock nodes to optimize image
- Clocks can also explicitly be set to “status=disabled” if desired
 - Allows external clock generators to be enabled or disabled by overlays
- Reuse “DT_INST_FOREACH_STATUS_OKAY” macro to define clock structures

Setting Rates with Clock Management

- All clock trees will define special clock output nodes
 - SOC agnostic method to set clock rates
- Clock management subsystem calls clock_configure with clock node specific data
 - For clock outputs, this will be a frequency
- Clock outputs will simply call clock_set_rate on parent clock, to configure a given frequency
- Same concept can be used by vendor drivers to provide a node for requesting a specific output frequency with other parameters (such as resolution or accuracy needed)
- SOC clock driver implementation is responsible for supporting set_rate API
 - SOC will be best aware of quirks and features of its clock tree
 - Allows SOC to optimize implementation however needed

```
fxcom0_output: fcom0-output {
    compatible = "clock-output";
    #clock-cells = <1>;
};

&flexcomm0 {
    status = "okay";
    clock-state-0 = <&fxcom0_output 12000000>;
    clock-state-names = "default";
};

/* Defines any data structure needed for mux */
#define Z_CLOCK_MGMT_CLOCK_OUTPUT_DATA_DEFINE(node_id, prop, idx)
/* Get value that will be passed to clock configure API */
#define Z_CLOCK_MGMT_CLOCK_OUTPUT_DATA_GET(node_id, prop, idx) \
    DT_PHA_BY_IDX(node_id, prop, idx, frequency)
```

Flash Usage Considerations



- Avoiding common data fields lets linker optimize better
 - Only having a clock hardware data field works best, as many clocks don't have runtime data
- Children registering for parent callback via linker sections required usage of "keep"
 - Large flash impact
 - Key takeaway: we can't have parents reference their clock children, since we need linker to discard clock children that aren't used to save flash space
- Flash usage is main reason I did not use device infrastructure
 - A simple hello world build for LPC55S69 includes *twenty-five* clock nodes
 - The "name" field alone would increase flash usage by 100 bytes

- Used preprocessor macros to generate optimized “clock setpoint” functions for each clock state
 - Heavy preprocessor optimization, errors could be hard to debug
 - Lots of macrobatics
 - No usage of device model- single instance
- Some good discussion there- a few key issues raised:
 - How to handle external clock producers (off chip)?
 - No good option in that framework
 - How to support set rate API?
 - Could theoretically provide nodes to do so at the SOC implementation, but it would be additional code overhead