


# Future of Zephyr Device Model

Ederson de Souza, Intel Corporation



# Agenda

- Current interest
- Deferred init
- Initialization order regarding services
- Device notifications
- Device unloading
- Questions



# Device model issues and pull requests abound

Support peripheral deallocation at runtime #20012

**kernel: device: add dts base device initialization priority #58296**

**Device initialization order that respects devicetree dependencies #22545**

**Refining Zephyr's Device Driver Model - Take II #22941**

init sequence rework #6

**init: specify dependencies between init entries #49310**

**Allow for selective deferment of**

**device initialization (aka manual init) #39896**

t circular DTS node refere


Ability to use an API with device without extending its own API #22415

Adding support for multi-functional-devices #48934

Device initialization improvements #34518


**Device (de)initialization use case #40**

**Refining Zephyr's Device Driver Model #6393**



# What people want?

- Deferring device initialization
- Unloading/de-initializing devices
  - Power management integration/conflicts
- Initialization order regarding "services"
- API extensibility
- And more...



# Please talk!




# Deferred initialization

A look into the recent time

- Status: merged
- Really lengthy discussion - #39896 is from Oct/21 - maybe other discussion before
- A few PRs to tackle it
  - o With more discussion!
  - o PR discussion deviates from initial issue

```
# device tree
(...)
    status = "okay";
    zephyr,deferred-init;
};


/* Code */
ret = device_init(SOME_DEVICE);
```



# Initialization order regarding services


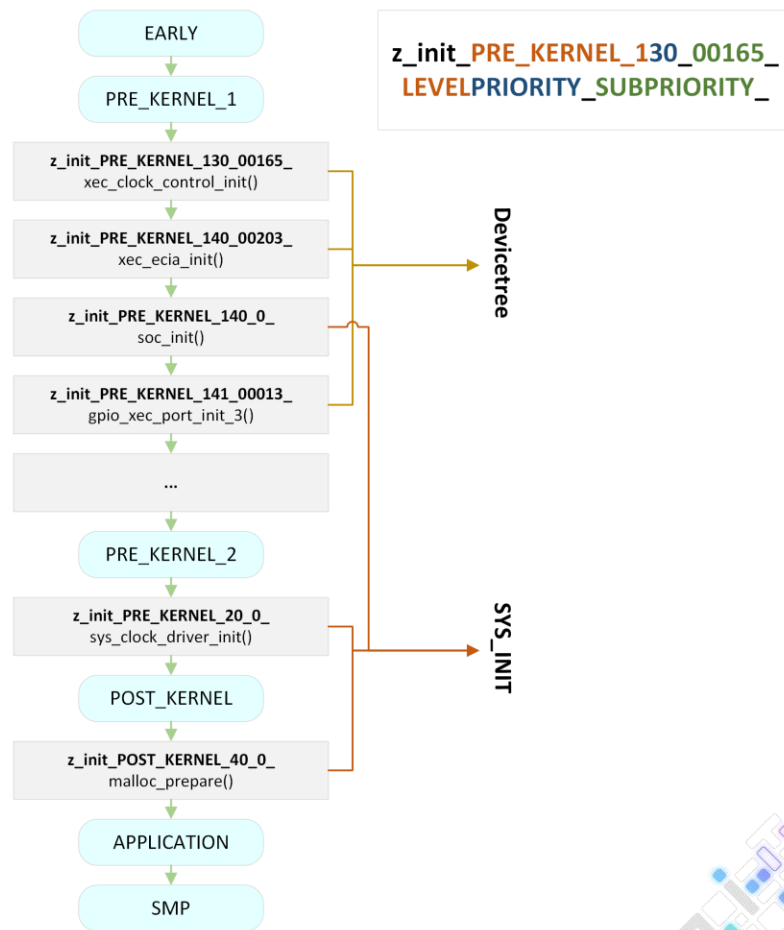
A look into the near future

- Device initialization order comes from devicetree
  - o There's also a priority, so GPIO can be initialized after interrupts, for instance
- Things not on devicetree – calling them "services" here – are initialized via SYS\_INIT
- Initialization level and priority inside that level are used to order SYS\_INIT calls
- Numbers make it hard to track dependencies
  - o Usually, priorities are tweaked after something doesn't work



# How it is today

- "Init entries" containing the device and a function to perform initialization are created
- Sorted by LEVEL, PRIORITY and SUBPRIORITY
  - `KEEP(*(SORT(.z_init_<LEVEL>?_*)))`;
- SUBPRIORITY keeps different instances of the same device in some order
- ``west build -t initilevels`` to see the order





# A proposal

Get rid of numbers, track dependencies by name


- Read info during build and generate a linker script snippet with entries topologically sorted
- Code defined dependencies
- Kconfig allows dependencies to be overwritten
- Integrate into current model, no need to rewrite the world

```
/* Code defined dependencies */
ZERVICE_DEFINE(soc_init, "microchip,xec-gpio-v2",
"microchip,xec-ecia");

# Kconfig allows dependencies to be overwritten
config ZERVICE_SOC_INIT_AFTER
    string "SoC initialization after"
    default "microchip,xec-ecia"

config ZERVICE_SOC_INIT_BEFORE
    string "SoC initialization before"
    default "microchip,xec-gpio-v2"

# prj.conf overwriting a dependency
CONFIG_ZERVICE_SOC_INIT_BEFORE=
"microchip,xec-gpio-v2#4"
```


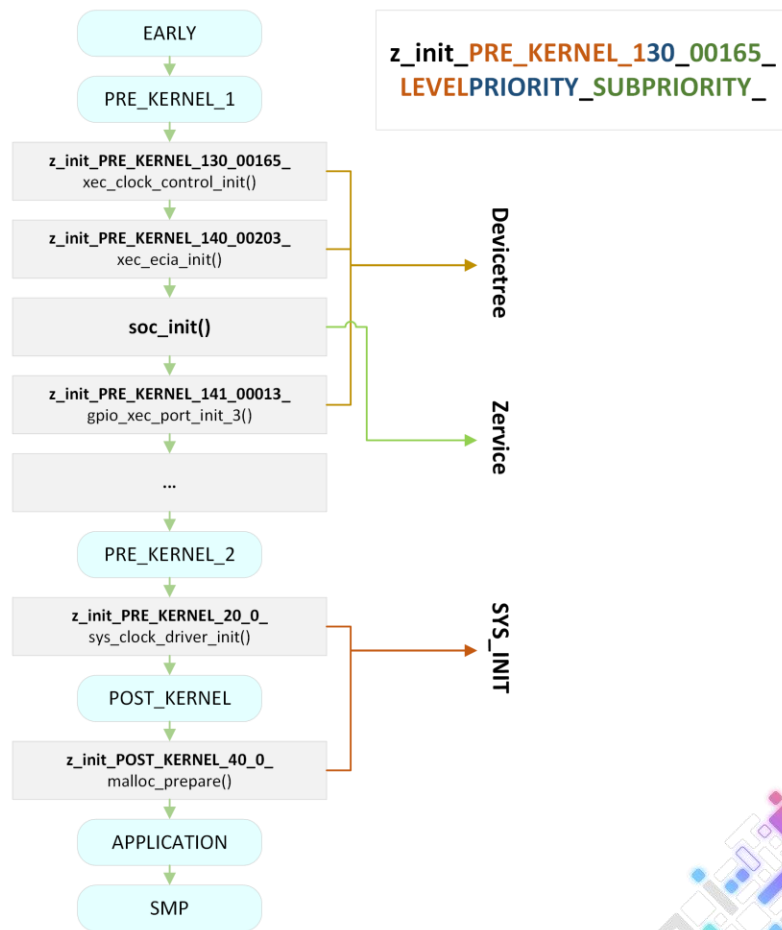


# ZERVICE?



# How it becomes (I)


- "Init entries" are still created for devices, SYS\_INIT and Zervices
- Devices and SYS\_INIT ones are still sorted as of today in the *first build step*
- A new build script reads the first build step ELF map, Kconfig and devicetree info to find the place for the Zervices
- A linker script snippet is generated with all "init entries" in their final ordering, which is used in the *next build step*



## How it becomes (II)


- One can even use current levels as dependencies
  - `ZERVICE_DEFINE(some_service, "APPLICATION", "i2c_init;PRE_KERNEL_2")`
- While seamless integration with current scheme is helpful, expectation is that eventually everything moves to the dependence-based mechanism
- It is possible to depend on a specific instance of a device, but currently this is done using the arbitrary devicetree instance number
  - Not good, as it doesn't relate with anything (#3 device may come before #0)
  - Device ordinal doesn't seem much better

```
/* Generated snippet sample */
__init_EARLY_start = .;
__init_PRE_KERNEL_1_start = .;
KEEP(*( .z_init_PRE_KERNEL_130_00165__init__device_dts_ord_165))
KEEP(*( .z_init_PRE_KERNEL_140_00203__init__device_dts_ord_203))
KEEP(*( .z_zervice_soc_init))
KEEP(*( .z_init_PRE_KERNEL_141_00013__init__device_dts_ord_13))
KEEP(*( .z_init_PRE_KERNEL_141_00014__init__device_dts_ord_14))
(...)
```



# Relevant PR

<https://github.com/zephyrproject-rtos/zephyr/pull/71470>




# Device notifications

A look into the not so near future

- Exploring ideas to notify about devices that failed to initialize – or were initialized late
  - Allow applications to track device state
- Used zbus to deliver notifications
  - Applications then subscribe to channels relating to a topic (device failed) or the device itself
- Why not?
  - Include zbus – but zbus should be light
  - Zbus is fully functional at Application level – is it really a problem?
- Help with another device issues?



[This Photo](#) by Unknown author is licensed under [CC BY-NC-ND](#).



# Device unloading/de-initializing


A look into the not so near future

- Devices can be initialized, why not the opposite?
  - Unloading/closing then on a Zephyr based bootloader before jumping to the next boot phase
- Overlap with Power Management?
  - Independent?
  - Interdependent?
  - One on top of the other? Which one?
- What to do about dependencies?
  - Mandatory and optional
- `device_get()/device_put()`
  - Changes across the board
  - Subtle errors
- Notification
  - Device notifies is going down, those (devices/application) who care subscribe and take appropriate action



# Relevant PR

<https://github.com/zephyrproject-rtos/zephyr/pull/71469>





# Questions? Comments? Suggestions?





# Zephyr<sup>®</sup> Project

## Developer Summit

