



# Ticket spinlocks in Zephyr RTOS: what, where & why

Evgenii Paltsev

Zephyr Developer Summit, 2024

# Evgenii Paltsev

- Sr software engineer in Synopsys
- Develop, port and maintain OSS projects for Synopsys ARC and RISC-V processors architectures
- Maintainer of ARC architecture in Zephyr RTOS
- Main focus - Zephyr RTOS, Linux kernel, U-Boot

# Rationale

- Multicore platform bring up – ARCV3, 8 / 12 cores
- Necessity of running something
  - SMP capable
  - allows to load the system
  - simple (simpler than Linux)

# Rationale

Zephyr is great candidate:

- Mature SMP support – in tree since v1.11
- POSIX subsystem – allows to run existing benchmarks / stress tests with minimal changes
- We plan to support Zephyr anyway



# Co-development

- The HW platform was not yet ready – so tests were prepared and run in simulation
- Synopsys nSIM simulator was used

# Issue – some stress tests fail

What to blame?

- SMP implementation for our architecture
- POSIX subsystem in Zephyr
- Stress tests itself
- Simulator
- Toolchain
- Universe
- ....



# Issue – some stress tests fail

- We have a livelock!
- It was possible due to combination of how simulation works and how Zephyr spinlocks are implemented

# Was this issue unic?

- At the same moment other team was investigating similar issue in Zephyr (but with HW platform)
- GitHub: [zephyrproject-rtos/zephyr#61541](https://github.com/zephyrproject-rtos/zephyr/issues/61541)
- This all ended in addition of spinlock fairness tests and alternative spinlock implementation to Zephyr



# Default Zephyr spinlock implementation

- Introduced with SMP addition in v1.11
- Unchanged since then
- No fairness guarantee

```
struct k_spinlock {
    atomic_t locked;
};

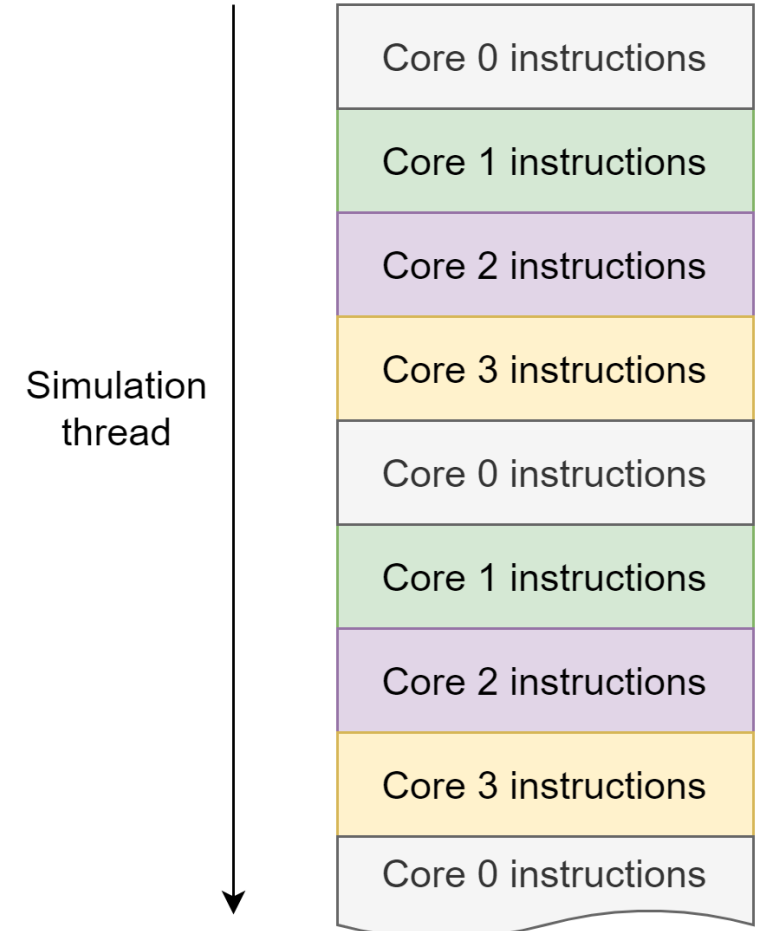
k_spinlock_key_t k_spin_lock(struct k_spinlock *l) {
    k_spinlock_key_t key = arch_irq_lock();

    while (!atomic_cas(&l->locked, 0, 1)) { }
    return key;
}

void k_spin_unlock(struct k_spinlock *l, k_spinlock_key_t key) {
    atomic_clear(&l->locked);
    arch_irq_unlock(key);
}
```

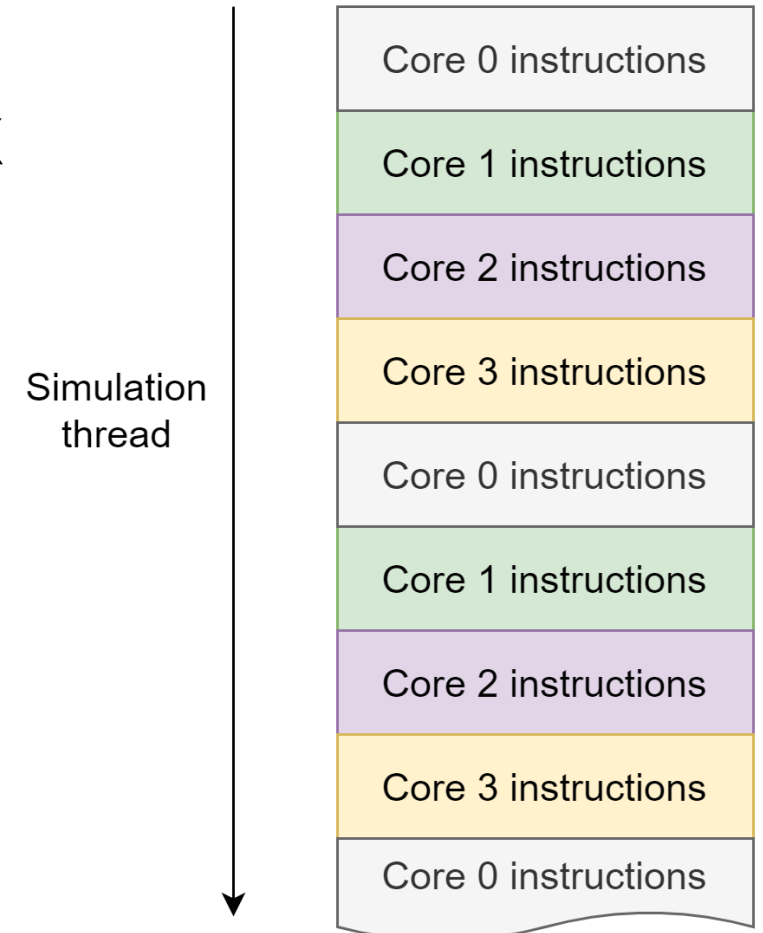
# Our simulator (nSIM)

- Sequentially execute blocks of instructions from each core in single simulation thread
- Simulation is deterministic:
  - Number of instructions in a block is fixed
  - No external events / interrupts due to lack of peripherals in model used



# Issue – our simulator (nSIM)

- We access spinlock from multiple cores
- Spinlock is locked & unlocked multiple times on Core X
- Simulator always switch to another core instructions when spinlock on Core X is locked (coincidence)
- Other cores aren't able to grab spinlock at all
- Not reproduced stably – minor code changes may mask issue



# Spinlock fairness test

- Test case in tests/kernel/spinlock
- Grab – hold – release spinlock in a loop on all cores
- Collect per-core attempts statistics

```
struct k_spinlock lock;

void test_thread(...) {
    int key = arch_irq_lock();
    synchronize_cores();

    do {
        k_spinlock_key_t s_key = k_spin_lock(&lock);

        spinlock_grabbed[core_id]++;
        busy_work();

        k_spin_unlock(&lock, s_key);
    } while (!finished);

    arch_irq_unlock(key);
}
```

# Fairness test - nSIM

- nSIM simulator, ARCV2 architecture, 4 cores (modified nsim\_hs\_smp config)
- Result:

CPU0 acquired spinlock 981 times, expected 1000

CPU1 acquired spinlock 1019 times, expected 1000

CPU2 acquired spinlock 990 times, expected 1000

CPU3 acquired spinlock 1010 times, expected 1000

# What about real SMP HW?


- HW board, ARCV2 architecture, 4 cores (hsdk config)
- What results would we get? Ideal:

CPU0 acquired spinlock 1000 times, expected 1000

CPU1 acquired spinlock 1000 times, expected 1000

CPU2 acquired spinlock 1000 times, expected 1000

CPU3 acquired spinlock 1000 times, expected 1000



# What about real SMP HW?

- HW board, ARCV2 architecture, 4 cores (hsdk config)
- What results would we get? Probable:

CPU0 acquired spinlock 999 times, expected 1000

CPU1 acquired spinlock 1001 times, expected 1000

CPU2 acquired spinlock 999 times, expected 1000

CPU3 acquired spinlock 1001 times, expected 1000

# What about real SMP HW?


- HW board, ARCV2 architecture, 4 cores (hsdk config)
- Actual results:

CPU0 acquired spinlock 0 times, expected 1000

CPU1 acquired spinlock 3997 times, expected 1000

CPU2 acquired spinlock 1 times, expected 1000

CPU3 acquired spinlock 2 times, expected 1000






# What about real SMP HW?

How does this happen?

- CAS – Compare And Swap atomic primitive
- Access '*locked*' var in loop from multiple cores simultaneously
- Our HW has private L1 caches for each core
- Core which owns cache line (in its L1 D\$) with '*locked*' var has more chances for CAS to succeed



```
k_spinlock_key_t k_spin_lock(k_spinlock *l) {  
    ...  
    while (!atomic_cas(&l->locked, 0, 1)) {  
    }  
    ...  
}
```

# Ticket spinlock

- Known algorithm – added to Linux kernel ~ 20 years ago

```
struct k_spinlock {
    atomic_t owner;
    atomic_t ticket;
};

k_spinlock_key_t k_spin_lock(struct k_spinlock *l) {
    k_spinlock_key_t key = arch_irq_lock();

    atomic_val_t ticket = atomic_inc(&l->ticket);
    while (atomic_get(&l->owner) != ticket) { }

    return key;
}

void k_spin_unlock(struct k_spinlock *l, k_spinlock_key_t key) {
    atomic_inc(&l->owner);
    arch_irq_unlock(key);
}
```

# Ticket spinlock

unlocked  
ticket = 0  
owner = 0

Spinlock state

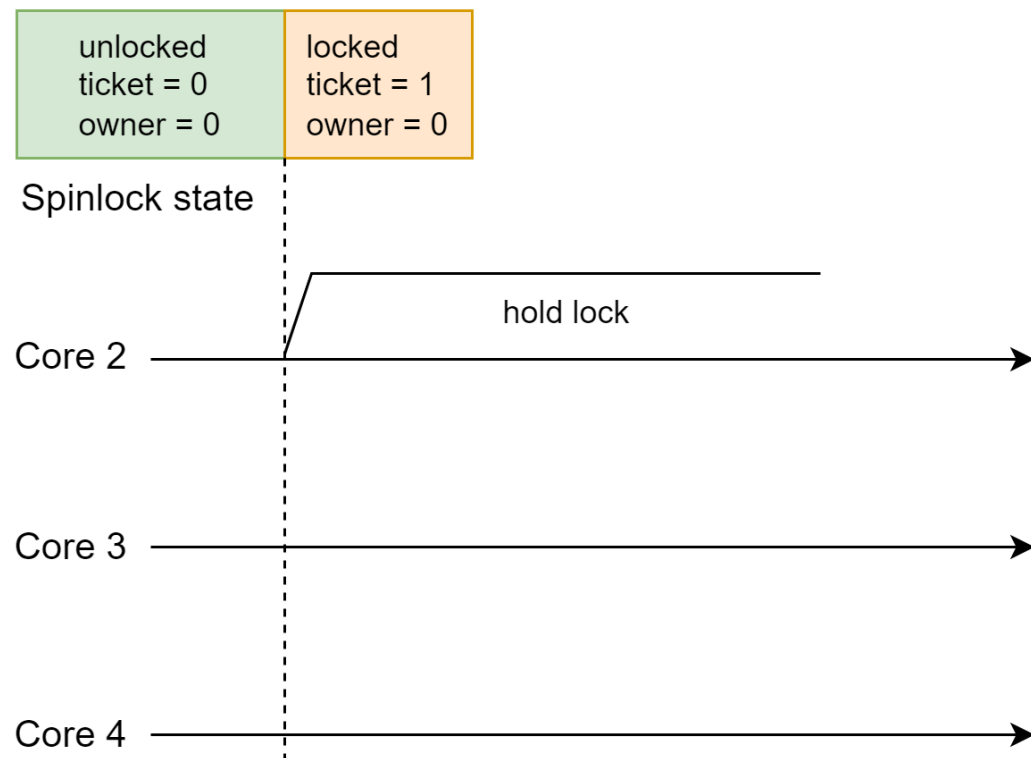
Core 2 →

Core 3 →

Core 4 →

```
struct k_spinlock {  
    atomic_t owner;  
    atomic_t ticket;  
};  
  
k_spinlock_key_t k_spin_lock(struct k_spinlock *l) {  
    k_spinlock_key_t key = arch_irq_lock();  
  
    atomic_val_t ticket = atomic_inc(&l->ticket);  
    while (atomic_get(&l->owner) != ticket) { }  
  
    return key;  
}  
  
void k_spin_unlock(struct k_spinlock *l, k_spinlock_key_t key) {  
    atomic_inc(&l->owner);  
    arch_irq_unlock(key);  
}
```

# Ticket spinlock



```
struct k_spinlock {
    atomic_t owner;
    atomic_t ticket;
};

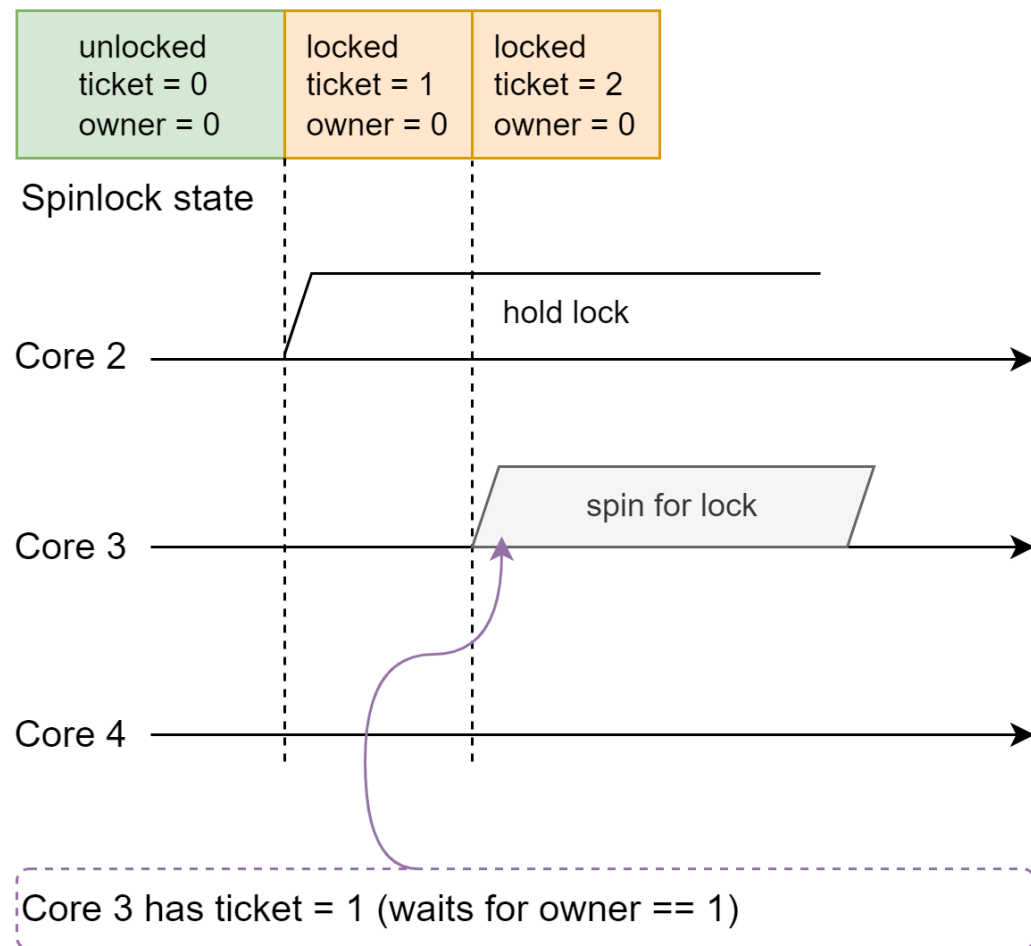
k_spinlock_key_t k_spin_lock(struct k_spinlock *l) {
    k_spinlock_key_t key = arch_irq_lock();

    atomic_val_t ticket = atomic_inc(&l->ticket);
    while (atomic_get(&l->owner) != ticket) { }

    return key;
}

void k_spin_unlock(struct k_spinlock *l, k_spinlock_key_t key) {
    atomic_inc(&l->owner);
    arch_irq_unlock(key);
}
```

# Ticket spinlock



```
struct k_spinlock {
    atomic_t owner;
    atomic_t ticket;
};

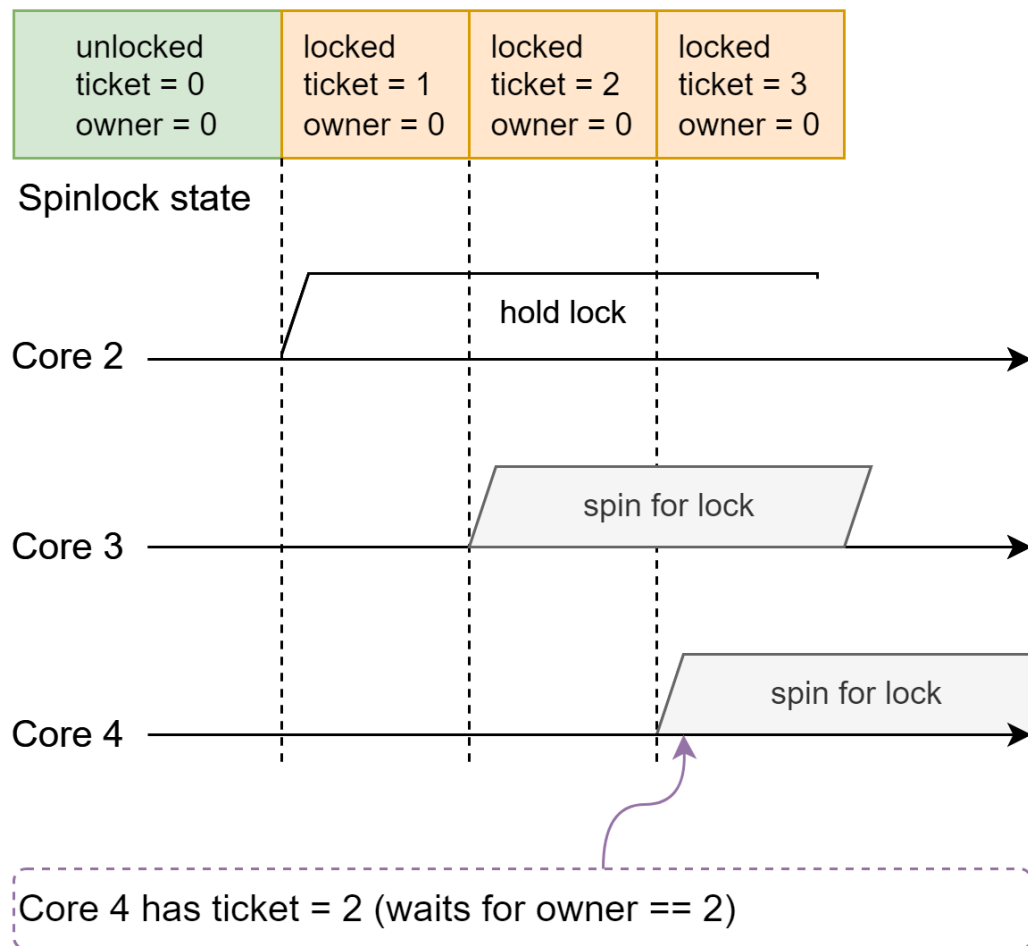
k_spinlock_key_t k_spin_lock(struct k_spinlock *l) {
    k_spinlock_key_t key = arch_irq_lock();

    atomic_val_t ticket = atomic_inc(&l->ticket);
    while (atomic_get(&l->owner) != ticket) { }

    return key;
}

void k_spin_unlock(struct k_spinlock *l, k_spinlock_key_t key) {
    atomic_inc(&l->owner);
    arch_irq_unlock(key);
}
```

# Ticket spinlock



```
struct k_spinlock {
    atomic_t owner;
    atomic_t ticket;
};

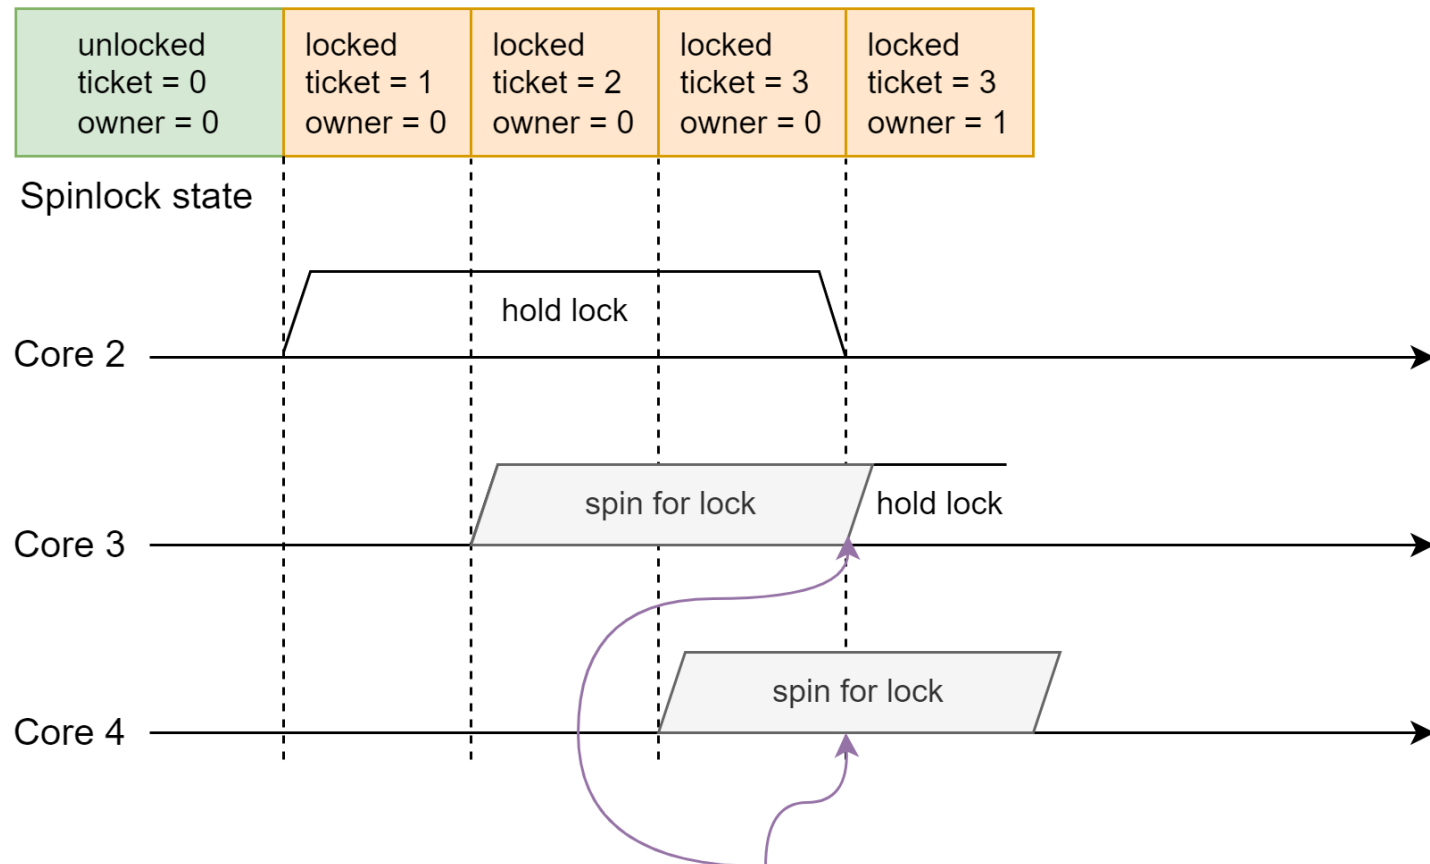
k_spinlock_key_t k_spin_lock(struct k_spinlock *l) {
    k_spinlock_key_t key = arch_irq_lock();

    atomic_val_t ticket = atomic_inc(&l->ticket);
    while (atomic_get(&l->owner) != ticket) { }

    return key;
}

void k_spin_unlock(struct k_spinlock *l, k_spinlock_key_t key) {
    atomic_inc(&l->owner);
    arch_irq_unlock(key);
}
```

# Ticket spinlock



Core 3 has ticket = 1 (waits for owner == 1), Core 4 has ticket = 2 (waits for owner == 2)

```
struct k_spinlock {
    atomic_t owner;
    atomic_t ticket;
};

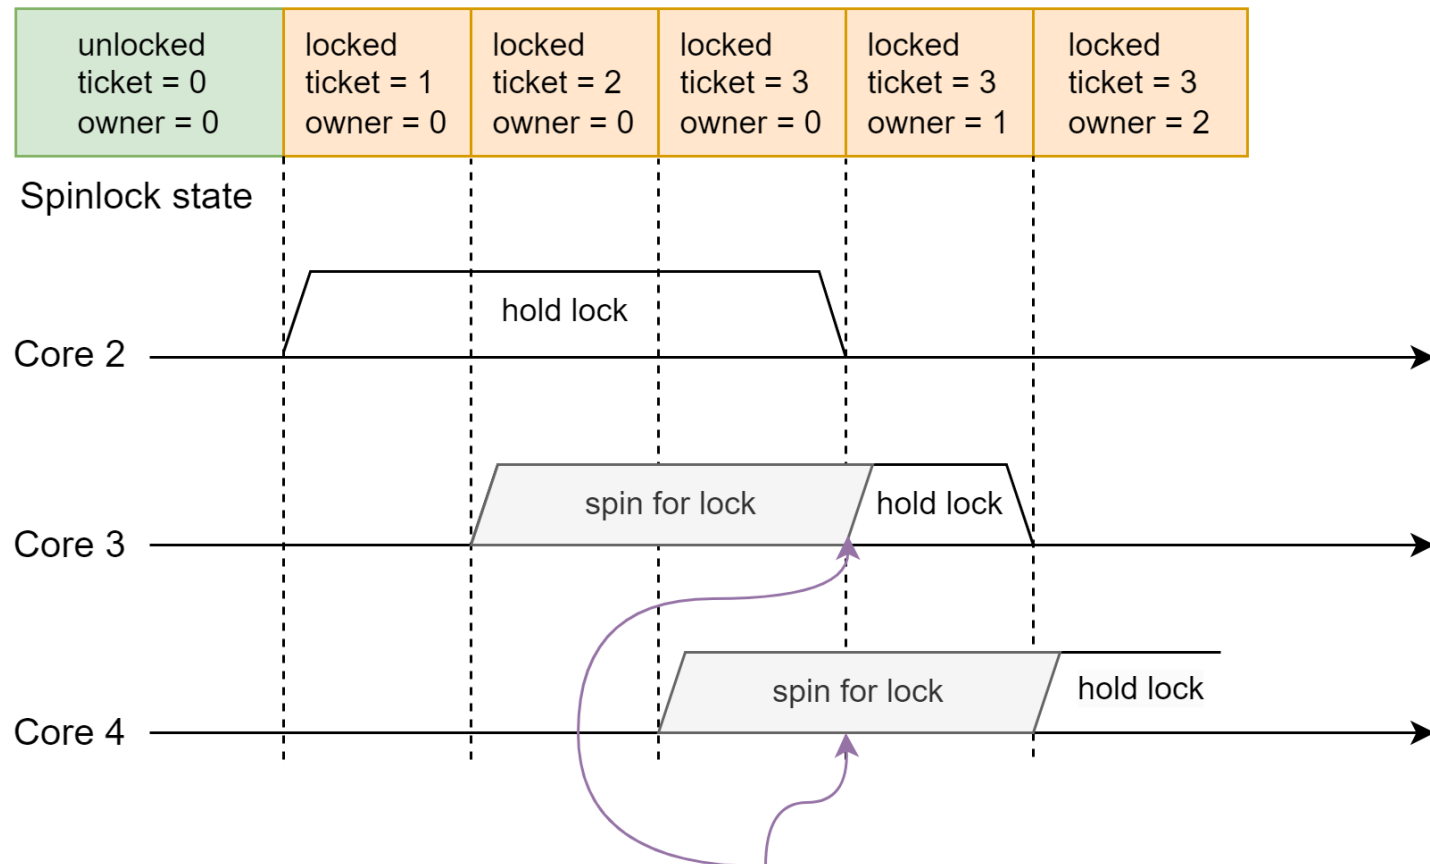
k_spin_lock(k_spinlock_t *l) {
    key = arch_irq_lock();

    ticket = atomic_inc(&l->ticket);
    while (atomic_get(&l->owner) != ticket)
    { }

    return key;
}

k_spin_unlock(k_spinlock_t *l, key_t key) {
    atomic_inc(&l->owner);
    arch_irq_unlock(key);
}
```

# Ticket spinlock



Core 3 has ticket = 1 (waits for owner == 1), Core 4 has ticket = 2 (waits for owner == 2)

```
struct k_spinlock {
    atomic_t owner;
    atomic_t ticket;
};

k_spin_lock(k_spinlock_t *l) {
    key = arch_irq_lock();

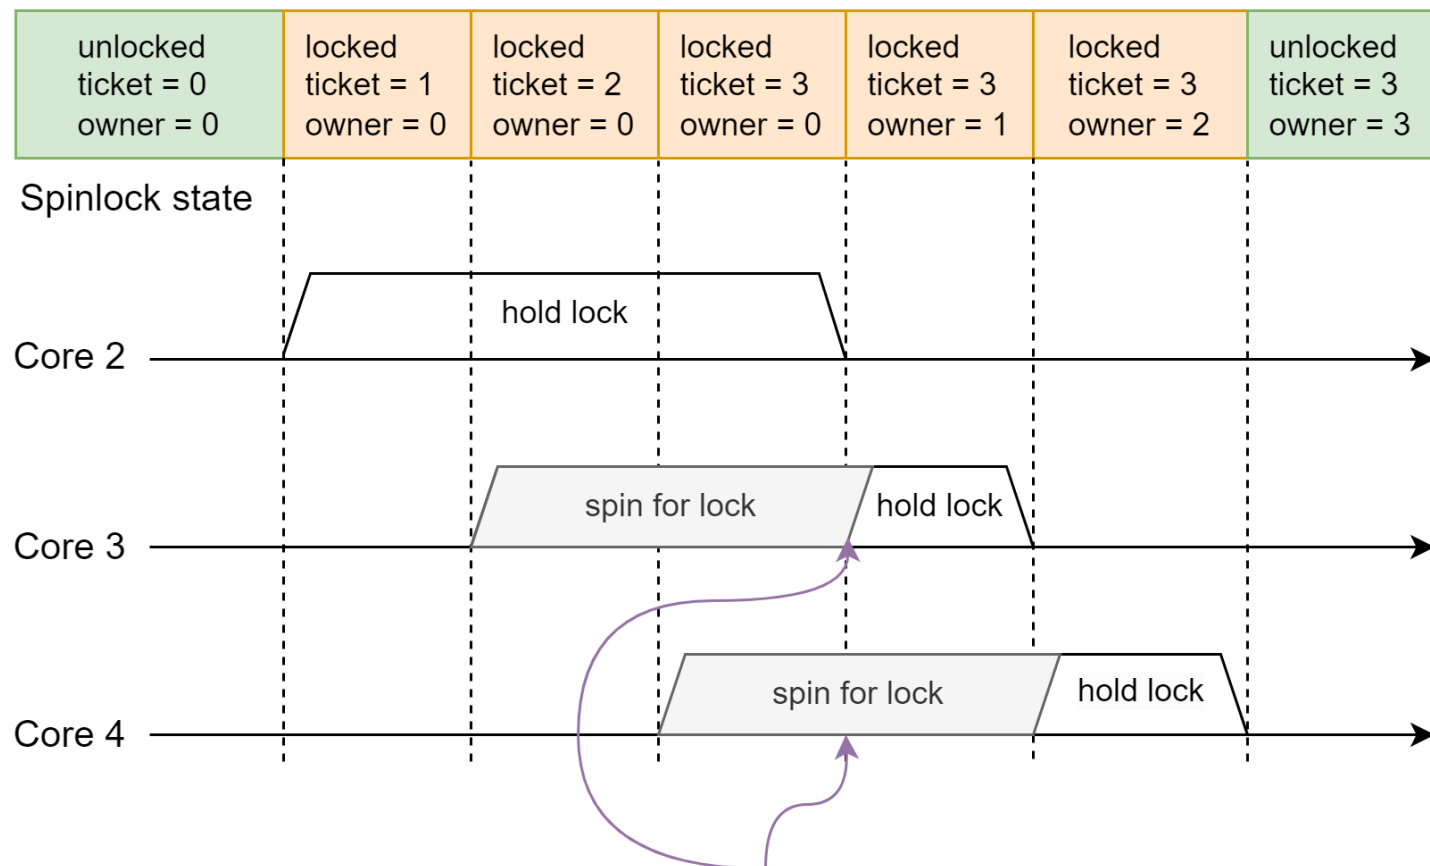
    ticket = atomic_inc(&l->ticket);
    while (atomic_get(&l->owner) != ticket)
    { }

    return key;
}

k_spin_unlock(k_spinlock_t *l, key_t key) {
    atomic_inc(&l->owner);
    arch_irq_unlock(key);
}
```



# Ticket spinlock



Core 3 has ticket = 1 (waits for owner == 1), Core 4 has ticket = 2 (waits for owner == 2)

```
struct k_spinlock {
    atomic_t owner;
    atomic_t ticket;
};

k_spin_lock(k_spinlock_t *l) {
    key = arch_irq_lock();

    ticket = atomic_inc(&l->ticket);
    while (atomic_get(&l->owner) != ticket)
    { }

    return key;
}

k_spin_unlock(k_spinlock_t *l, key_t key) {
    atomic_inc(&l->owner);
    arch_irq_unlock(key);
}
```

# Fairness test with ticket spinlock - nSIM

- nSIM simulator, ARCV2 architecture, 4 cores (modified nsim\_hs\_smp config)
- Works nicely:

CPU0 acquired spinlock 1000 times, expected 1000

CPU1 acquired spinlock 1000 times, expected 1000

CPU2 acquired spinlock 1000 times, expected 1000

CPU3 acquired spinlock 1000 times, expected 1000

# Fairness test with ticket spinlock - HW

- HW board, ARCV2 architecture, 4 cores (hsdk config)
- Everything is OK, results:

CPU0 acquired spinlock 1000 times, expected 1000

CPU1 acquired spinlock 1000 times, expected 1000

CPU2 acquired spinlock 1000 times, expected 1000

CPU3 acquired spinlock 1000 times, expected 1000

# Fairness test with ticket spinlock - QEMU

- QEMU, ARM cortex A53 architecture, 4 cores (modified qemu\_cortex\_a53\_smp config)
- Everything is OK, results:

CPU0 acquired spinlock 1000 times, expected 1000

CPU1 acquired spinlock 1000 times, expected 1000

CPU2 acquired spinlock 1000 times, expected 1000

CPU3 acquired spinlock 1000 times, expected 1000

# Fairness test with ticket spinlock - QEMU

- But what if our host is a bit busy?

```
for i in $(seq $(($(getconf _NPROCESSORS_ONLN) - 2))); do
    cat /dev/random > /dev/null &
done
```

```
1  [|||||100.0%]    4  [|||||100.0%]    7  [|||||100.0%]    10 [|||    19.9%]
2  [|||||91.1%]    5  [|||||100.0%]    8  [|||    12.3%]    11 [|||||100.0%]
3  [|||||95.6%]    6  [|||||100.0%]    9  [|||    12.4%]    12 [|||||100.0%]
Mem[|||||26.4G/62.5G]
Swp[|||||0K/976M]
Tasks: 343, 3094 thr; 10 running
Load average: 8.79 8.02 5.24
Uptime: 5 days, 06:08:39
```

# Fairness test with ticket spinlock - QEMU

- QEMU, ARM cortex A53 architecture, 4 cores (modified qemu\_cortex\_a53\_smp config)
- Ooops:

CPU0 acquired spinlock 99670 times, expected 100000

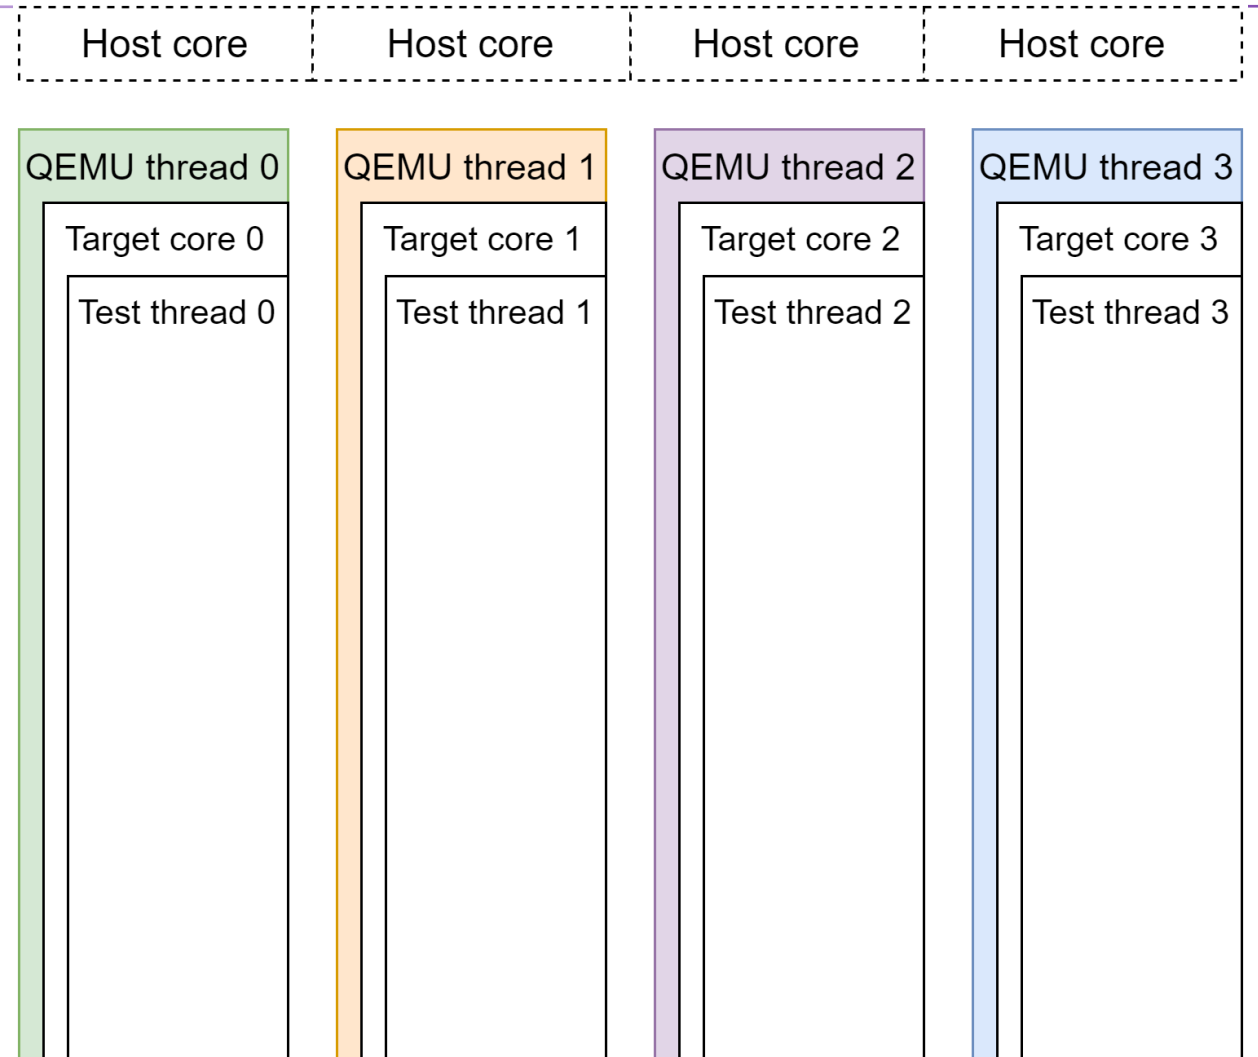
CPU1 acquired spinlock 100344 times, expected 100000

CPU2 acquired spinlock 100264 times, expected 100000

CPU3 acquired spinlock 99722 times, expected 100000

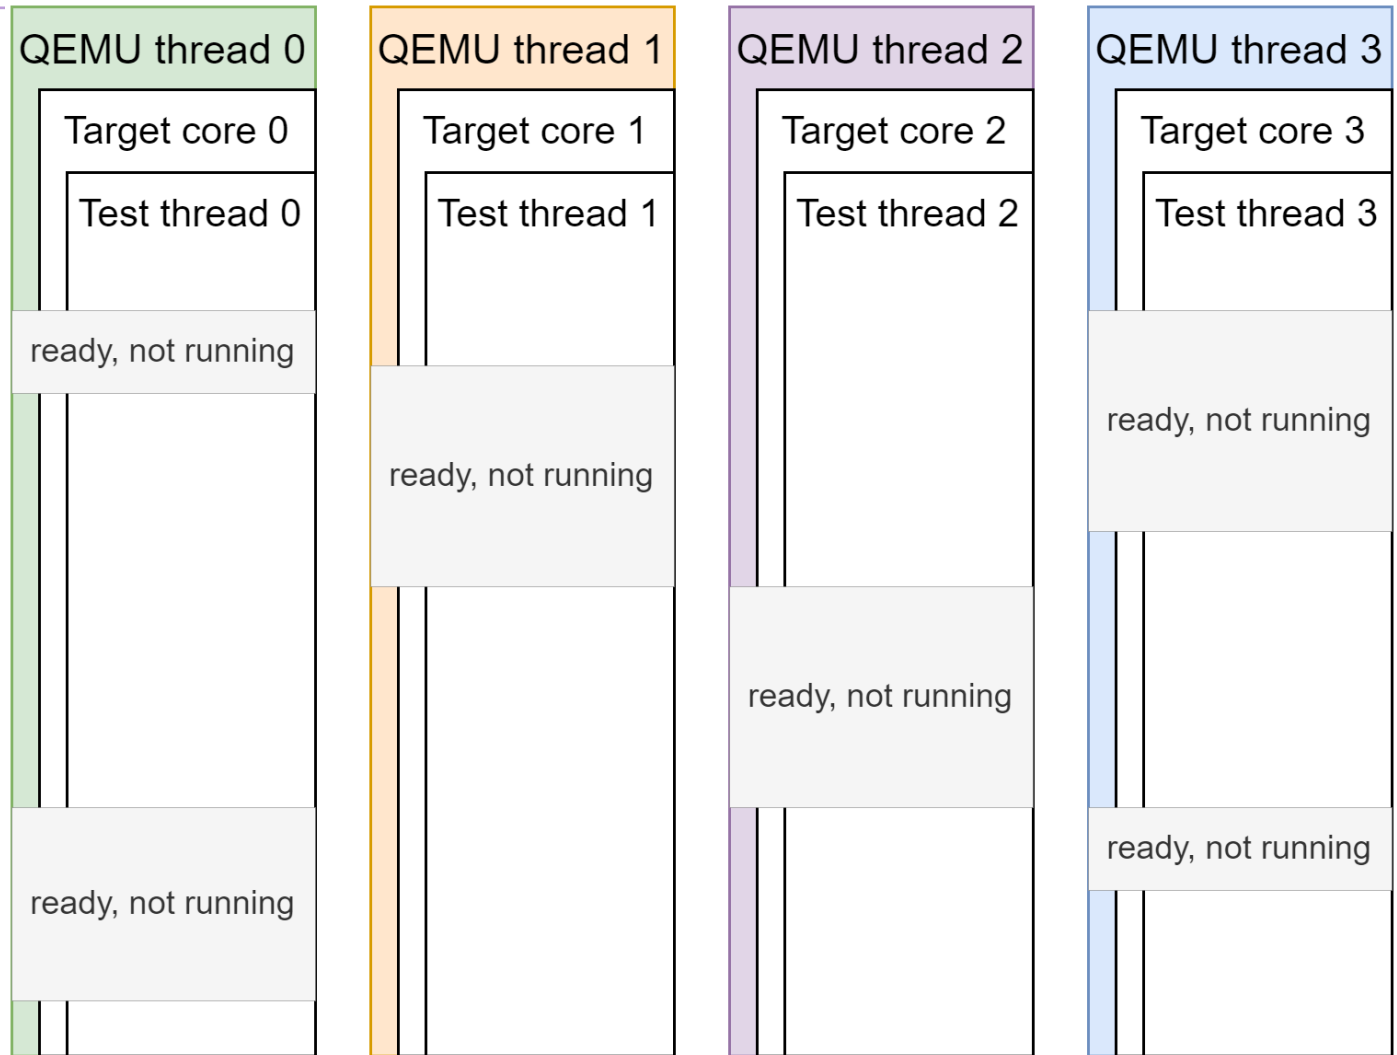
# Fairness test with ticket spinlock - QEMU

- `qemu -smp cpus=4 ...`
- QEMU spawns several threads



# Fairness test with ticket spinlock - QEMU


- High host load - QEMU threads more likely to be scheduled out
- If Zephyr thread request spinlock it'll get it in proper order
- But Zephyr thread may not request it as QEMU thread isn't running





# Ticket spinlock

- We've improved situation with spinlock fairness
- Any drawbacks?



# Ticket spinlock – drawbacks - slowdown

`k_spin_lock()`

- + atomic\_get (memory read, 2 memory barriers)

`k_spin_unlock()`

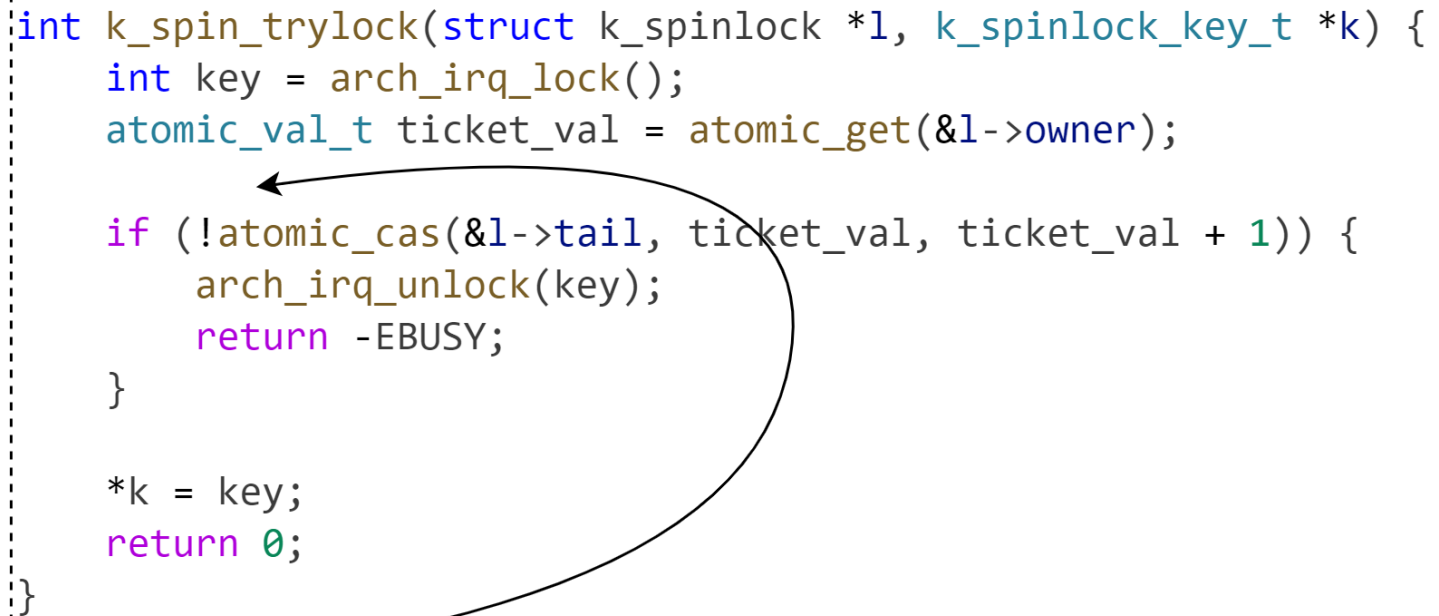
- memory write (in atomic\_clear)

- + atomic RMW (in atomic\_inc)

# Ticket spinlock – drawbacks - wrap around

- Spinlock taken  $0xffffffff + 1$  times
  - requires  $0xffffffff + 2$  number of CPUs
- spinlock taken and released  $0xffffffff$  times and taken again
  - Needs to be done while we execute several instructions in interrupts locked context

```
int k_spin_trylock(struct k_spinlock *l, k_spinlock_key_t *k) {  
    int key = arch_irq_lock();  
    atomic_val_t ticket_val = atomic_get(&l->owner);  
  
    if (!atomic_cas(&l->tail, ticket_val, ticket_val + 1)) {  
        arch_irq_unlock(key);  
        return -EBUSY;  
    }  
  
    *k = key;  
    return 0;  
}
```



# Ticket spinlock - wrap around

- Why it may happen

# Ticket spinlock - wrap around

- Potential solution – use double atomic primitives
  - Potential solution – store ticket and owner in halves of one atomic variable
  - Instead of existing `atomic_inc()` we need to implement our own increments based on CAS
- 
- Not (yet) implemented in Zephyr

Thanks!  
Questions?

# Contacts

## **Evgenii Paltsev**

[Eugeniy.Paltsev@synopsys.com](mailto:Eugeniy.Paltsev@synopsys.com)

[PaltsevEvgeniy@gmail.com](mailto:PaltsevEvgeniy@gmail.com)

<https://www.linkedin.com/in/evgeniy-paltsev>

# Images licensing notes

- The slides [4-6, 14, 16, 33, 36, 37] include images which are based on or includes content from xkcd.com.
- Content from xkcd.com is licensed under the Creative Commons Attribution-NonCommercial 2.5 license